

# شبکه های مولد تخاصمی

# Generative Adversarial Networks (GAN)



CLASS.  
VISION

AI & DEEP LEARNING COURSES

Alireza Akhavanpour

[Akhavanpour.ir](http://Akhavanpour.ir)  
[CLASS.VISION](http://CLASS.VISION)



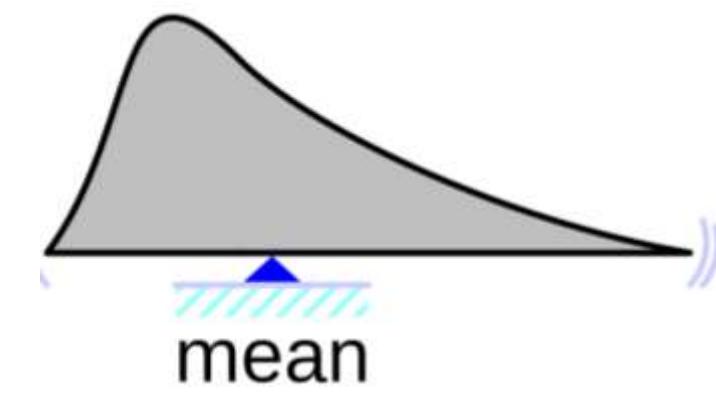
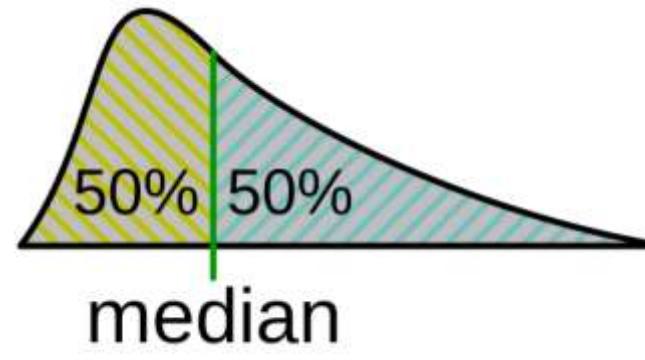
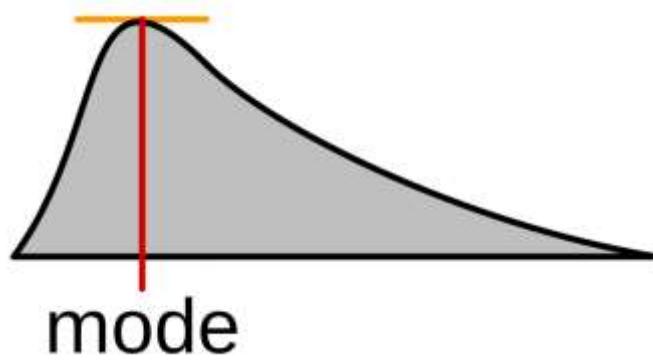
# Mode Collapse

# Outline

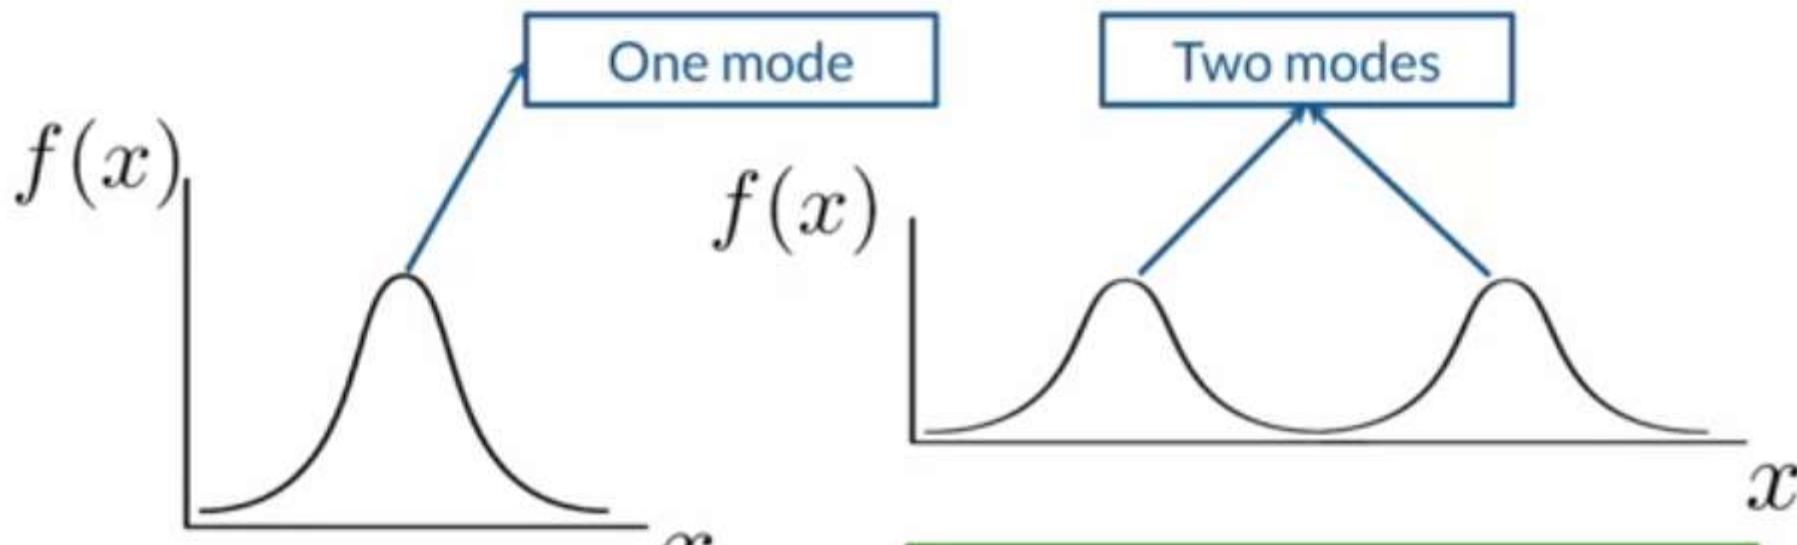
- Modes in distributions
- Mode collapse in GANS
- Intuition behind it during training

# Mode

The **mode** is the value that appears most often in a set of data values. ([Wikipedia](#))



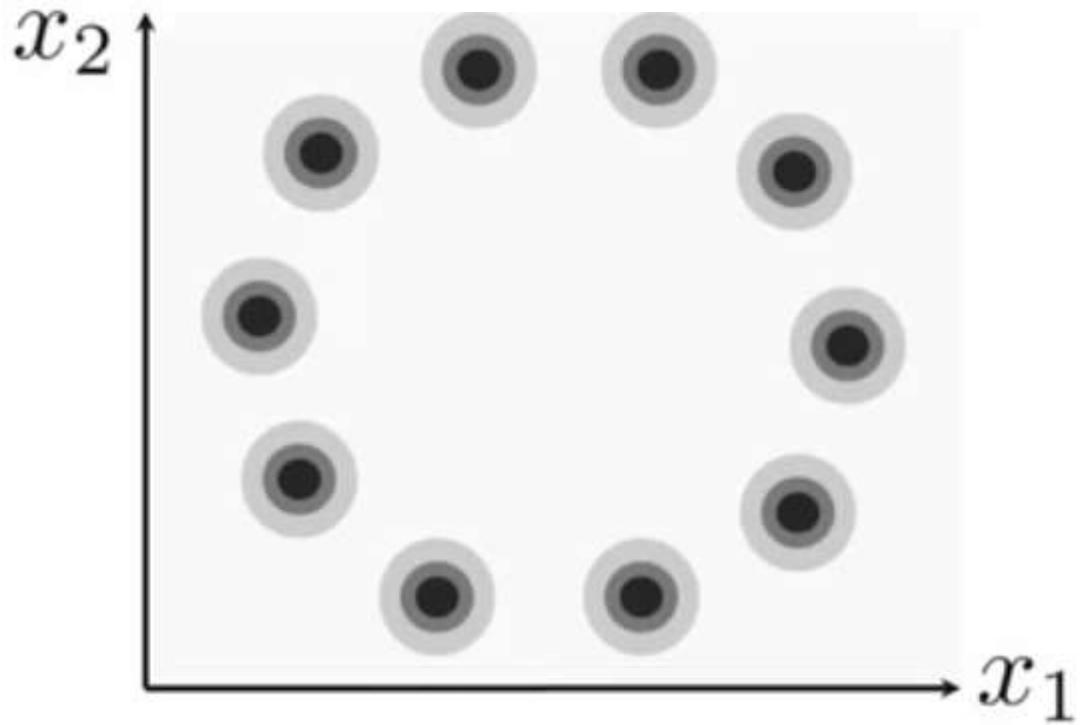
# Mode collapse



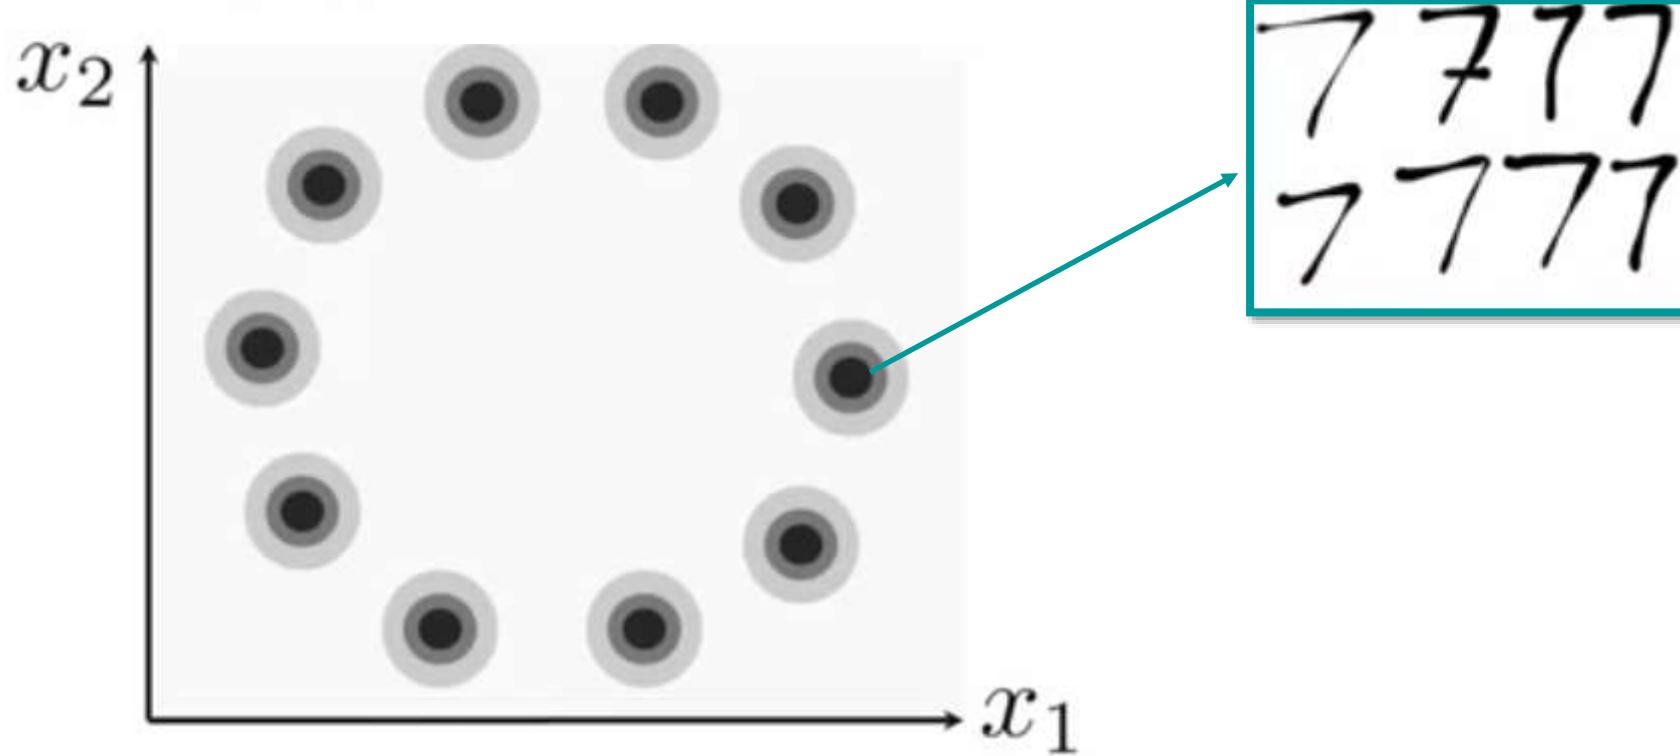
# Mode collapse



# Mode collapse



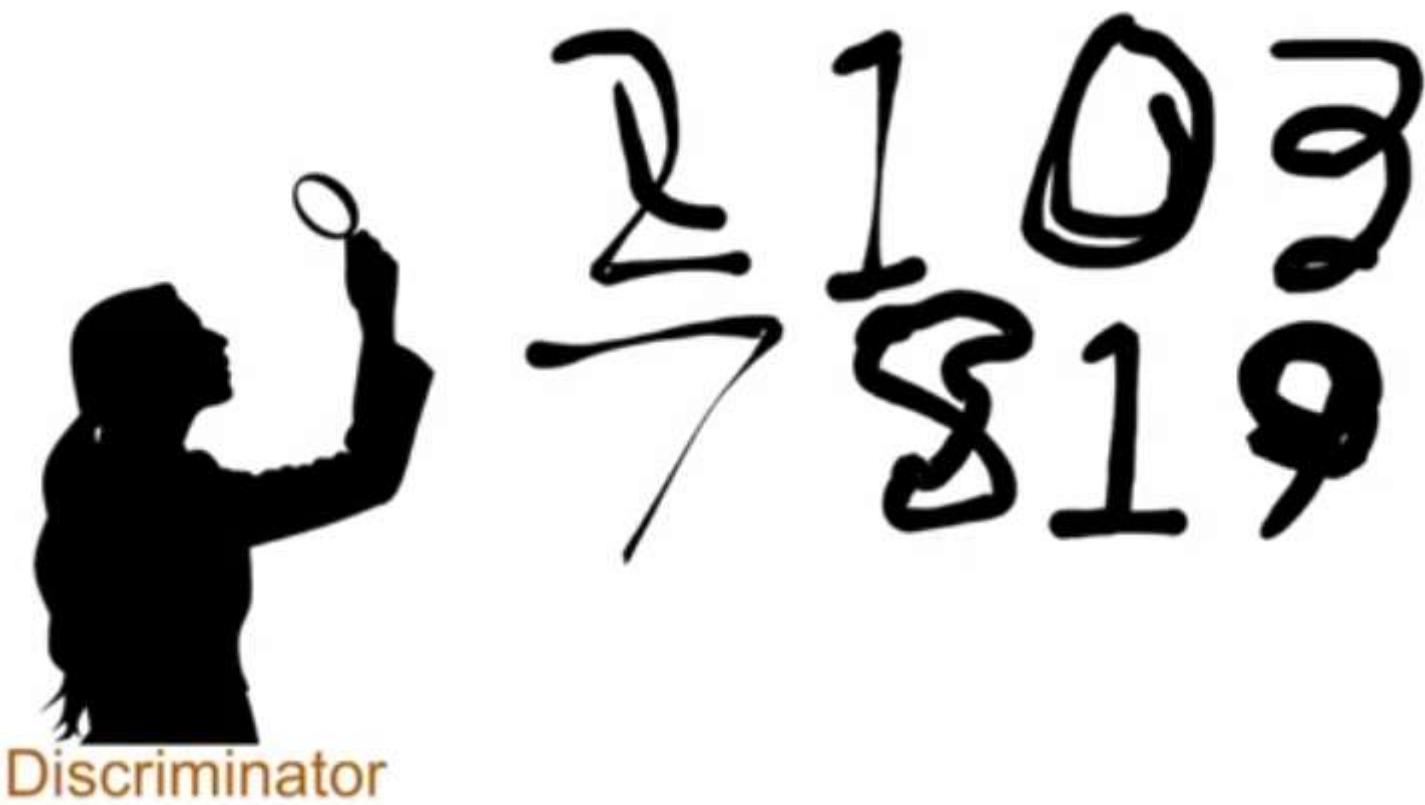
# Mode collapse



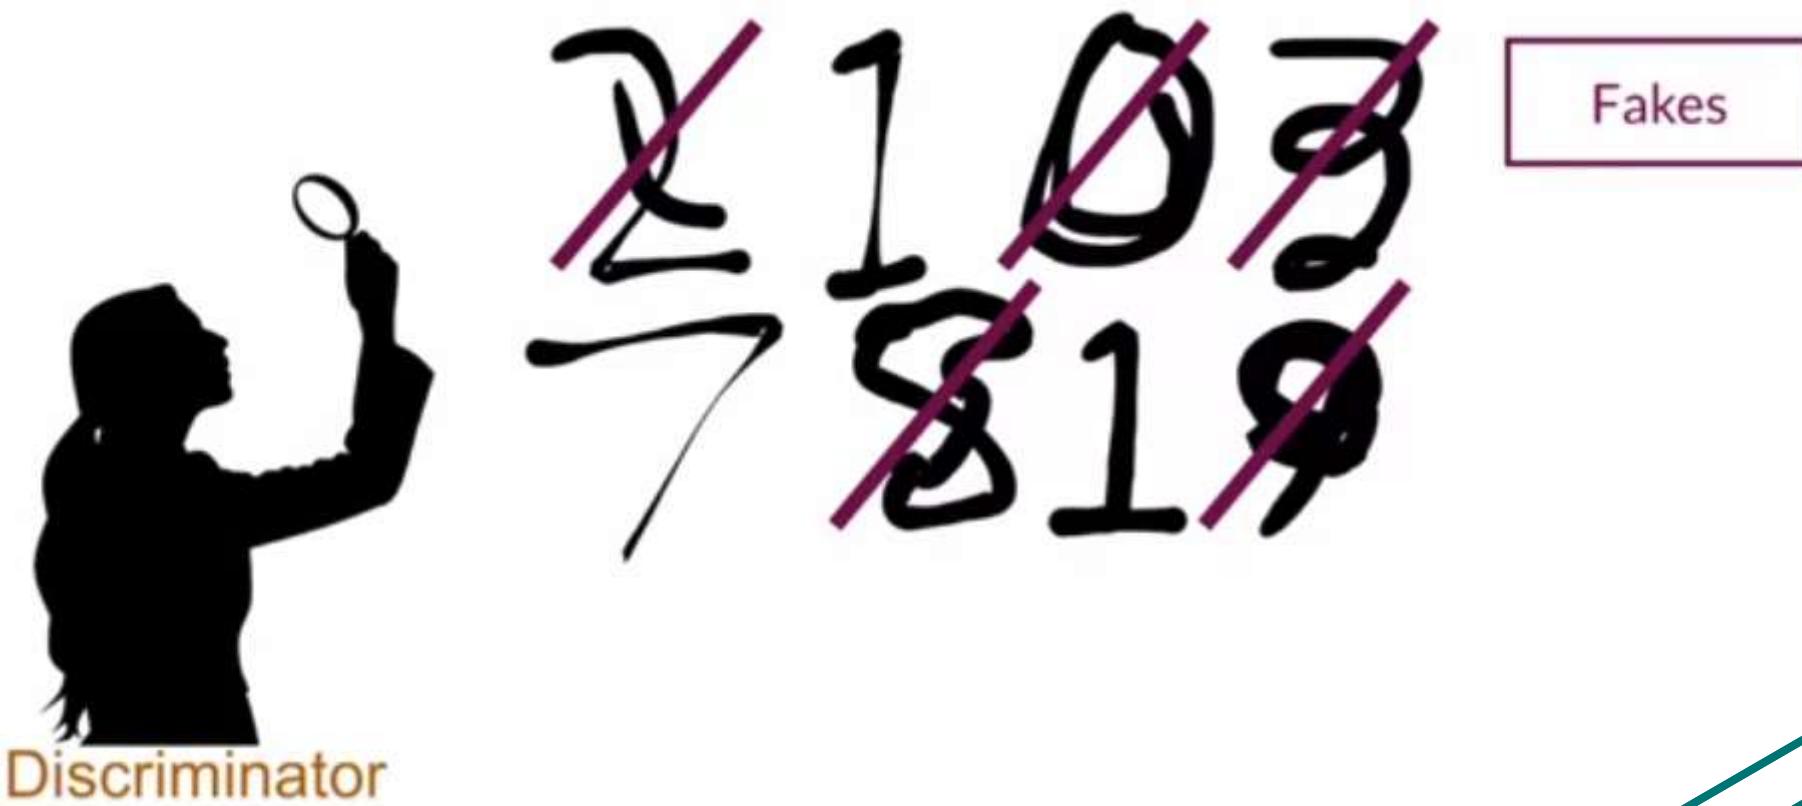
# Mode collapse

۲۱۰۳  
۷۸۱۹

# Mode collapse



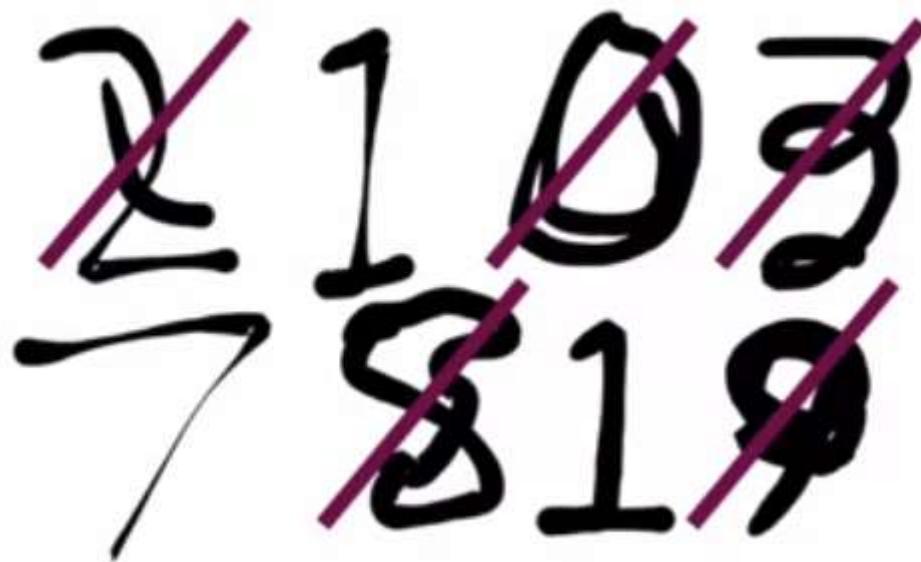
# Mode collapse



# Mode collapse

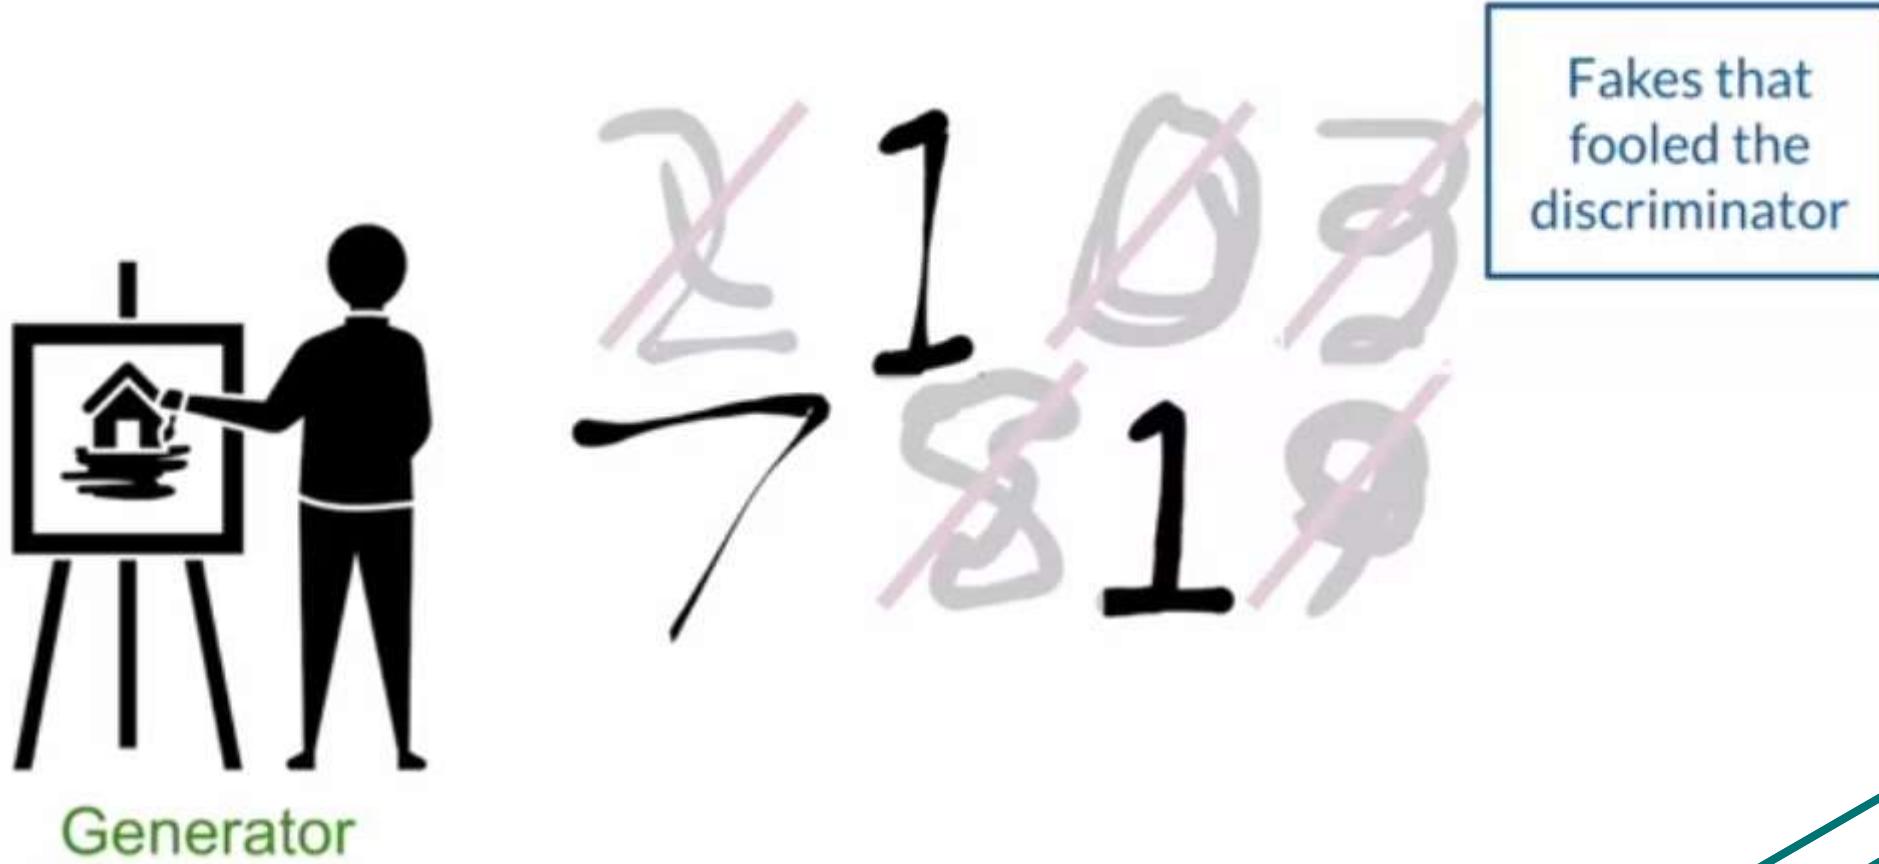


Generator



Fakes

# Mode collapse



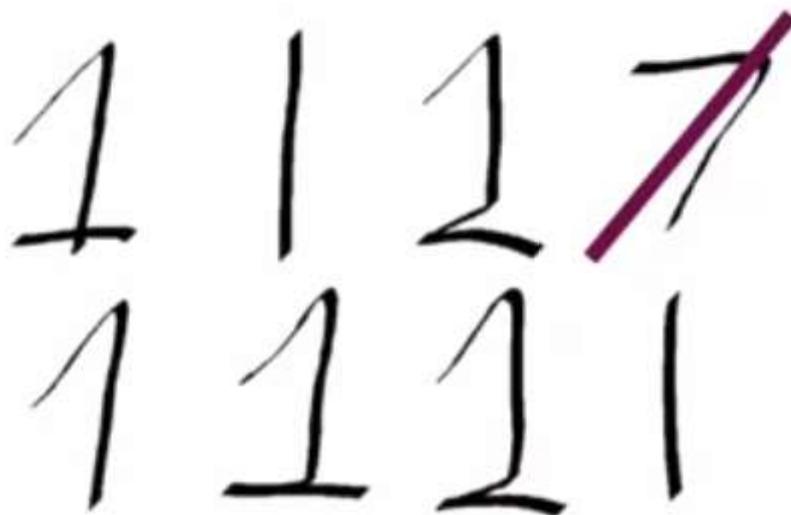
# Mode collapse



Generator

1 1 1 7  
1 1 1 1

# Mode collapse

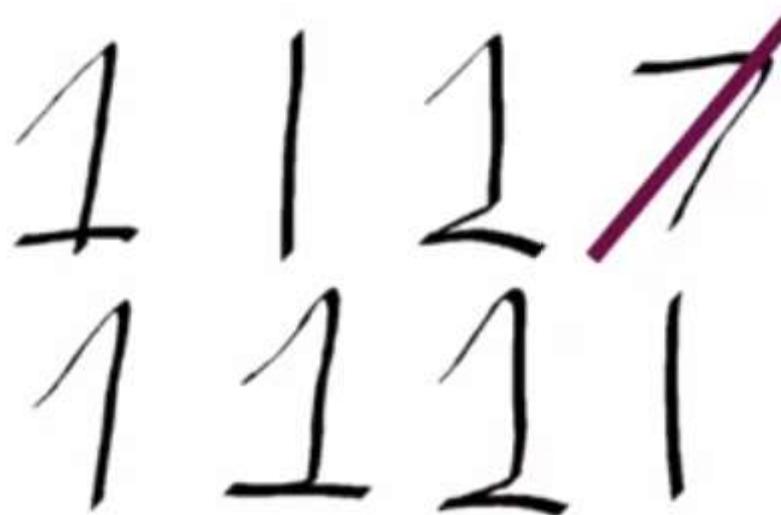


CLASS.  
vision

# Mode collapse



Generator



Fakes

# Mode collapse



Generator

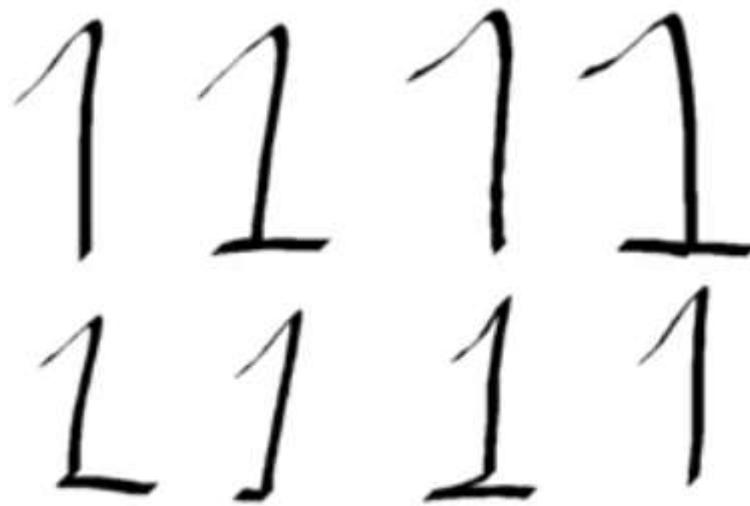


Fakes that  
fooled the  
discriminator

# Mode collapse



Generator



# Summary

- Modes are peaks in the distribution of features
- Typical with real-world datasets
- Mode collapse happens when the generator gets stuck in one mode

# Problem with BCE Loss

# Outline

- BCE Loss and the end objective in GANs
- Problem with BCE Loss

# BCE Loss in GANs

$$J(\theta) = - \frac{1}{m} \sum_{i=1}^m \left[ y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta)) \right]$$

Prediction      Label      Features  
↓                  ↓                  ↓  
Parameters

# BCE Loss in GANs

$$J(\theta) = - \frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

Prediction

Label

Features

Parameters



Generator

Maximize  
cost

# BCE Loss in GANs

$$J(\theta) = - \frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

Prediction

Label

Features

Parameters



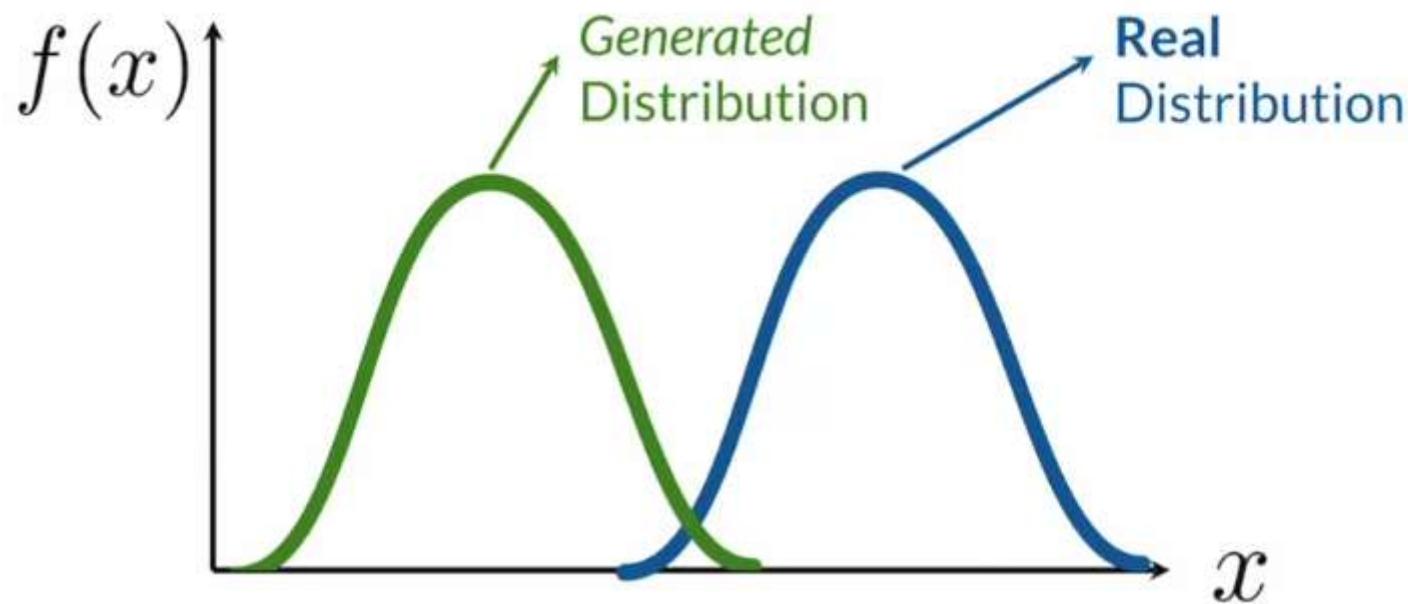
Maximize  
cost



Minimize  
cost

# Objective in GANs

Make the generated and real distributions look similar



# Problem with BCE Loss

Often, the discriminator gets better than the generator



Discriminator

Single output

Easier to train  
than the  
generator



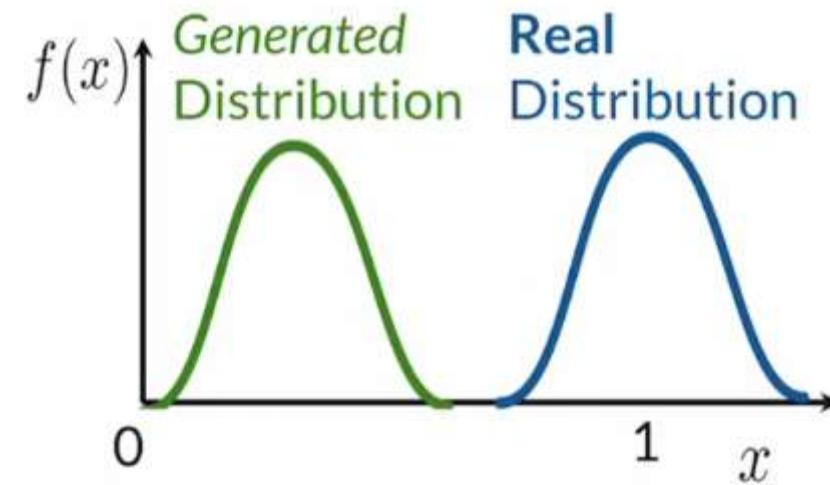
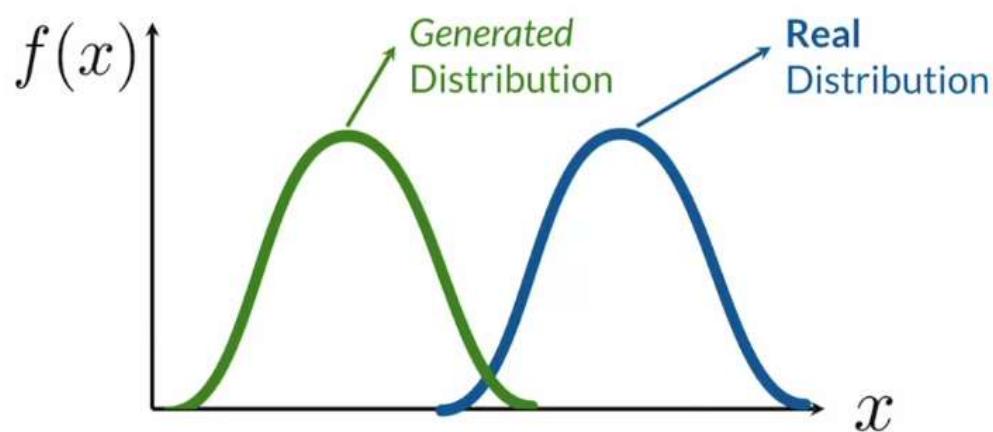
Generator

Complex  
output

Difficult to  
train

# Problem with BCE Loss

- Difference between distributions
- Gradient vanishing!!!



# Summary

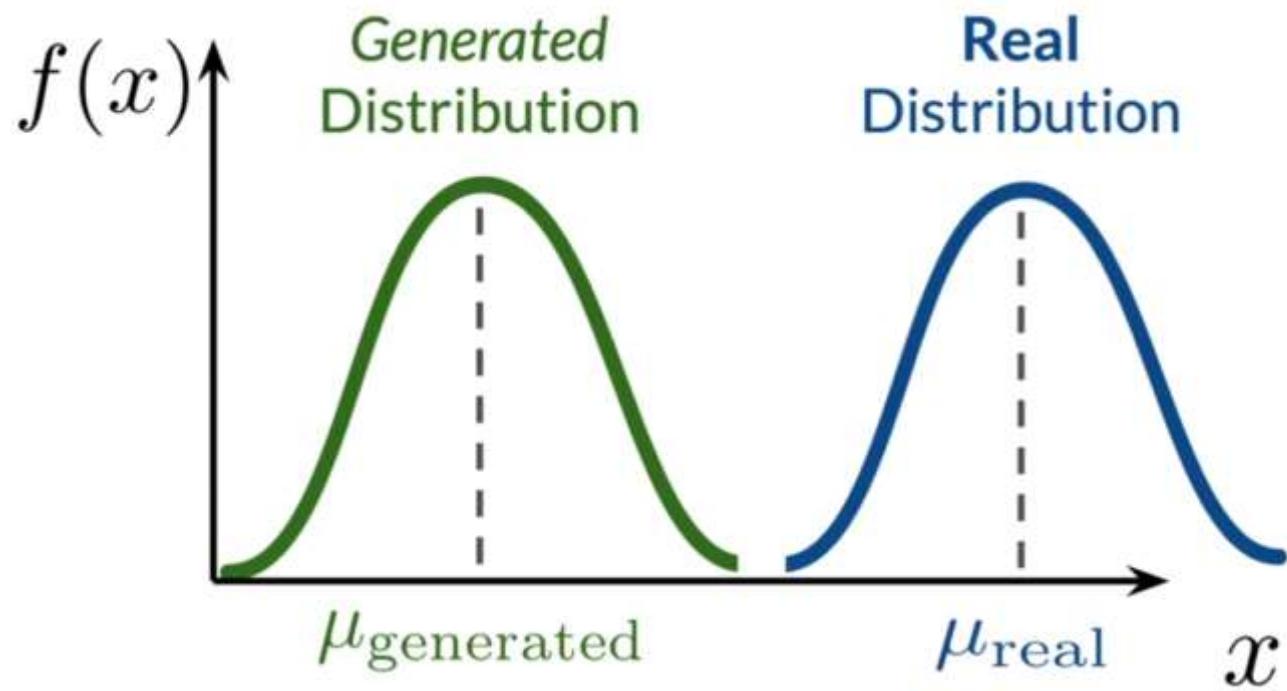
- GANs try to make the real and generated distributions look similar
- When the discriminator improves too much, the function approximated by BCE Loss will contain flat regions
- Flat regions on the cost function = **vanishing gradients**

# Earth Mover's Distance

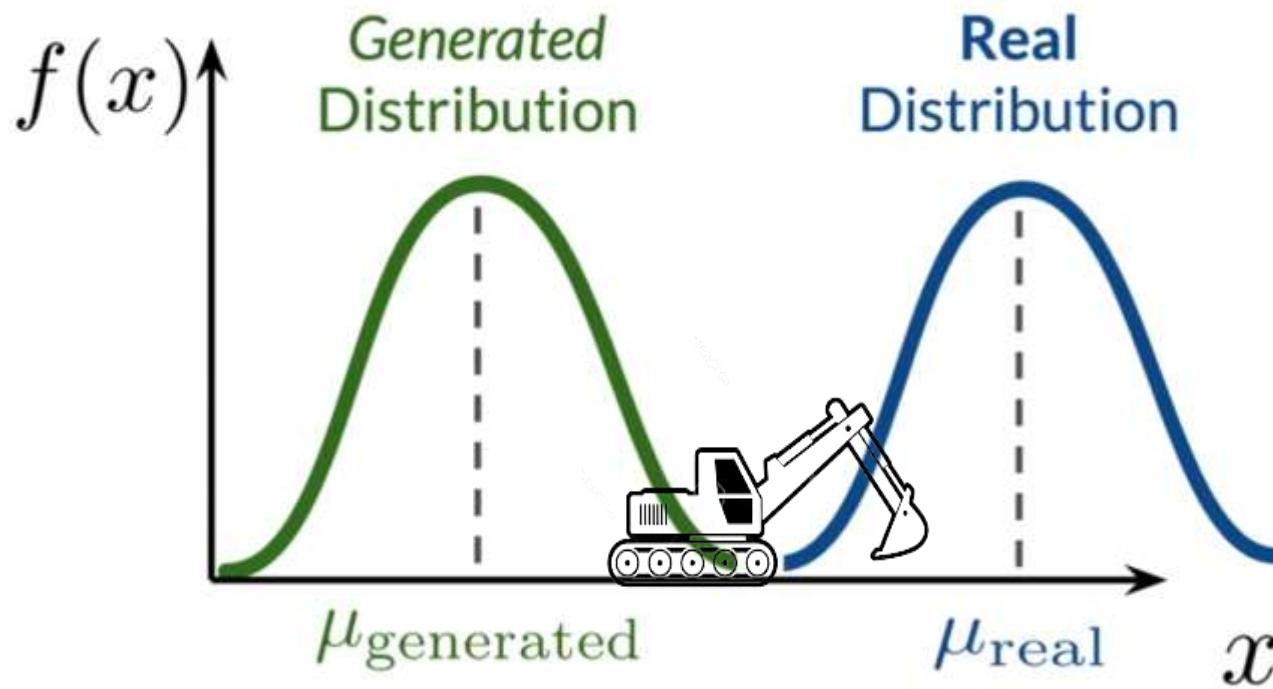
# Outline

- Earth Mover's Distance (EMD)
- Why it solves the vanishing gradient problem of BCE Loss

# Earth Mover's Distance

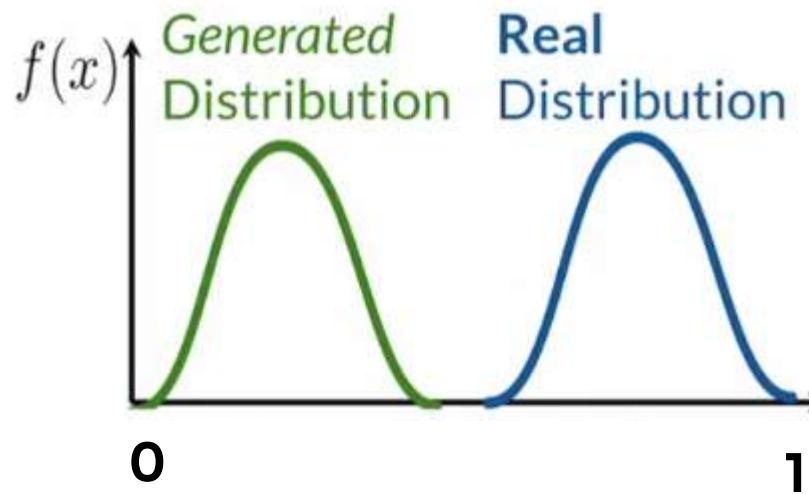


# Earth Mover's Distance



**Effort** to make the generated distribution equal to the real distribution

# Earth Mover's Distance

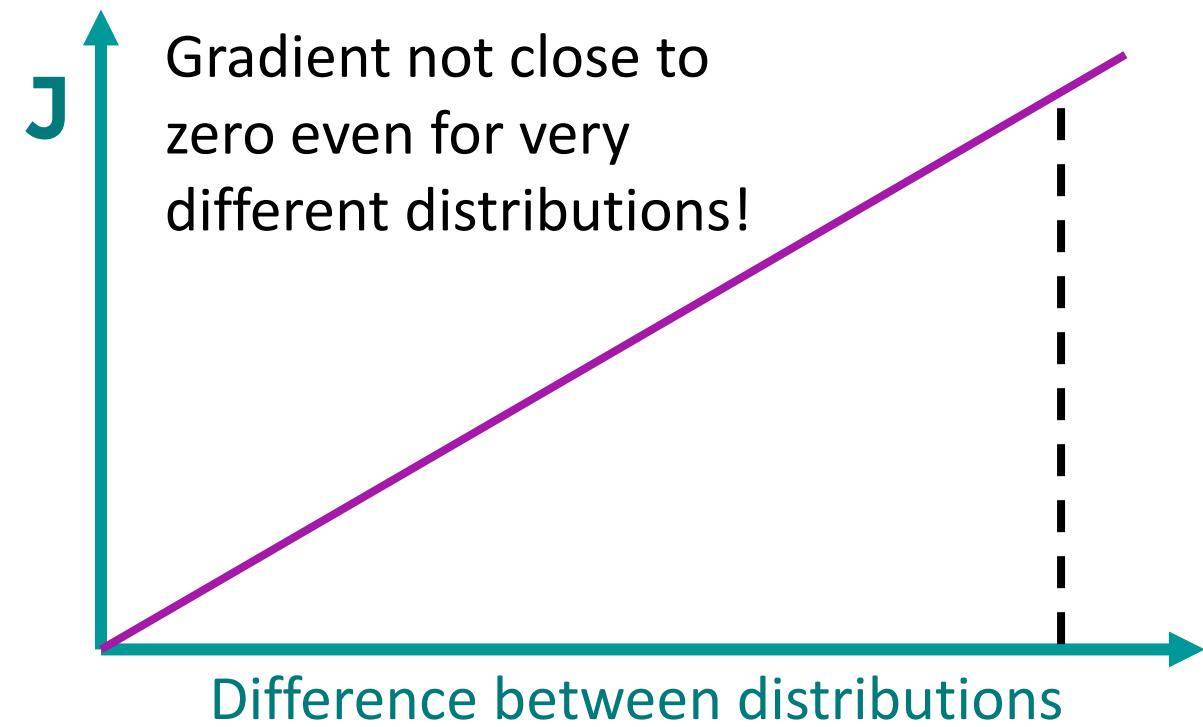
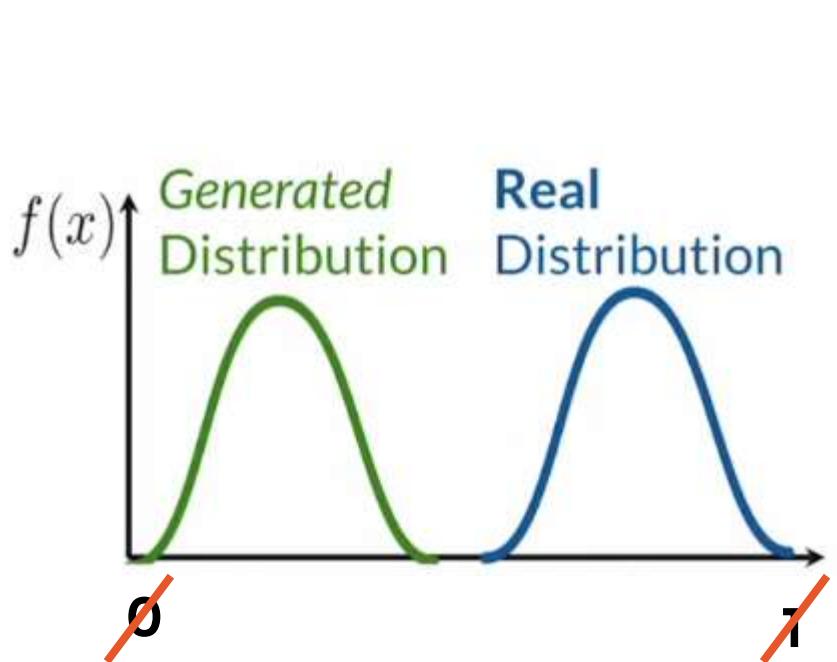


Cross Entropy



CLASS.  
vision

# Earth Mover's Distance



# Summary

- Earth mover's distance (EMD) is a function of amount and distance
- Doesn't have flat regions when the distributions are very different
- Approximating EMD solves the problems associated with BCE

# Wasserstein Loss

# How To Pronounce 😊



**Yann LeCun**  
@ylecun

...

Hint: "Wasserstein", as in "Wasserstein distance" is pronounced "vassershtaeen" not "wassersteen". It rhymes with "Einstein" not "saltine".

<https://twitter.com/ylecun/status/991375003626233858?lang=en>

# Outline

- BCE Loss Simplified
- W-Loss and its comparison with BCE Loss

# BCE Loss Simplified

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \left[ y^{(i)} \log(h(x^{(i)}, \theta)) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta)) \right]$$

# BCE Loss Simplified

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h(x^{(i)}, \theta)) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$\max_g$



Minimize  
cost

شبکه  
علیرضا



Maximize  
cost

# BCE Loss Simplified

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h(x^{(i)}, \theta)) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

minmax  
 $d$        $g$



Minimize  
cost

شبکه  
علیرضا



Maximize  
cost

# BCE Loss Simplified

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h(x^{(i)}, \theta)) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$$\min_d \max_g -[\mathbb{E}(\quad) + \mathbb{E}(\quad)]$$



Minimize  
cost

شبکه  
علیرضا



Maximize  
cost

# BCE Loss Simplified

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h(x^{(i)}, \theta)) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$$\min_d \max_g -[\mathbb{E}(\log(d(x))) + \mathbb{E}(1 - \log(d(g(z))) )]$$



Minimize  
cost

شبکه  
علیرضا



Maximize  
cost

# W-Loss

- W-Loss approximates the **Earth Mover's Distance**

# W-Loss

- W-Loss approximates the **Earth Mover's Distance**

$$\min_g \max_c \left[ \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) \right]$$

# W-Loss

- W-Loss approximates the **Earth Mover's Distance**

$$\min_g \max_c \left[ \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) \right]$$



Maximize  
the  
distance

# W-Loss

- W-Loss approximates the **Earth Mover's Distance**

$$\min_g \max_c \left[ \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) \right]$$



Generator

Minimize  
the  
distance

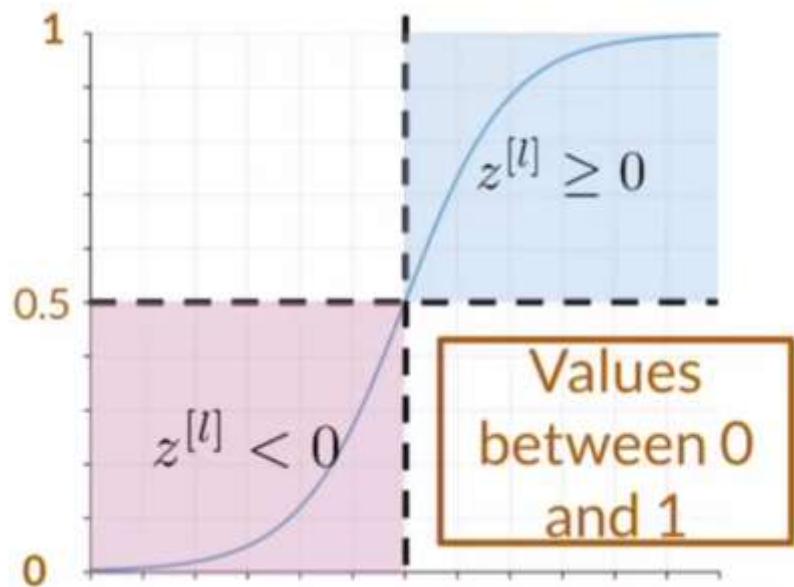


Critic

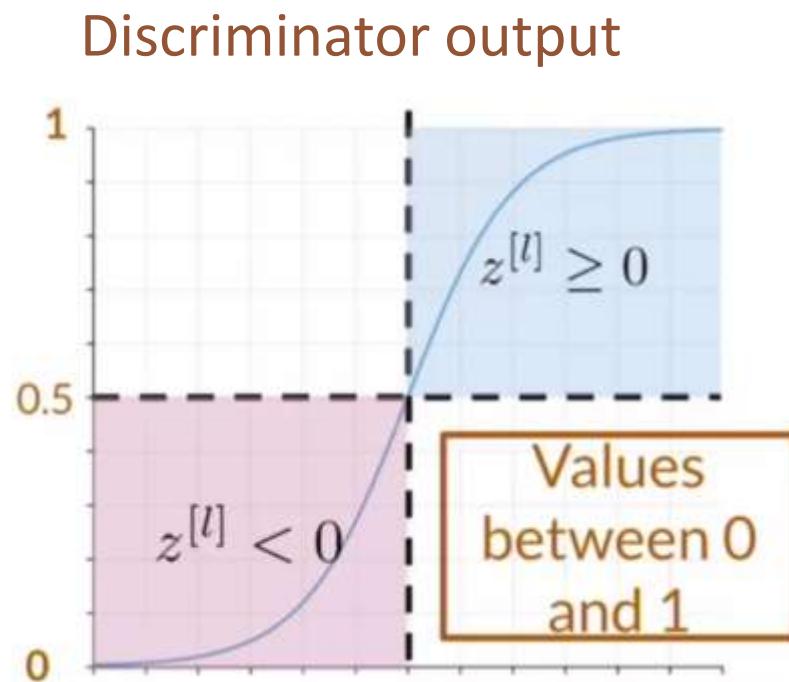
Maximize  
the  
distance

# Discriminator Output

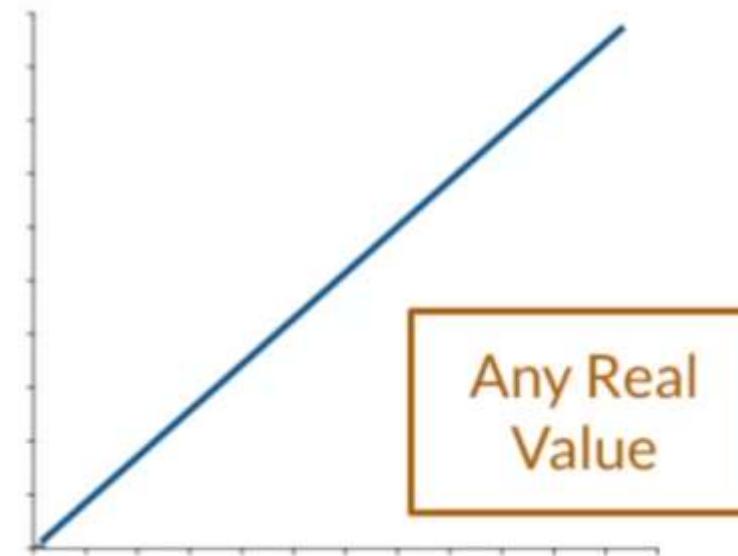
Discriminator output



# Discriminator Output



Critic  
Discriminator output



# W-Loss vs BCE Loss

BCE Loss	W-Loss
Discriminator outputs between 0 and 1	Critic outputs any number

# W-Loss vs BCE Loss

BCE Loss	W-Loss
Discriminator outputs between 0 and 1	Critic outputs any number
$-\left[ \mathbb{E}(\log(d(x))) + \mathbb{E}(1 - \log(d(g(z)))) \right]$	$\left[ \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) \right]$

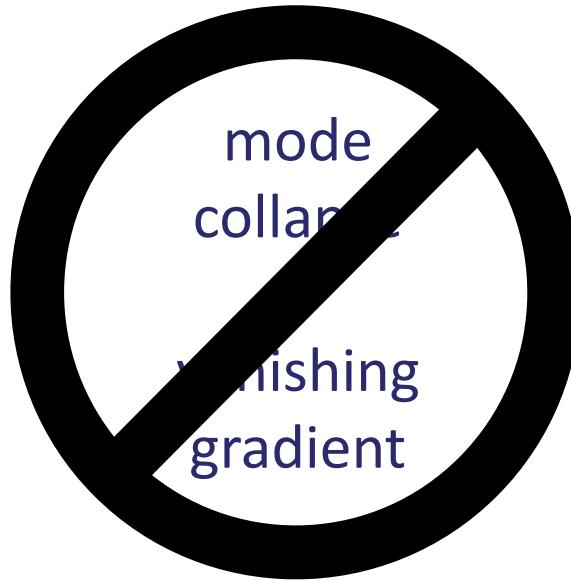
# W-Loss vs BCE Loss

BCE Loss	W-Loss
Discriminator outputs between 0 and 1	Critic outputs any number
$-\left[ \mathbb{E}(\log(d(x))) + \mathbb{E}(1 - \log(d(g(z)))) \right]$	$\left[ \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) \right]$

W-Loss helps with mode collapse and vanishing gradient problems

# Summary

- W-Loss looks very similar to BCE Loss
- W-Loss prevents mode collapse and vanishing gradient problems



# Condition on Wasserstein Critic

# Outline

- Continuity condition on the critic's neural network
- Why this condition matters

# Condition on W-loss

$$\min_g \max_c \left[ \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) \right]$$

# Condition on W-loss

$$\min_g \max_c \left[ \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) \right]$$

# Condition on W-loss

$$\min_g \max_c \left[ \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) \right]$$

# Condition on W-loss

$$\min_g \max_c \left[ \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) \right]$$

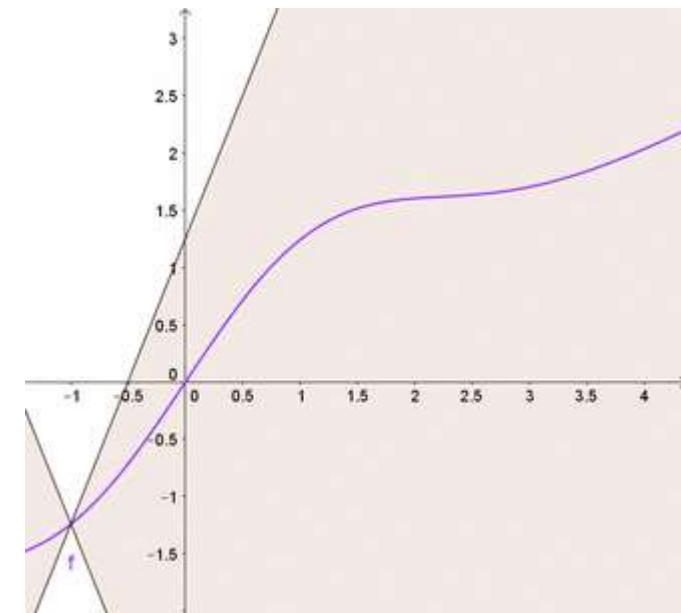
Needs to be 1-Lipschitz Continues

# Condition on W-loss

- Critic needs to be 1-Lipschitz continuous
- The norm of the gradient should be at most 1 for every point

# What is Lipschitz continuity

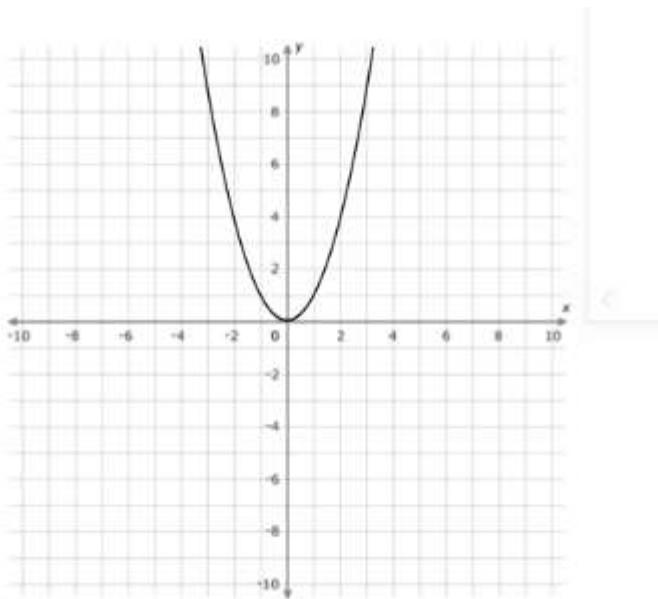
- Lipschitz continuous function is limited in how fast it can change
- For a Lipschitz continuous function, there exists a double cone (white) whose origin can be moved along the graph so that the whole graph always stays outside the double cone



[https://en.wikipedia.org/wiki/Lipschitz\\_continuity](https://en.wikipedia.org/wiki/Lipschitz_continuity)

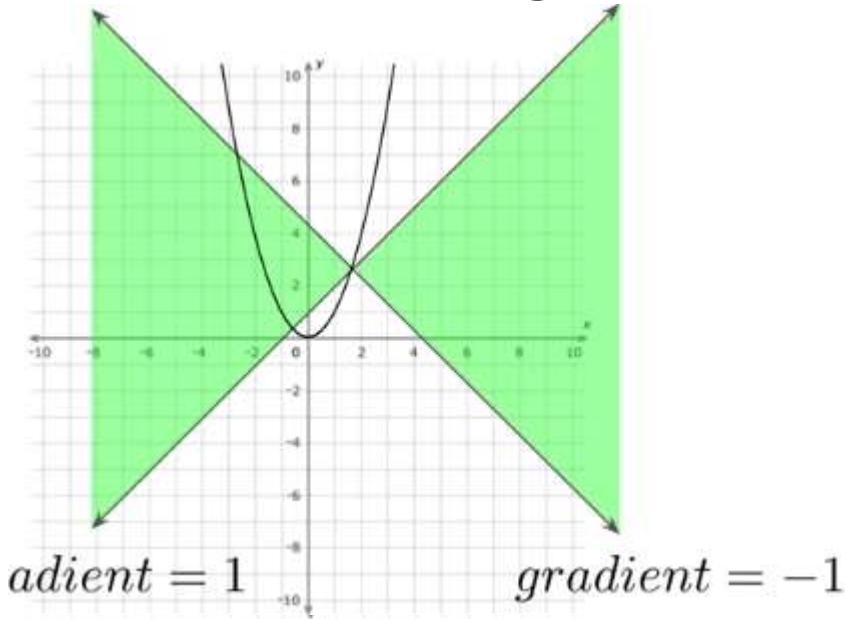
# Condition on W-loss

- Critic needs to be 1-Lipschitz continuous
- The norm of the gradient should be at most 1 for every point



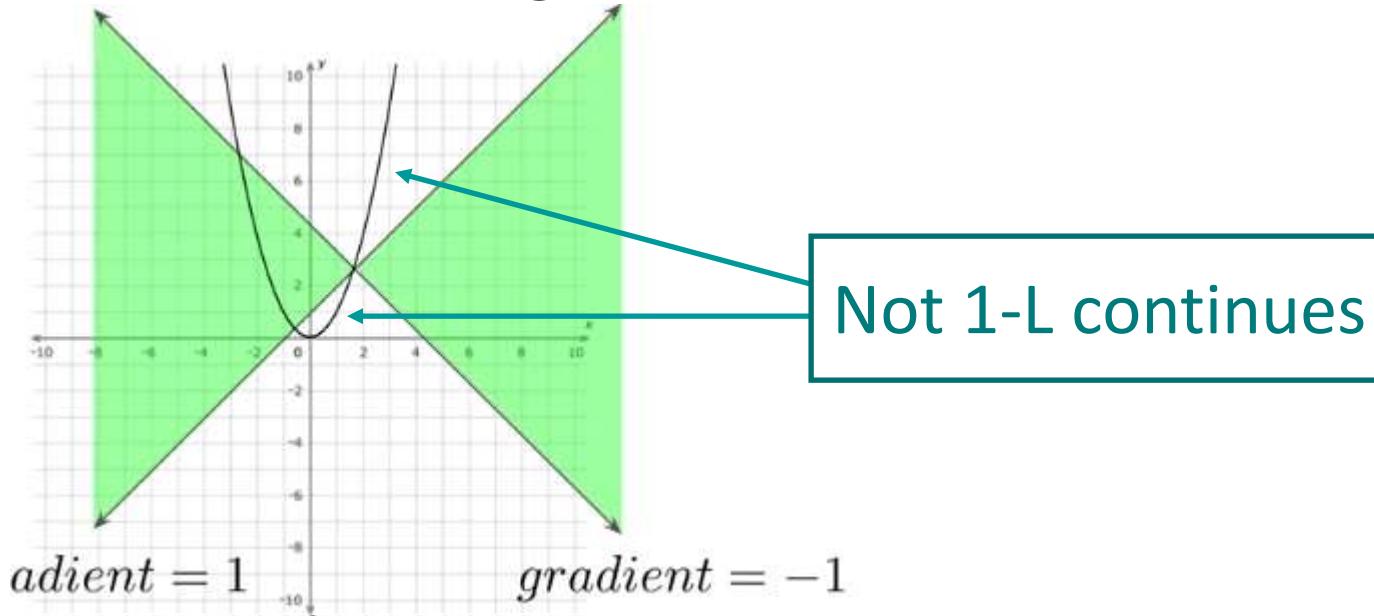
# Condition on W-loss

- Critic needs to be 1-Lipschitz continuous
- The norm of the gradient should be at most 1 for every point



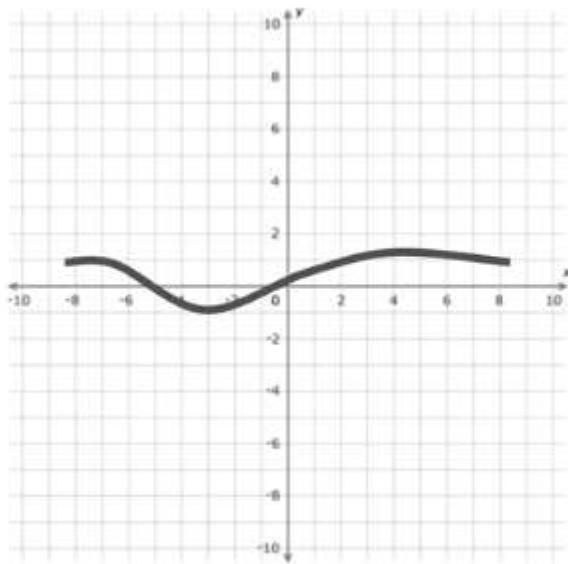
# Condition on W-loss

- Critic needs to be 1-Lipschitz continuous
- The norm of the gradient should be at most 1 for every point



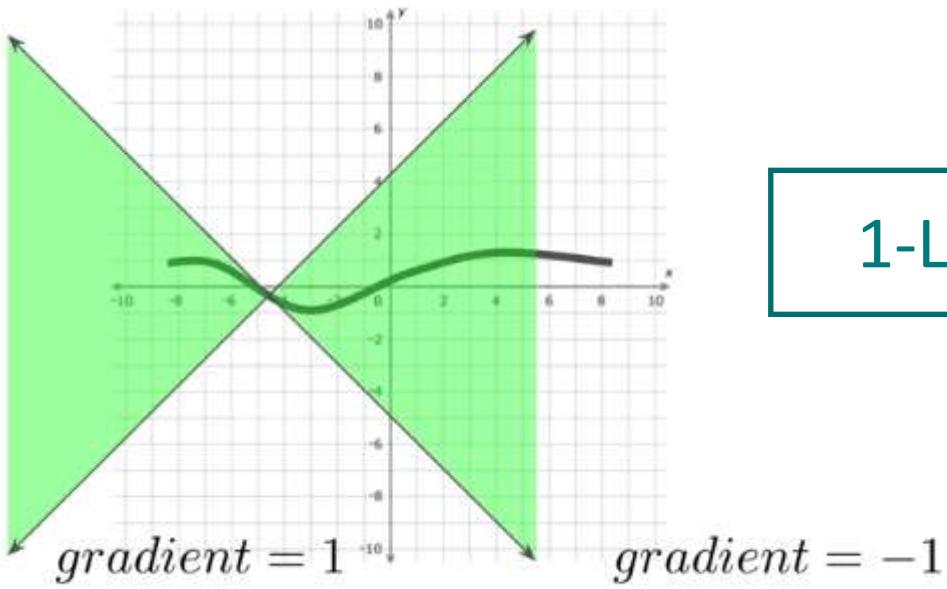
# Condition on W-loss

- Critic needs to be 1-Lipschitz continuous
- The norm of the gradient should be at most 1 for every point



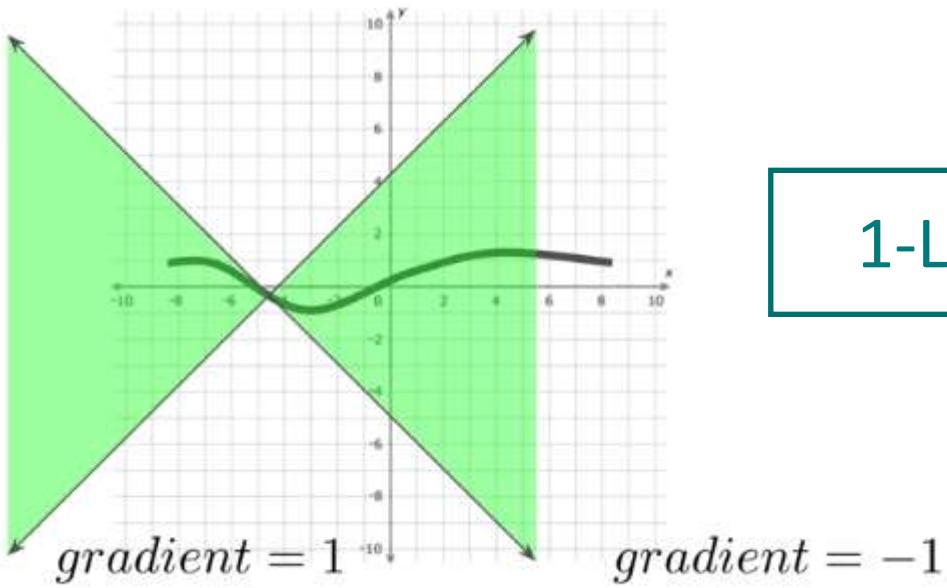
# Condition on W-loss

- Critic needs to be 1-Lipschitz continuous
- The norm of the gradient should be at most 1 for every point



# Condition on W-loss

- Critic needs to be 1-Lipschitz continuous
- The norm of the gradient should be at most 1 for every point



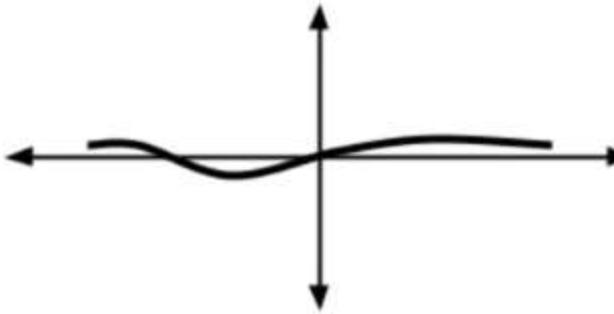
1-L continues

So W-loss is Valid.  
Needed for training stable  
neural networks with W-loss



# Summary

- Critic's neural network needs to be 1-L continuous when using w-loss
- This condition ensures that W-Loss is validly approximating Earth Mover's Distance



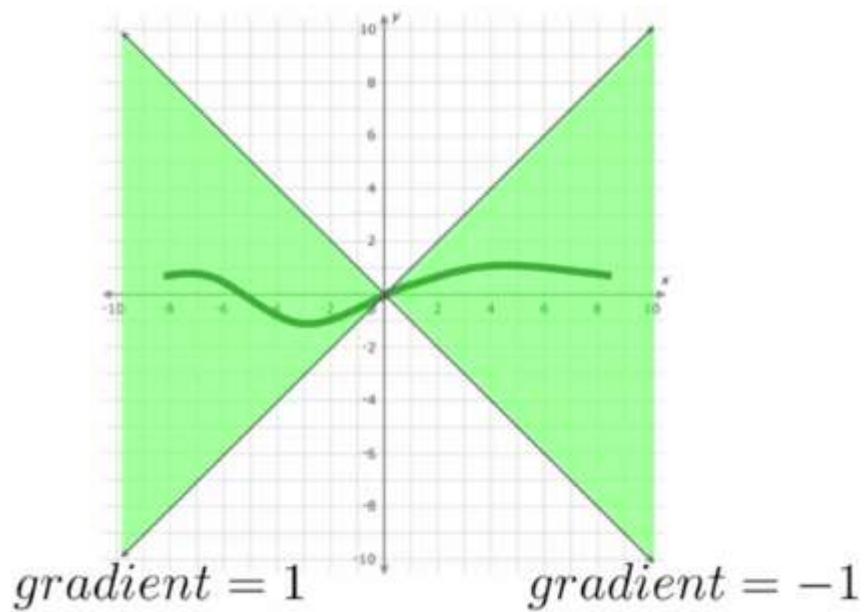
# 1-Lipschitz Continuity Enforcement

# Outline

- Weight clipping and gradient penalty
- Advantages of gradient penalty

# Enforcement

- Critic needs to be 1-L continuous



Norm of the gradient at most 1

$$\|\nabla f(x)\|_2 \leq 1$$



CLASS.  
vision  
www.class-vision.com

# Enforcement: Weight clipping

Weight clipping forces the weights of the critic to a fixed interval



# Enforcement: Weight clipping (Tensorflow implementation A)

```
class WeightsClip(tf.keras.constraints.Constraint):  
    def __init__(self, min_value=-0.01, max_value=0.01):  
        self.min_value = min_value  
        self.max_value = max_value  
    def __call__(self, w):  
        return tf.clip_by_value(w, self.min_value, self.max_value)
```

# Enforcement: Weight clipping (Tensorflow implementation A)

```
model = tf.keras.Sequential(name='critics')
model.add(Conv2D(16, 3, strides=2, padding='same',
                 kernel_constraint=WeightsClip(),
                 bias_constraint=WeightsClip()))
model.add(BatchNormalization(
                 beta_constraint=WeightsClip(),
                 gamma_constraint=WeightsClip()))
```

# Enforcement: Weight clipping (Tensorflow implementation B)

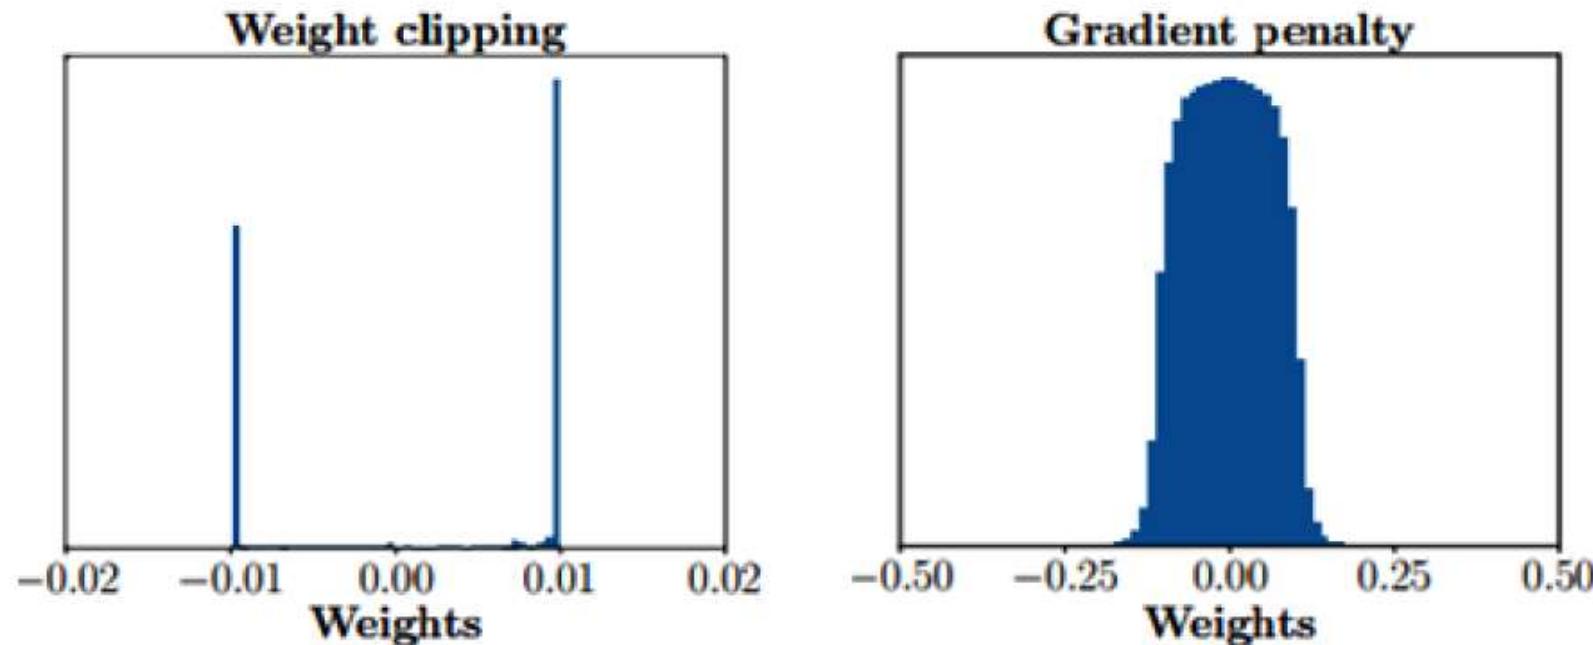
```
for layer in critic.layers:  
    weights = layer.get_weights()  
    weights = [tf.clip_by_value(w, -0.01, 0.01) for  
              w in weights]  
    layer.set_weights(weights)
```

# Enforcement: Weight clipping

Weight clipping forces the weights of the critic to a fixed interval



# Enforcement: Weight clipping VS Gradient penalty



Left: Weight clipping pushes weights toward two values. Right: Gradients produced by gradient penalty. Source: I. Gulrajani et al, 2017, Improved Training of Wasserstein GANs

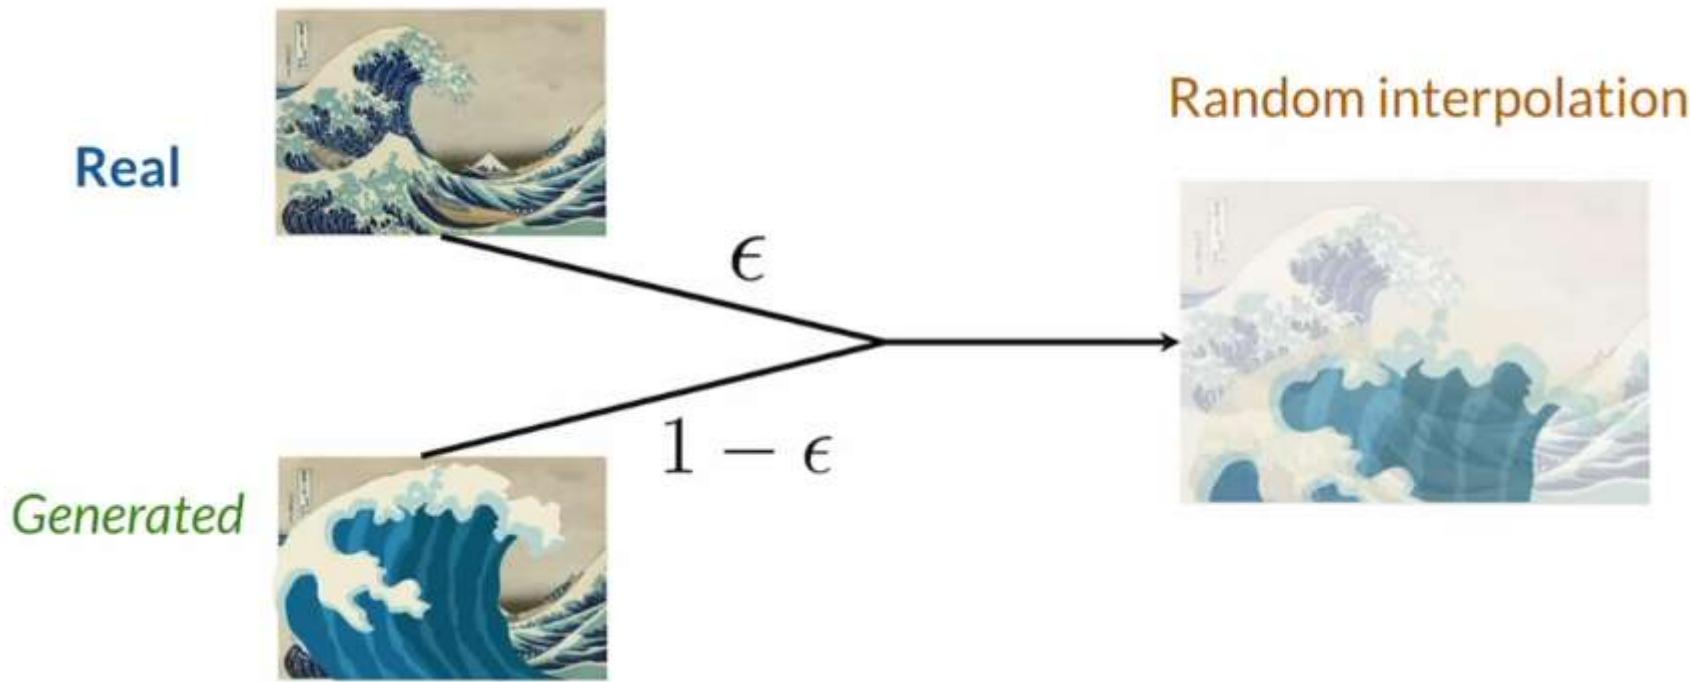
# Enforcement: Gradient penalty

$$\min_g \max_c \left[ \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) \right] + \lambda reg$$

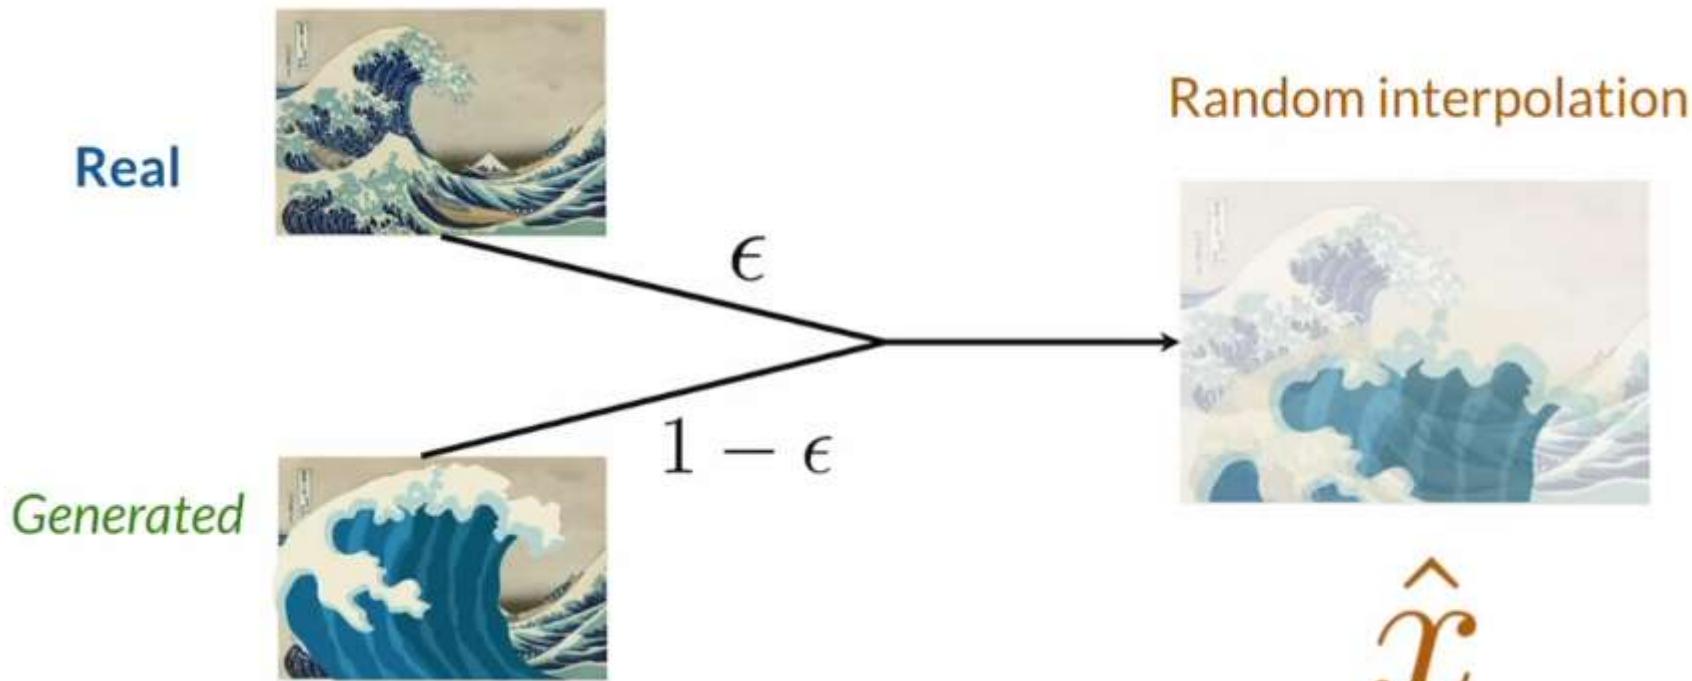


Regularization of critic's gradient

# Enforcement: Gradient penalty



# Enforcement: Gradient penalty



# Enforcement: Gradient penalty

$$(\|\nabla c(\hat{x})\|_2 - 1)^2$$

Regularization term

# Enforcement: Gradient penalty

$$(\|\nabla c(\hat{x})\|_2 - 1)^2 \quad \text{Regularization term}$$

$\downarrow$

$$\epsilon x + (1 - \epsilon) g(z) \quad \text{Interpolation}$$

# Enforcement: Gradient penalty

$$(\|\nabla c(\hat{x})\|_2 - 1)^2 \quad \text{Regularization term}$$

$\downarrow$

$$\epsilon \boxed{x} + (1 - \epsilon) \boxed{g(z)} \quad \text{Interpolation}$$

Real                      Generated

## 4 Gradient penalty

We now propose an alternative way to enforce the Lipschitz constraint. A differentiable function is 1-Lipschitz if and only if it has gradients with norm at most 1 everywhere, so we consider directly constraining the gradient norm of the critic’s output with respect to its input. To circumvent tractability issues, we enforce a soft version of the constraint with a penalty on the gradient norm for random samples  $\hat{x} \sim \mathbb{P}_{\hat{x}}$ . Our new objective is

$$L = \underbrace{\mathbb{E}_{\tilde{x} \sim \mathbb{P}_g} [D(\tilde{x})] - \mathbb{E}_{x \sim \mathbb{P}_r} [D(x)]}_{\text{Original critic loss}} + \underbrace{\lambda \mathbb{E}_{\hat{x} \sim \mathbb{P}_{\hat{x}}} [(\|\nabla_{\hat{x}} D(\hat{x})\|_2 - 1)^2]}_{\text{Our gradient penalty}}. \quad (3)$$

**Sampling distribution** We implicitly define  $\mathbb{P}_{\hat{x}}$  sampling uniformly along straight lines between pairs of points sampled from the data distribution  $\mathbb{P}_r$  and the generator distribution  $\mathbb{P}_g$ . This is motivated by the fact that the optimal critic contains straight lines with gradient norm 1 connecting coupled points from  $\mathbb{P}_r$  and  $\mathbb{P}_g$  (see [Proposition 1](#)). Given that enforcing the unit gradient norm constraint everywhere is intractable, enforcing it only along these straight lines seems sufficient and experimentally results in good performance.

**Penalty coefficient** All experiments in this paper use  $\lambda = 10$ , which we found to work well across a variety of architectures and datasets ranging from toy tasks to large ImageNet CNNs.

# Putting It All Together

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) + \lambda \mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$$

# Putting It All Together

$$\min_g \max_c \boxed{\mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))} + \lambda \mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$$

Makes the GAN less prone to **mode collapse** and **vanishing gradient**

# Putting It All Together

$$\min_g \max_c \quad \boxed{\mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))} + \boxed{\lambda \mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2}$$

Makes the GAN less prone to **mode collapse** and **vanishing gradient**

Tries to make the critic be 1-L Continuous, for the loss function to be **continuous and differentiable**

# Summary

- Weight clipping and gradient penalty are ways to enforce 1-L continuity
- Gradient penalty tends to work better

# Implementation

- [https://keras.io/examples/generative/wgan\\_gp/](https://keras.io/examples/generative/wgan_gp/)

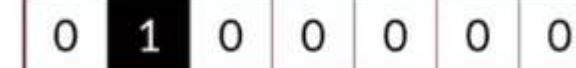
# Conditional Generation

# Conditional Generation: Inputs

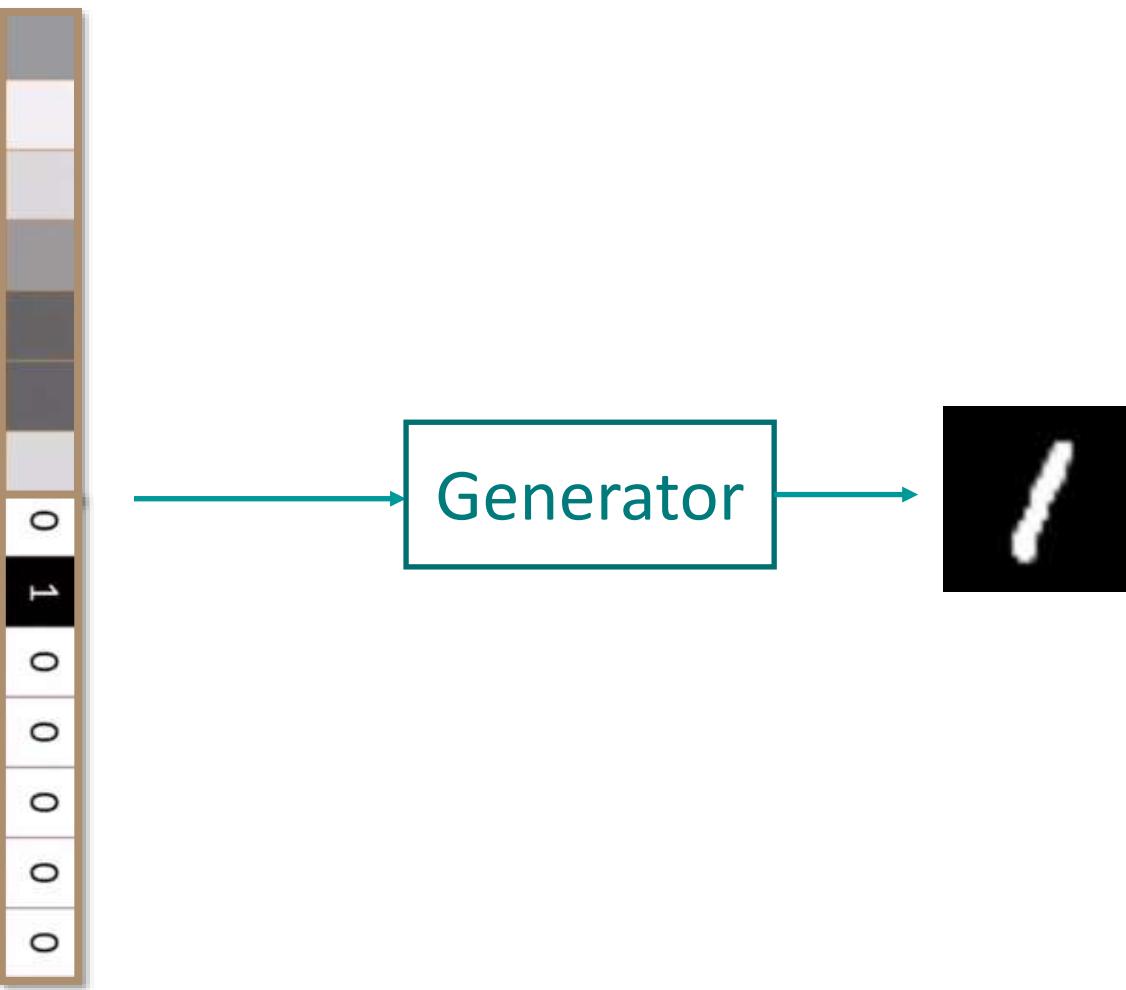
Noise vector



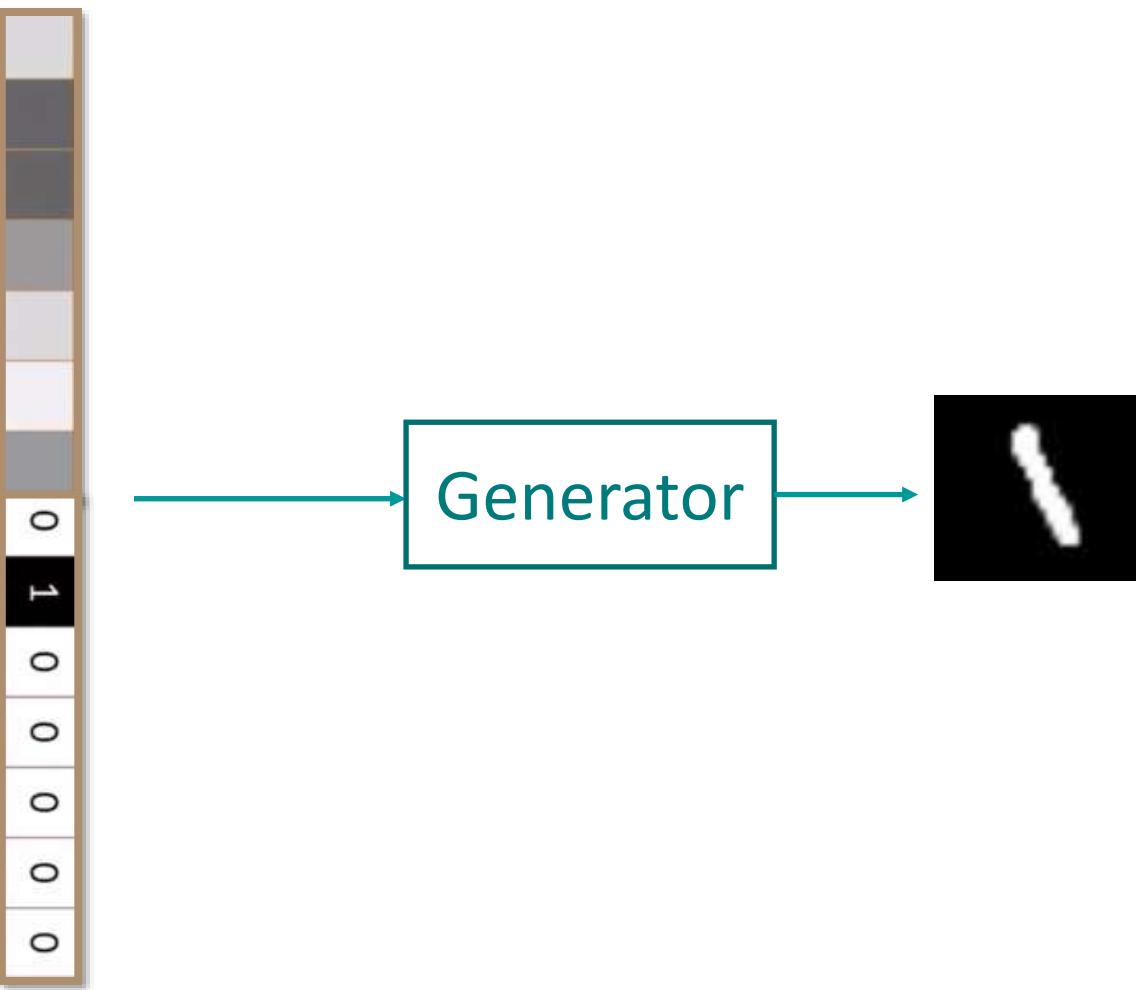
Class (one-hot) vector



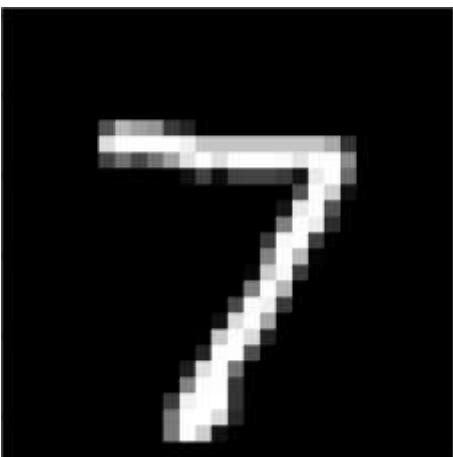
# Conditional Generation: Inputs



# Conditional Generation: Inputs



# Discriminator Input

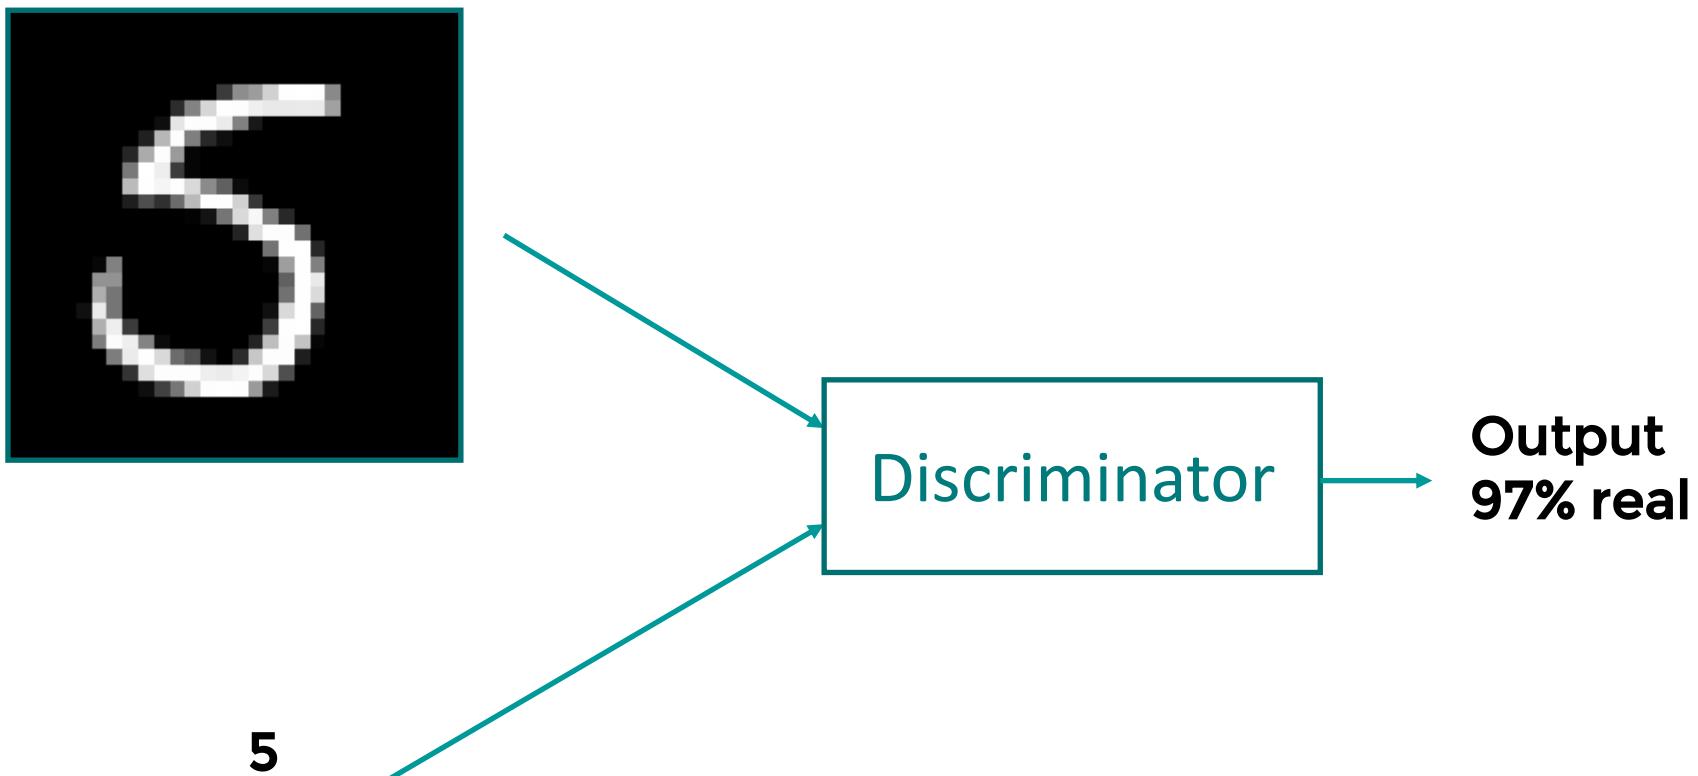


5

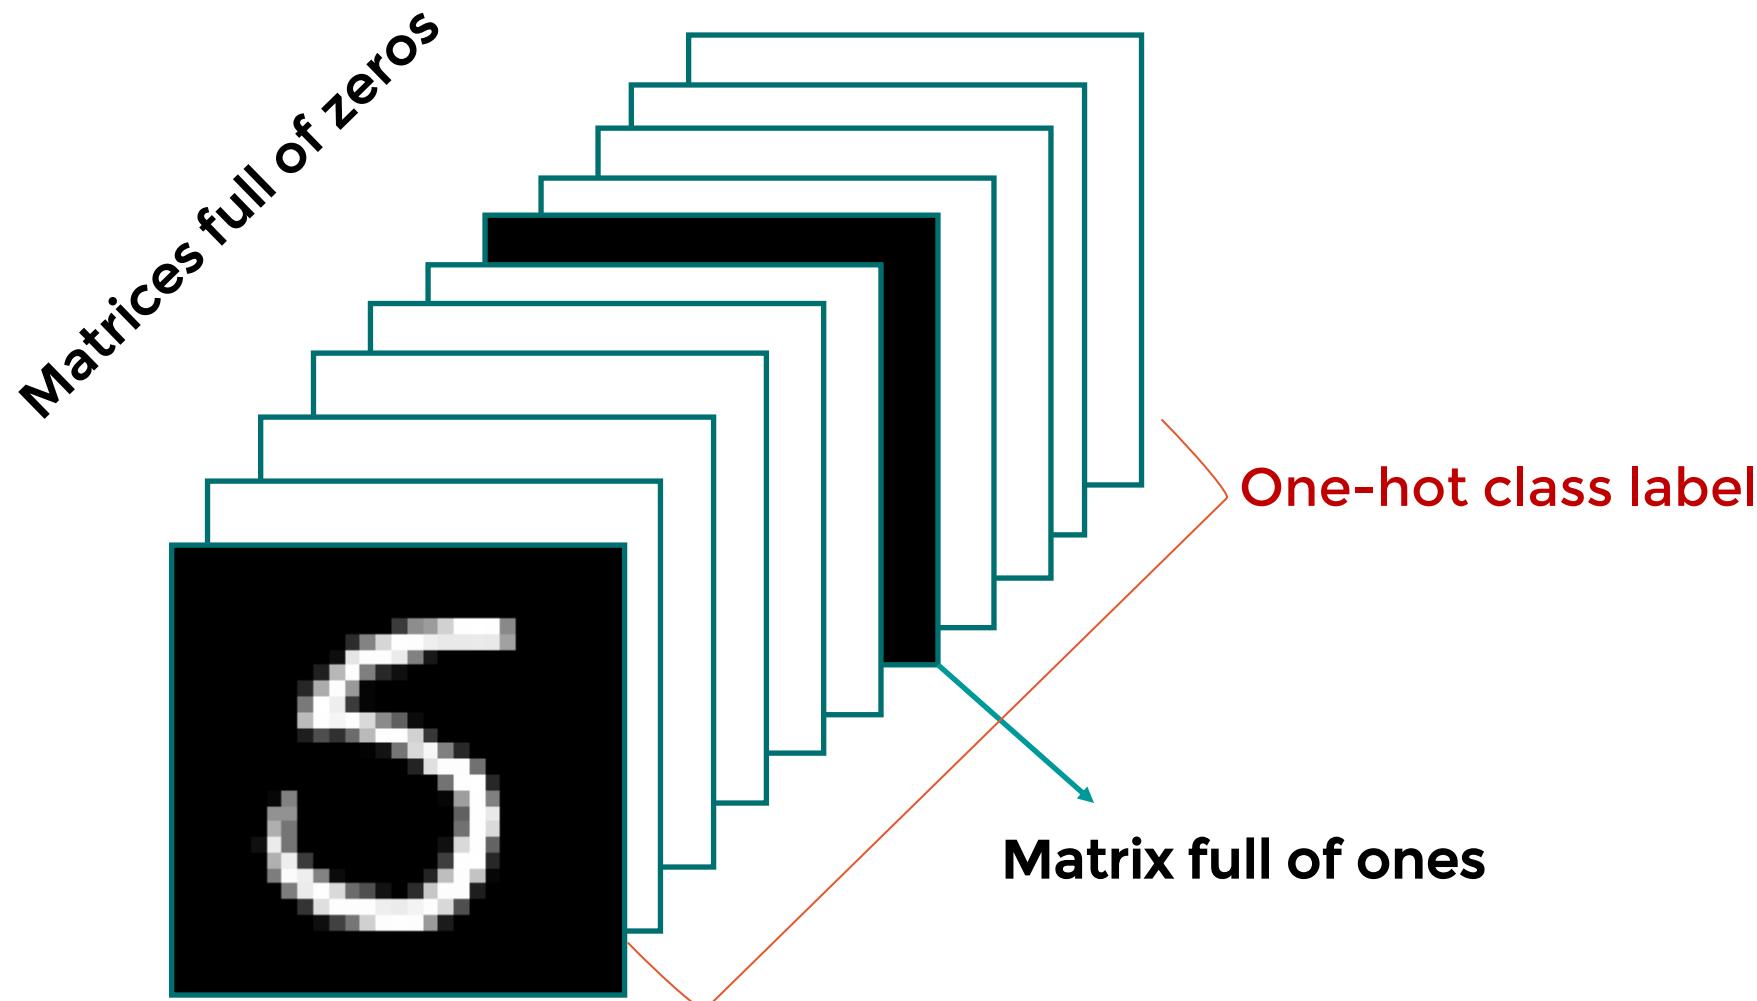


Output  
3% real

# Discriminator Input



# Discriminator Input



# Summary

- The class is passed to the generator as one-hot vectors
- The class is passed to the discriminator as one-hot matrices
- The size of the vector and the number of matrices represent the number of classes

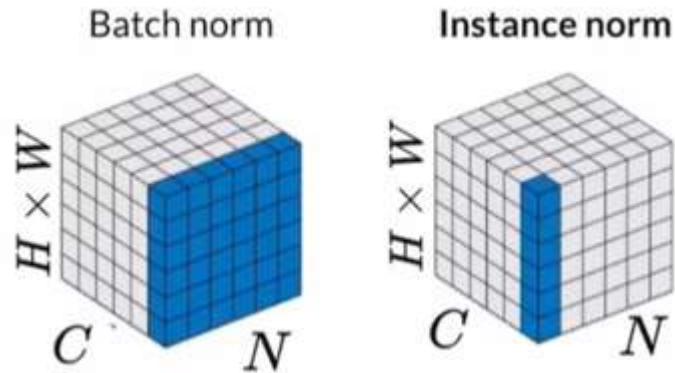
# Implementation...

- [https://keras.io/examples/generative/conditional\\_gan/](https://keras.io/examples/generative/conditional_gan/)

# AdaIN

# Instance Norm

$$\frac{x_i - \mu(x_i)}{\delta(x_i)}$$



# AdaIN

$$AdaIN(x_i, y) = y_{s,i} \frac{x_i - \mu(x_i)}{\delta(x_i)} + y_{b,i}$$

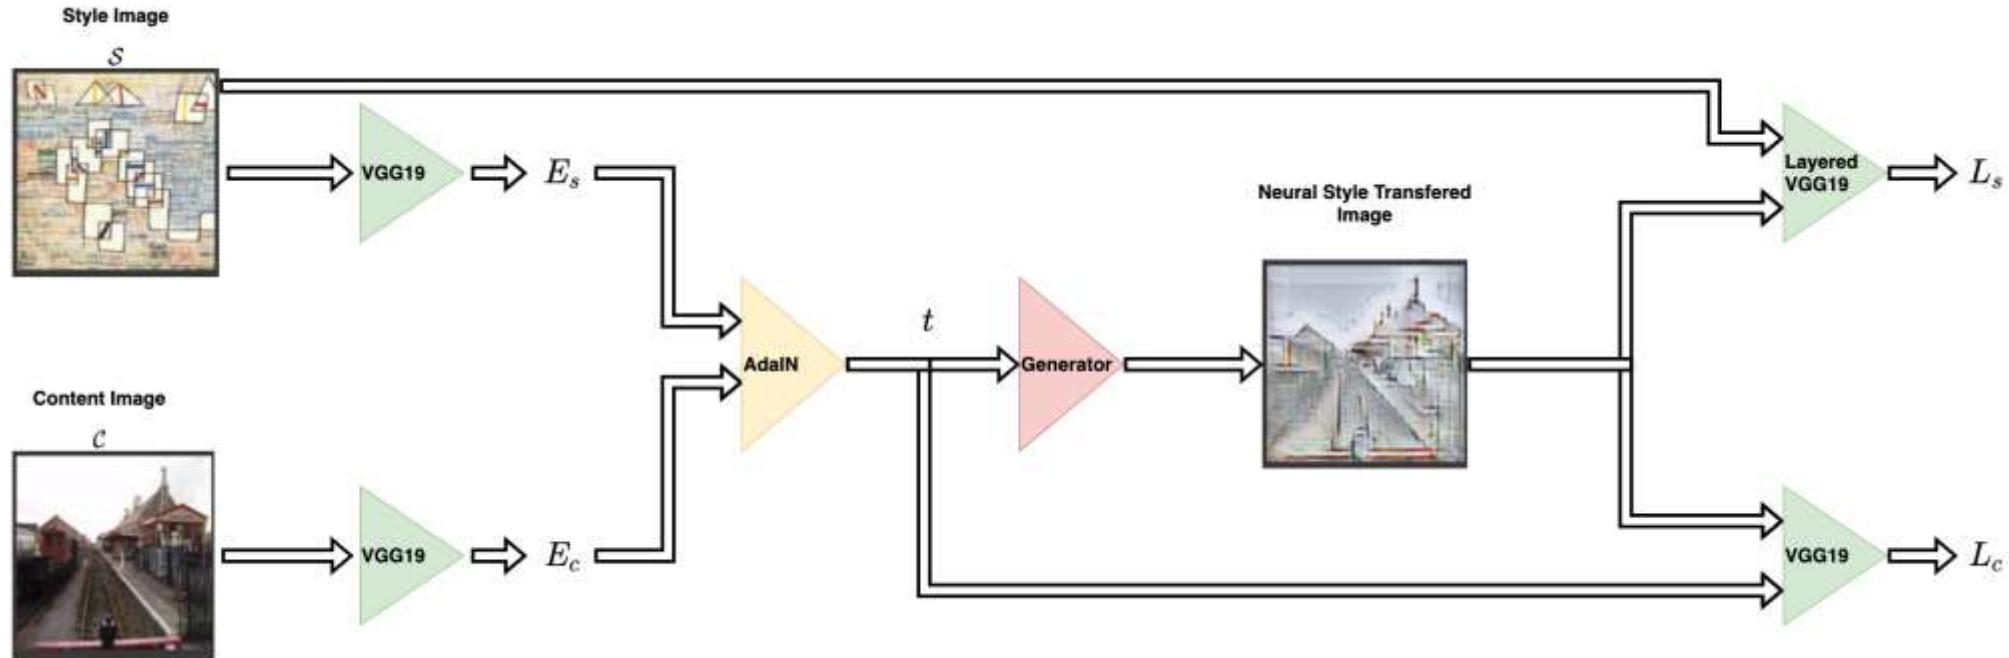
Step 1: Instance  
Normalization

# AdaIN

$$AdaIN(x_i, y) = y_{s,i} \frac{x_i - \mu(x_i)}{\delta(x_i)} + y_{b,i}$$

Step 2: Adaptive styles

# Architecture



# Loss

$$\mathcal{L}_t = \mathcal{L}_c + \lambda \mathcal{L}_s$$

# Loss

$$\mathcal{L}_t = \mathcal{L}_c + \lambda \mathcal{L}_s$$



$$\mathcal{L}_c = \|f(g(t)) - t\|_2$$

# Loss

$$\mathcal{L}_t = \mathcal{L}_c + \lambda \mathcal{L}_s$$



$$\mathcal{L}_s = \sum_{i=1}^L \|\mu(\phi_i(g(t))) - \mu(\phi_i(s))\|_2 + \sum_{i=1}^L \|\sigma(\phi_i(g(t))) - \sigma(\phi_i(s))\|_2$$

# Implementation

- <https://keras.io/examples/generative/adain/>

# GAN Improvements

# 4.5 years of GAN progress on face generation.



Ian Goodfellow  
@goodfellow\_ian

...

4.5 years of GAN progress on face generation.

[arxiv.org/abs/1406.2661](https://arxiv.org/abs/1406.2661) [arxiv.org/abs/1511.06434](https://arxiv.org/abs/1511.06434)

[arxiv.org/abs/1606.07536](https://arxiv.org/abs/1606.07536) [arxiv.org/abs/1710.10196](https://arxiv.org/abs/1710.10196)

[arxiv.org/abs/1812.04948](https://arxiv.org/abs/1812.04948)



[https://twitter.com/goodfellow\\_ian/status/1084973596236144640](https://twitter.com/goodfellow_ian/status/1084973596236144640)

GANs

Alireza Akhavanpour

شبکه های مولد

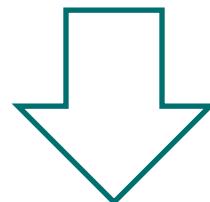
علیرضا اخوان پور



CLASS.  
VISION

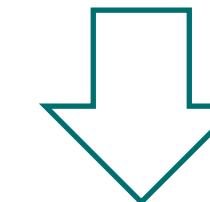
# Main improvements: (1) Stability

4 3 7



High std

4 4 4

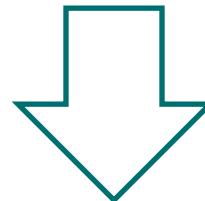


Low std

# Main improvements:

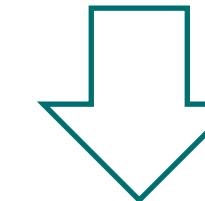
## (1) Stability

4 3 7



High std

4 4 4



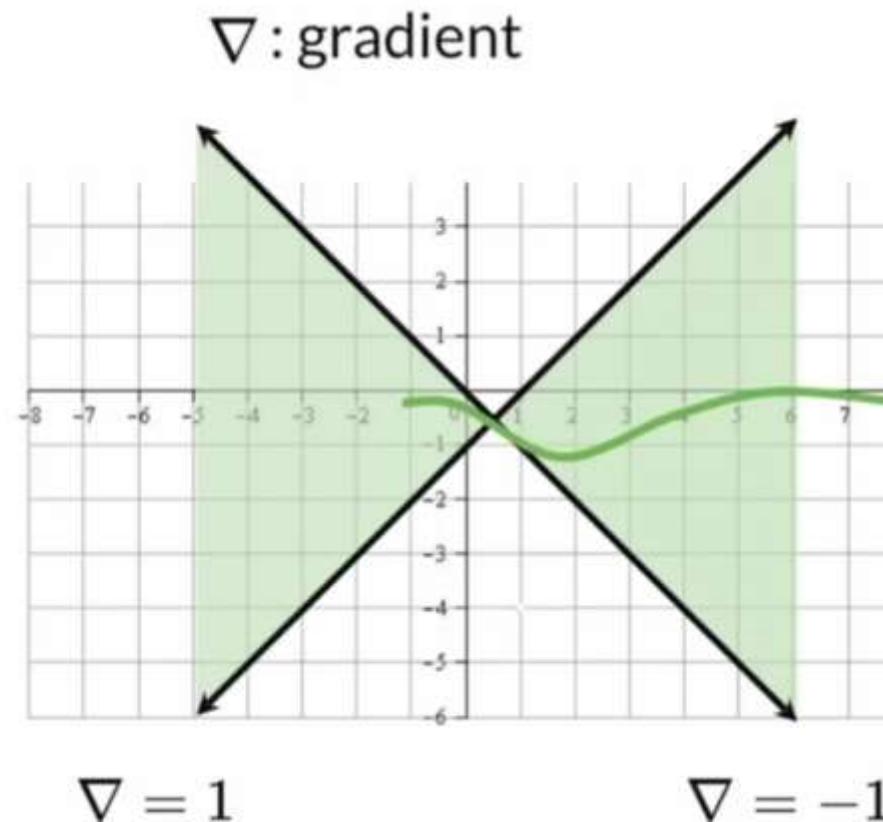
Low std

Use standard deviation in batch to encourage diversity.

# Main improvements: (1) Stability

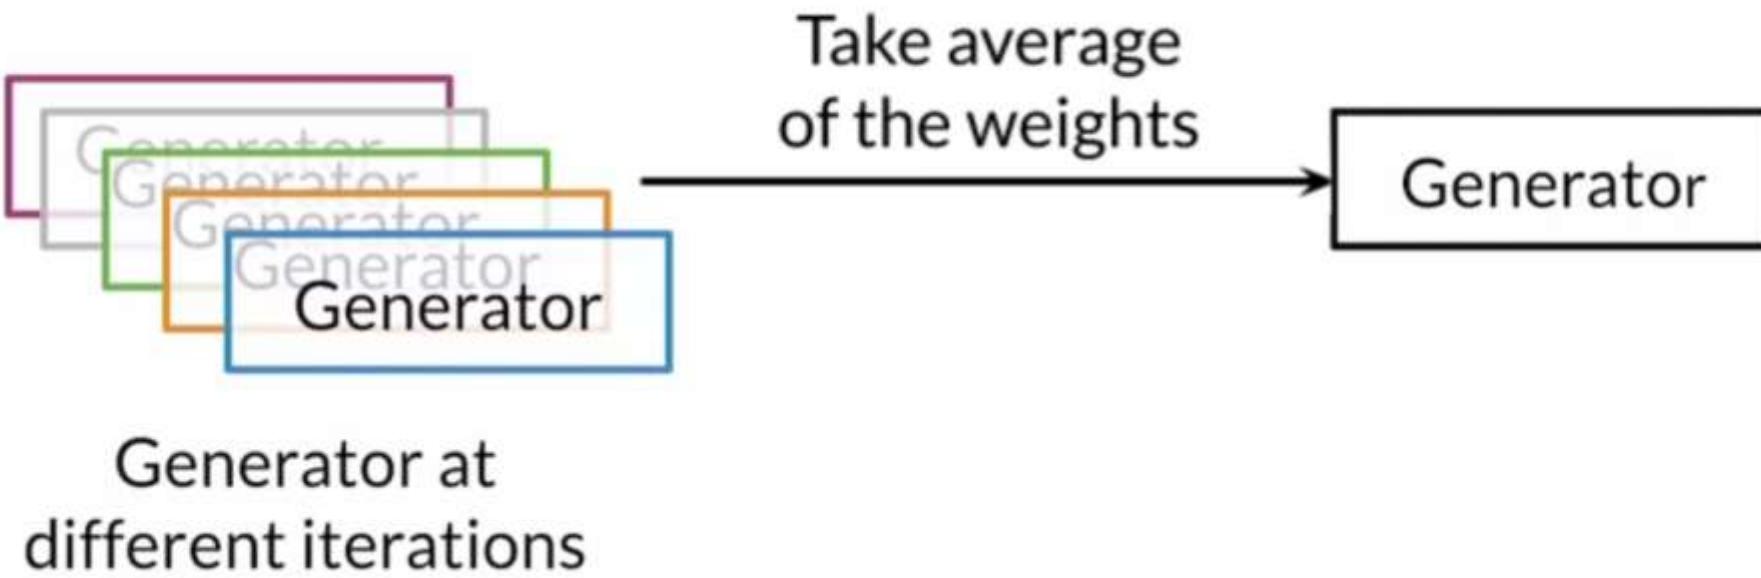
Improve stability by enforcing  
1-Lipschitz continuity

E.g. **WGAN-GP** and **Spectral  
Normalization**



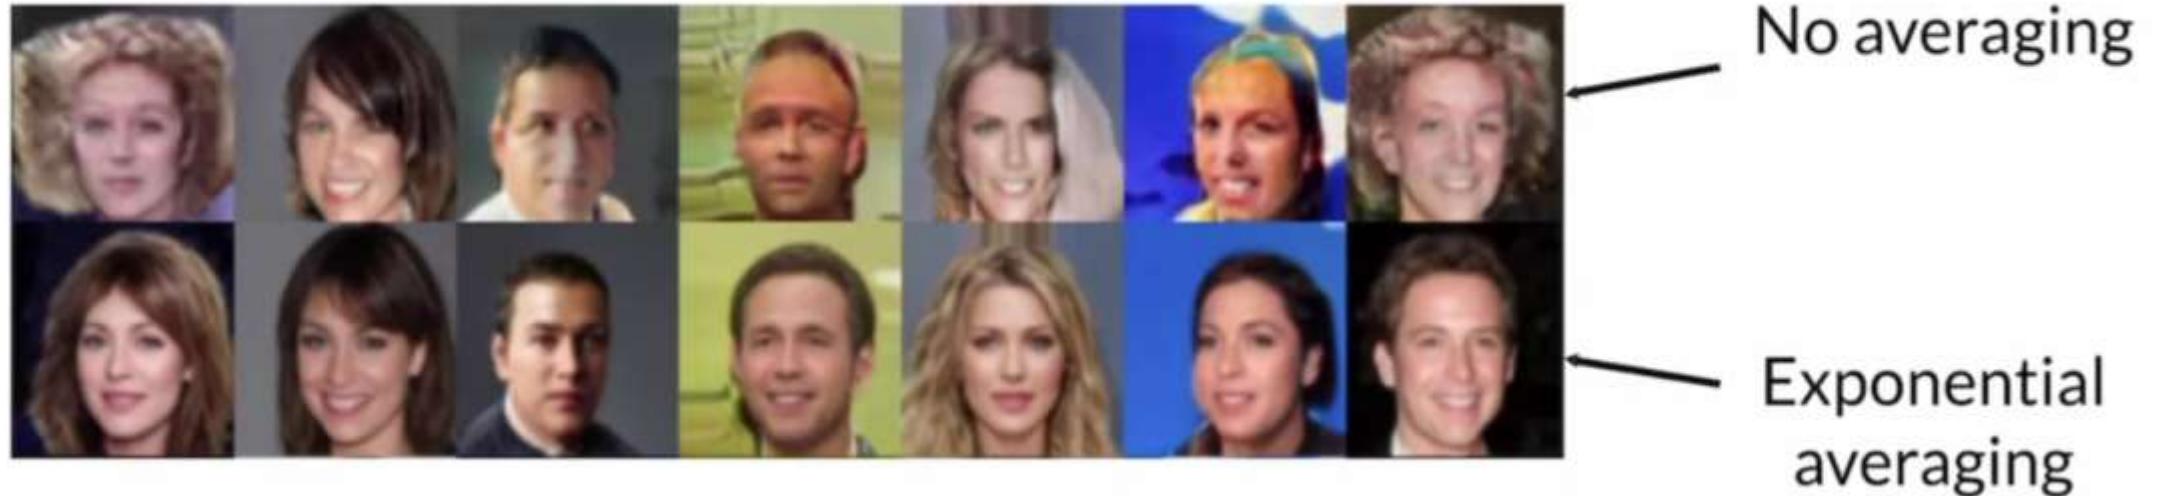
# Main improvements:

## (1) Stability



# Main improvements:

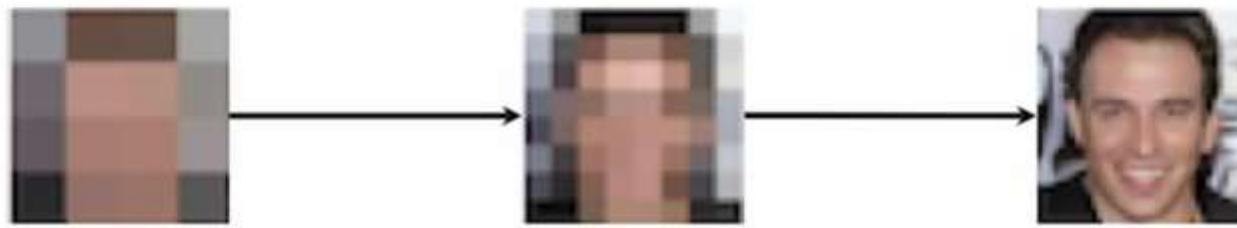
## (1) Stability



Use moving average for smoother results

The Unusual Effectiveness of Averaging in GAN Training:  
<https://arxiv.org/pdf/1806.04498v2.pdf>

# Main improvements: (1) Stability



Progressive growing gradually trains  
GAN at increasing resolutions

Progressive Growing of GANs for Improved Quality, Stability, and Variation:  
<https://arxiv.org/pdf/1710.10196.pdf>

# Main improvements: (2) Capacity



Larger models can use  
higher resolution images

# Main improvements: (2) Diversity



A Style-Based Generator Architecture for Generative Adversarial Networks

<https://arxiv.org/pdf/1812.04948.pdf>

GANs

Alireza Akhavanpour

شبکه های مولد

علیرضا اخوان پور



CLASS.  
VISION

# Summary

- GANs have improved because of:
  - **Stability** - longer training and better images
  - **Capacity** - larger models and higher resolution images
  - **Diversity** - increasing variety in generated images

# Progressive Growing

# ProGAN

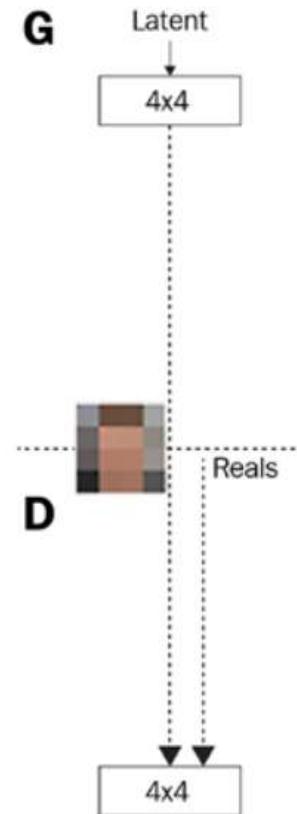


<https://arxiv.org/pdf/1710.10196.pdf>

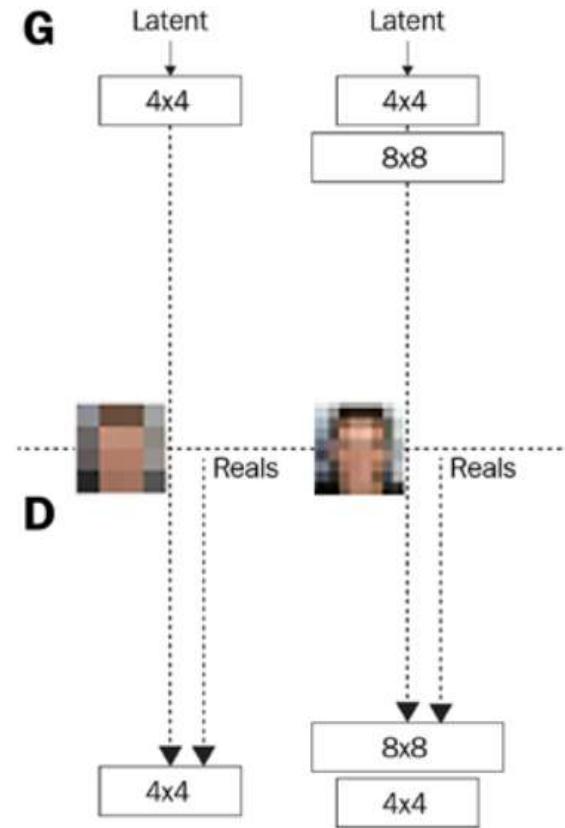
# ProGAN

- Progressive growing
- Pixel normalization
- Minibatch statistics
- Equalized learning rate

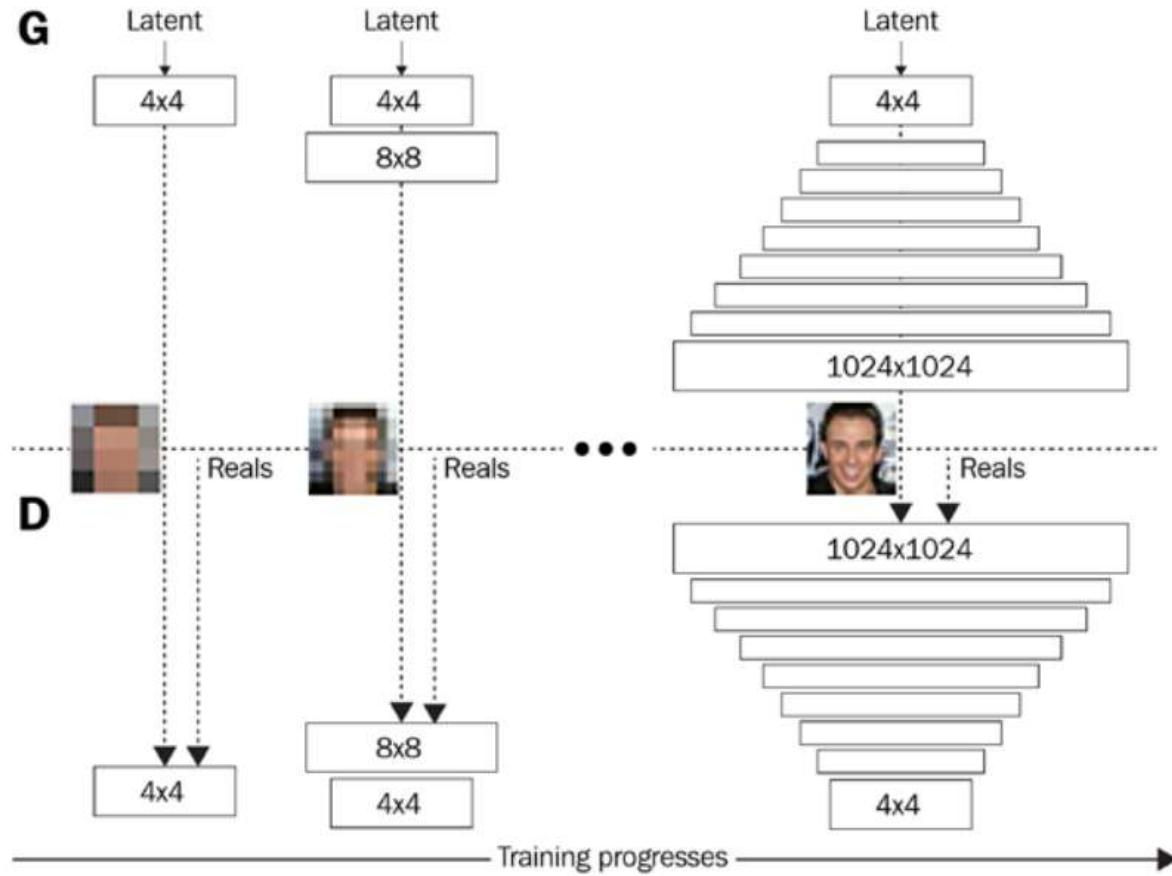
# Progressive growing



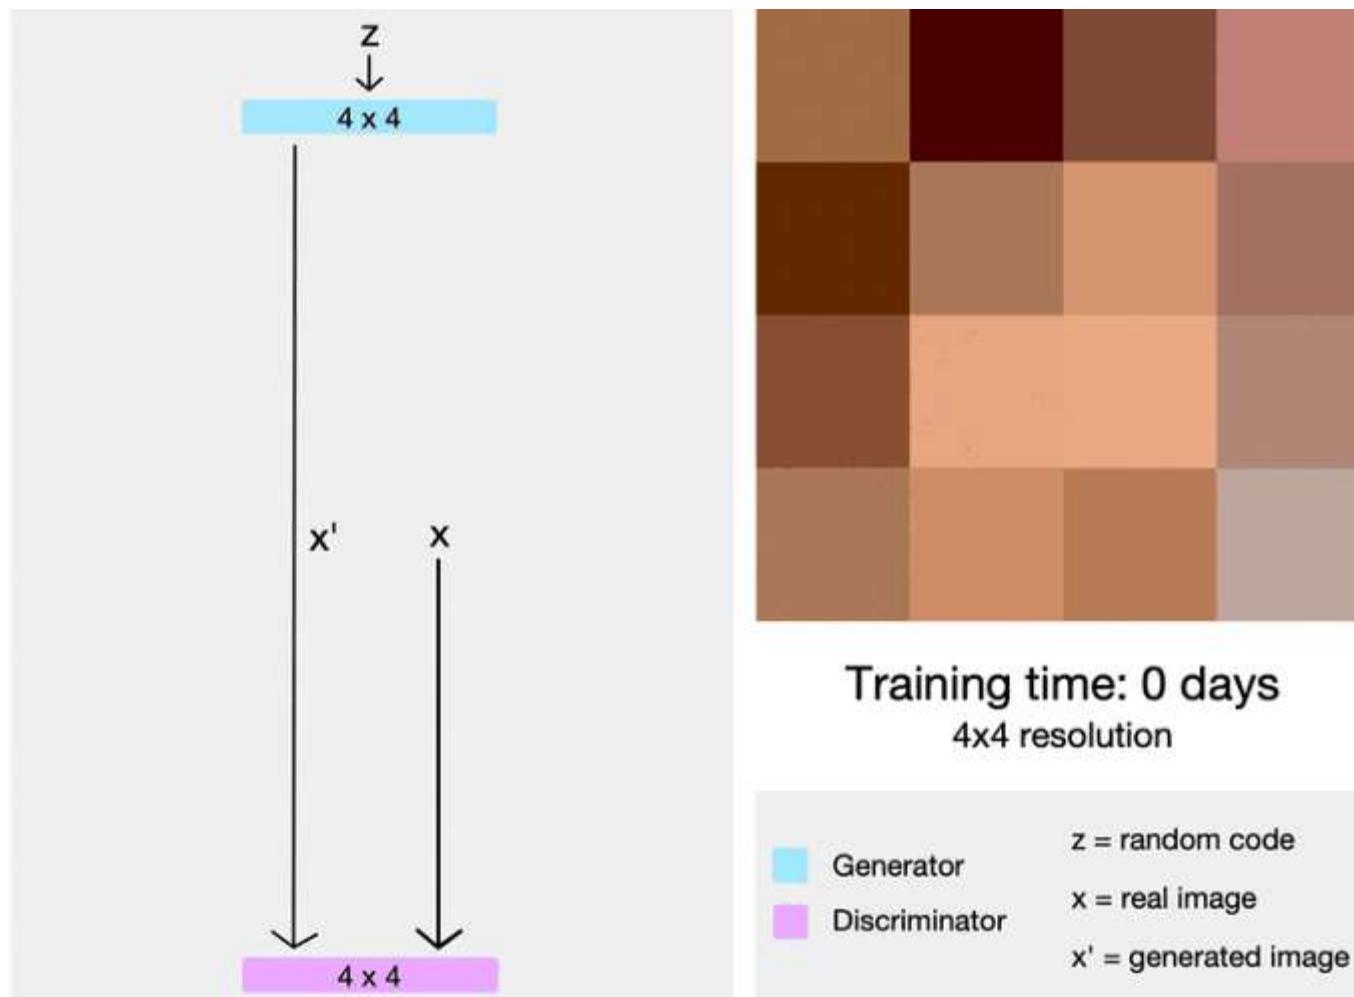
# Progressive growing



# Progressive growing



# Progressive Growing in Action



The **ProGAN** starts out generating very low resolution images. When training stabilizes, a new layer is added and the resolution is doubled. This continues until the output reaches the desired resolution. By progressively growing the networks in this fashion, high-level structure is learned first, and training is stabilized.

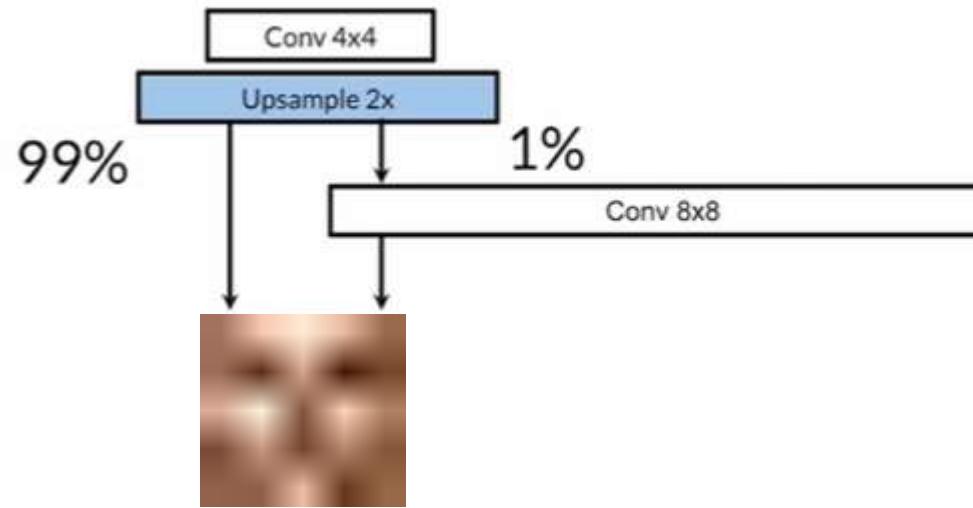


# Progressive Growing in Action



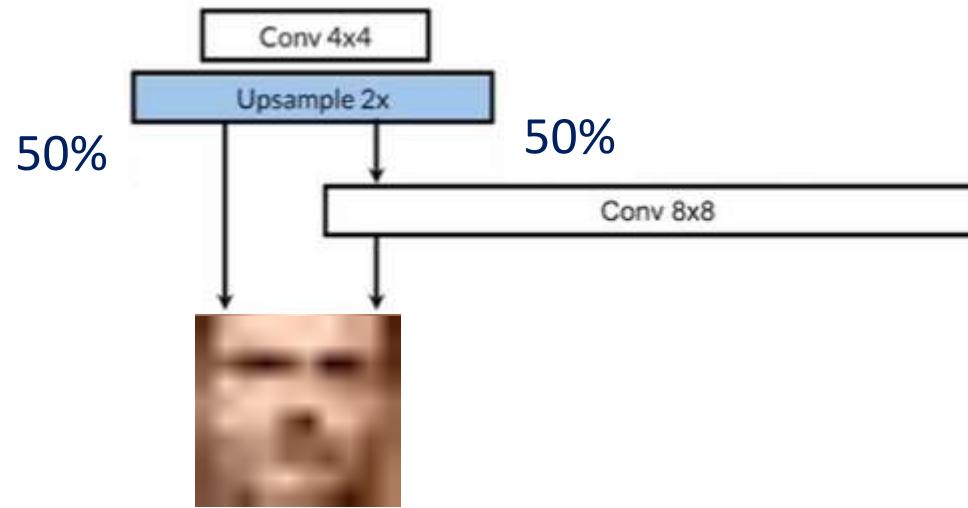
<https://www.gwern.net/images/gan/stylegan/2019-03-16-stylegan-facestraining.mp4>

# Progressive Growing: Generator



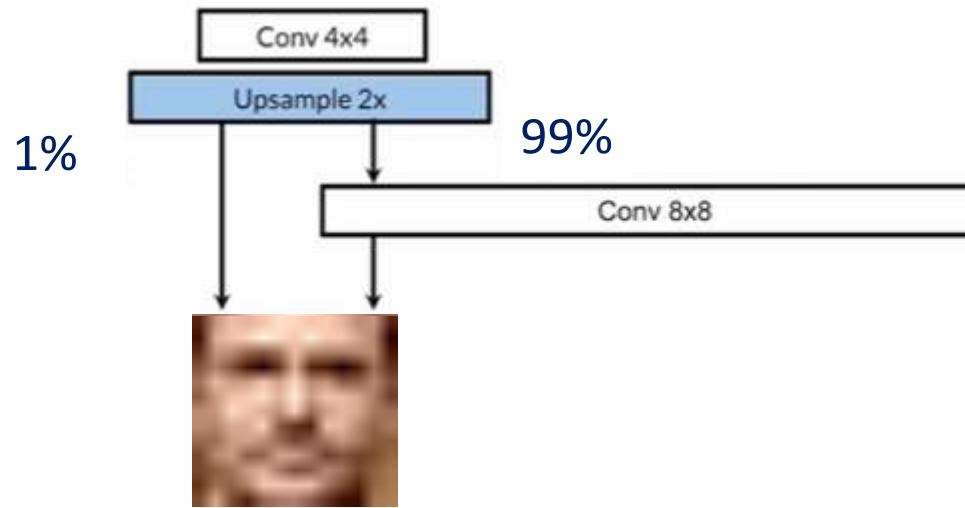
Source: <https://bit.ly/2TggMR8>

# Progressive Growing: Generator



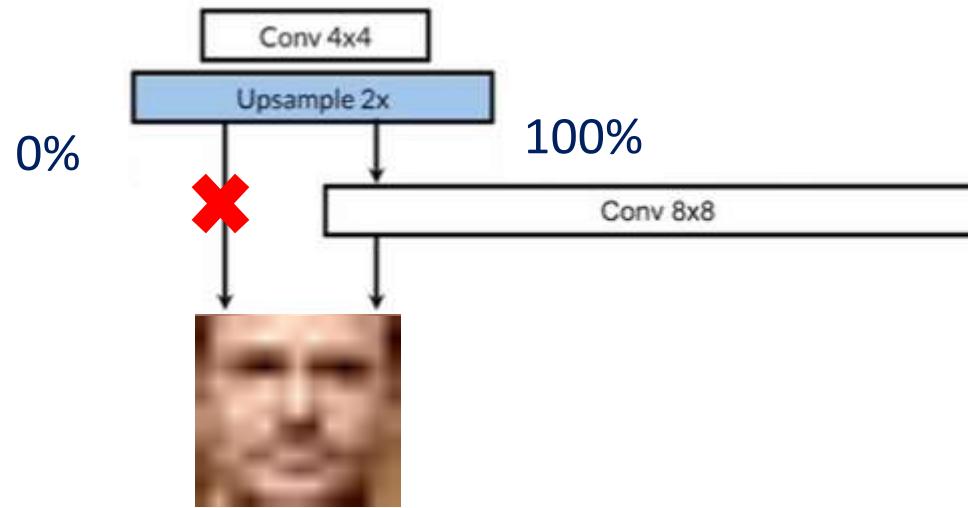
Source: <https://bit.ly/2TggMR8>

# Progressive Growing: Generator



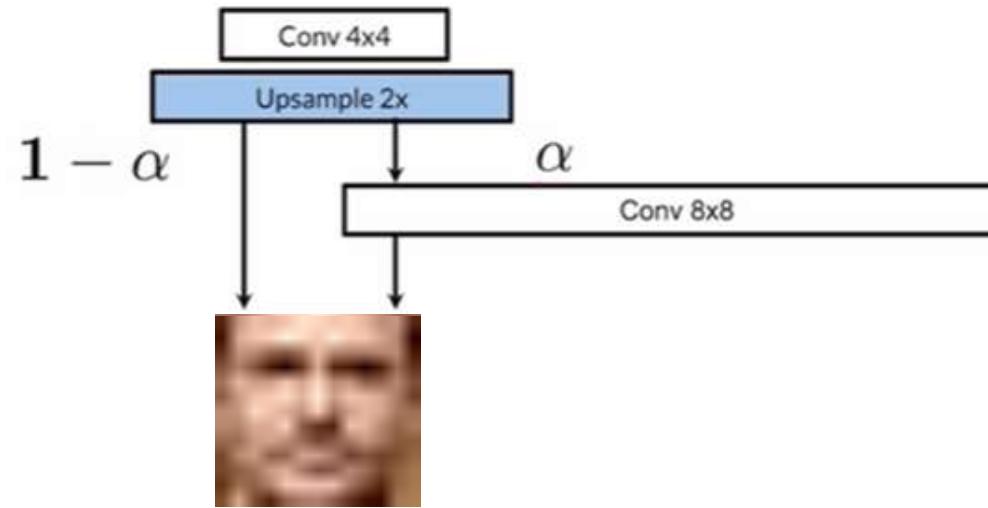
Source: <https://bit.ly/2TggMR8>

# Progressive Growing: Generator



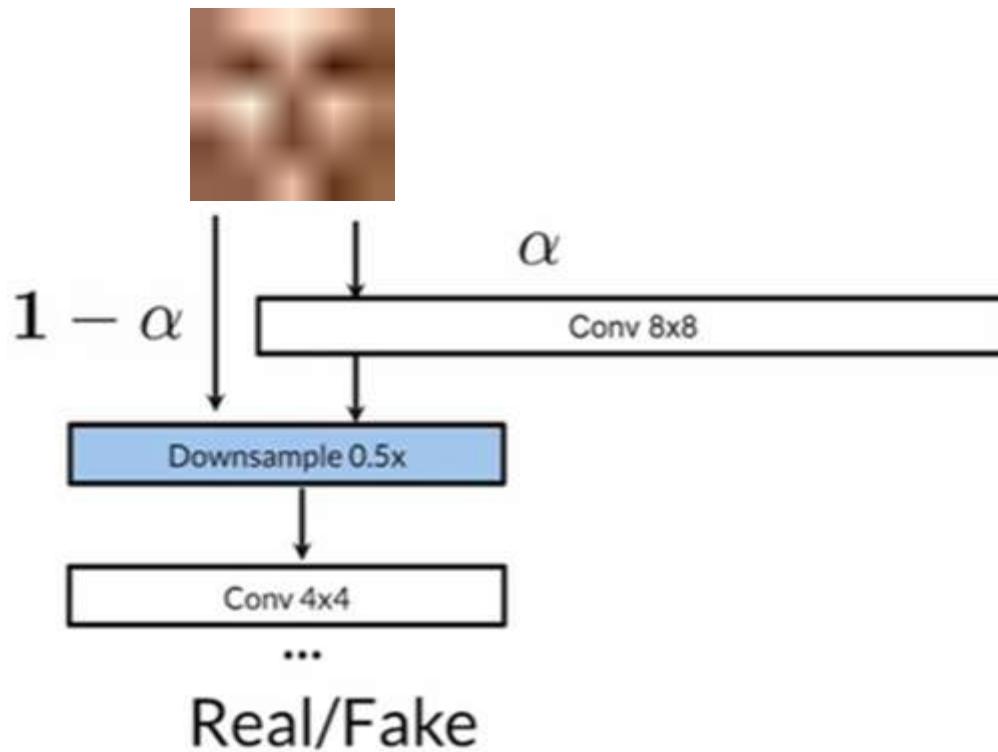
Source: <https://bit.ly/2TggMR8>

# Progressive Growing: Generator



Source: <https://bit.ly/2TggMR8>

# Progressive Growing: Discriminator



Source: <https://bit.ly/2TggMR8>

# Progressive Growing: FadeIn

```
class FadeIn(Layer):
    @tf.function
    def call(self, input_alpha, a, b):
        alpha = tf.reduce_mean(input_alpha)
        y = alpha * a + (1. - alpha) * b
        return y
```

# Progressive Growing: FadeIn

```
class FadeIn(Layer):
    @tf.function
    def call(self, input_alpha, a, b):
        alpha = tf.reduce_mean(input_alpha)
        y = alpha * a + (1. - alpha) * b
        return y
```

Why  
reduce\_mean?!

# Progressive Growing: FadeIn

```
class FadeIn(Layer):  
    @tf.function  
    def call(self, input_alpha, a, b):  
        alpha = tf.reduce_mean(input_alpha)  
        y = alpha * a + (1. - alpha) * b  
        return y
```

Why  
reduce\_mean?!

alpha is an  
input!!  
 $x = \text{FadeIn}(\text{alpha}, x, y)$

# Progressive Growing: FadeIn

```
class FadeIn(Layer):
    @tf.function
    def call(self, input_alpha, a, b):
        alpha = tf.reduce_mean(input_alpha)
        y = alpha * a + (1. - alpha) * b
        return y
```

Why  
reduce\_mean?!

alpha is an  
input!!

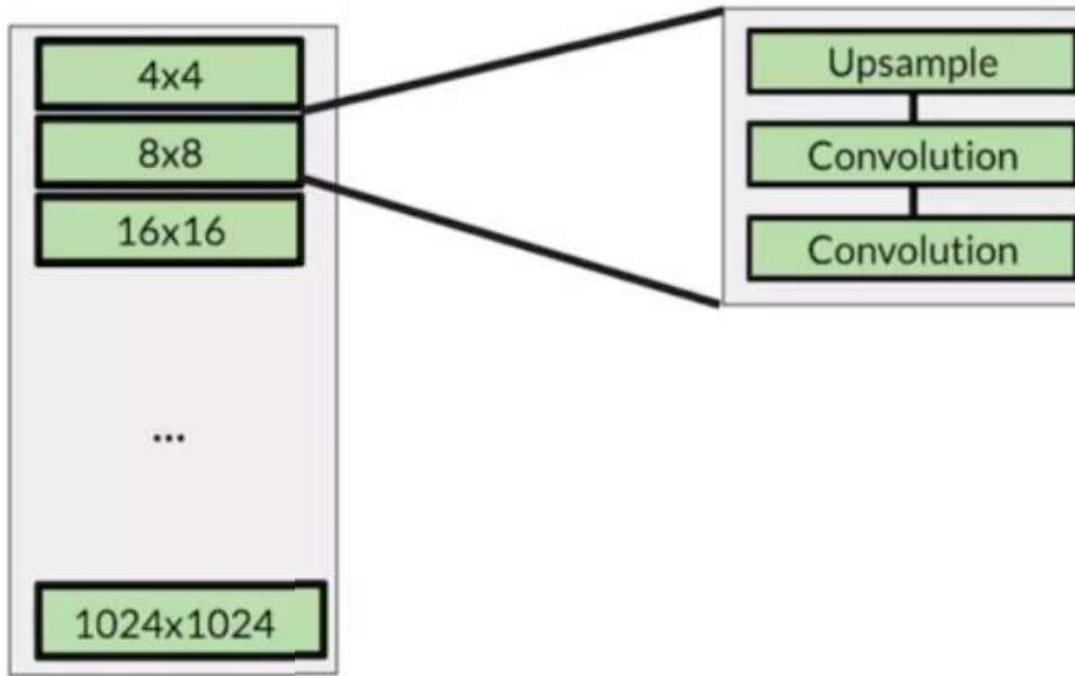
x = FadeIn()(alpha  
, x, y)

In this implementation, alpha is an input to our model.  
As the model input is assumed to be a minibatch, `tf.reduce_mean()` in `FadeIN()` is meant to convert the batched value to a scalar value.

# Progressive Growing in Context



# Progressive Growing in Context



# Progressive Growing

The implementation uses two functions that build the blocks to convert into/from RGB images.

- ❑ Both blocks use a 1x1 convolutional layer
- ❑ **to\_rgb** block uses a filter size of 3 to match the RGB channels
- ❑ **from\_rgb** blocks use a filter size that matches the input activation of the discriminator block at that scale.

# Progressive Growing

```
def build_to_rgb(self, res, filter_n):
    return Sequential([Input(shape=(res, res, filter_n)),
                      Conv2D(3, 1, gain=1,
                             activation='tanh')])

def build_from_rgb(self, res, filter_n):
    return Sequential([Input(shape=(res, res, 3)),
                      Conv2D(filter_n, 1),
                      LeakyReLU(0.2)])
```

# Progressive Growing: Generator

Now, we can look at the following steps to grow the generator to 32x32 (in this example!)

1. Add a 4x4 generator  
the input is a latent vector
2. In a loop, add subsequent generators of gradually increasing resolutions until the one before the target resolution  
in our example, 16x16
3. Add `to_rgb` from 16x16 and upsample it to 32x32.
4. Add the 32x32 generator block.
5. Fade in the two images to create one final RGB image.





```

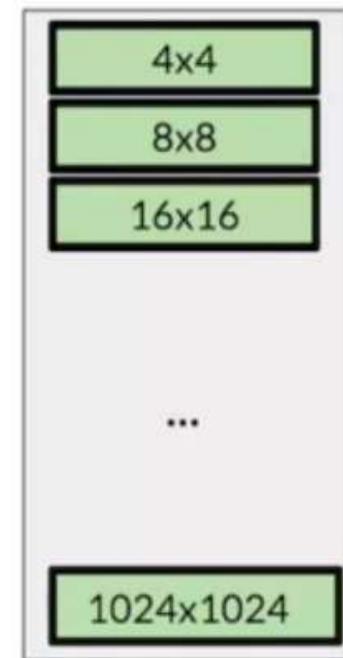
def grow_generator(self, log2_res):
    res = 2**log2_res
    alpha = Input(shape=(1))

    x = self.generator_blocks[2].input
    for i in range(2, log2_res):
        x = self.generator_blocks[i](x)

    old_rgb = self.to_rgb[log2_res-1](x)
    old_rgb = UpSampling2D((2,2))(old_rgb)

    x = self.generator_blocks[log2_res](x)
    new_rgb = self.to_rgb[log2_res](x)
    rgb = FadeIn()(alpha, new_rgb, old_rgb)
    self.generator = Model([self.generator_blocks[2].input, alpha],
                          [rgb,
                           name=f'generator_{res}_x_{res}'])

```



```
def build_all_generators(self):
    # build all the generator block
    self.to_rgb = {}
    self.generator_blocks = {}
    self.generator_blocks[2] = self.build_generator_base(self.z_dim)
    self.to_rgb[2] = self.build_to_rgb(4, self.log2_res_to_filter_size[2])

    for log2_res in range(3, self.log2_resolution+1):
        res = 2**log2_res
        filter_n = self.log2_res_to_filter_size[log2_res]
        self.to_rgb[log2_res] = self.build_to_rgb(res, filter_n)

        input_shape = self.generator_blocks[log2_res-1].output[0].shape
        gen_block = self.build_generator_block(log2_res, input_shape)
        self.generator_blocks[log2_res] = gen_block
```

```
def build_all_generators(self):
    # build all the generator block
    self.to_rgb = {}
    self.generator_blocks = {}
    self.generator_blocks[2] = self.build_generator_base(self.z_dim)
    self.to_rgb[2] = self.build_to_rgb(4, self.log2_res_to_filter_size[2])

    for log2_res in range(3, self.log2_resolution+1):
        res = 2**log2_res
        filter_n = self.log2_res_to_filter_size[log2_res]
        self.to_rgb[log2_res] = self.build_to_rgb(res, filter_n)

        input_shape = self.generator_blocks[log2_res-1].output[0].shape
        gen_block = self.build_generator_block(log2_res, input_shape)
        self.generator_blocks[log2_res] = gen_block
```

```
def build_generator_base(self, input_shape):  
    input_tensor = Input(shape=input_shape)  
    x = PixelNorm()(input_tensor)  
    x = Dense(8192, gain=1./8)(x)  
    x = Reshape((4, 4, 512))(x)  
    x = LeakyReLU(0.2)(x)  
    x = PixelNorm()(x)  
  
    x = Conv2D(512, 3,  
              name='gen_4x4_conv1')(x)  
    x = LeakyReLU(0.2)(x)  
    x = PixelNorm()(x)  
  
    return Model(input_tensor, x,  
                name='generator base')
```

The input is a latent vector

```
se(self.z_dim)  
s_to_filter_size[2])  
es]  
filter_n)  
.output[0].shape  
, input_shape)
```

```
def build_all_generators(self):
    # build all the generator block
    self.to_rgb = {}
    self.generator_blocks = {}
    self.generator_blocks[2] = self.build_generator_base(self.z_dim)
    self.to_rgb[2] = self.build_to_rgb(4, self.log2_res_to_filter_size[2])

    for log2_res in range(3, self.log2_resolution+1):
        res = 2**log2_res
        filter_n = self.log2_res_to_filter_size[log2_res]
        self.to_rgb[log2_res] = self.build_to_rgb(res, filter_n)

        input_shape = self.generator_blocks[log2_res-1].output[0].shape
        gen_block = self.build_generator_block(log2_res, input_shape)
        self.generator_blocks[log2_res] = gen_block
```

```
def build_all_generator(self):
    self.log2_res_to_filter_size = {
        0: 512,
        1: 512,
        2: 512, # 4x4
        3: 512, # 8x8
        4: 512, # 16x16
        5: 512, # 32x32
        6: 256, # 64x64
        7: 128, # 128x128
        8: 64, # 256x256
        9: 32, # 512x512
        10: 16} # 1024x1024
    self.generator_blocks[log2_res] = gen_block(
        generator_base(self.z_dim),
        self.log2_res_to_filter_size[2])
    for i in range(log2_res+1):
        self.generator_blocks[i] = gen_block(
            filter_size=log2_res_to_filter_size[log2_res],
            id_to_rgb(res, filter_n))
    self.generator_blocks[log2_res-1].output[0].shape = blocks[log2_res-1].output[0].shape
    self.generator_blocks[log2_res-1].block(log2_res, input_shape)
```

```
def build_all_generators(self):
    # build all the generator block
    self.to_rgb = {}
    self.generator_blocks = {}
    self.generator_blocks[2] = self.build_generator_base(self.z_dim)
    self.to_rgb[2] = self.build_to_rgb(4, self.log2_res_to_filter_size[2])

    for log2_res in range(3, self.log2_resolution+1):
        res = 2**log2_res
        filter_n = self.log2_res_to_filter_size[log2_res]
        self.to_rgb[log2_res] = self.build_to_rgb(res, filter_n)

        input_shape = self.generator_blocks[log2_res-1].output[0].shape
        gen_block = self.build_generator_block(log2_res, input_shape)
        self.generator_blocks[log2_res] = gen_block
```

```
def build_all_generators(self):
    # build all the generator block
    self.to_rgb = {}
    self.generator_blocks = {}
    self.generator_blocks[2] = self.build_generator_base(self.z_dim)
    self.to_rgb[2] = self.build_to_rgb(resolution, filter_size[2])
    self.log2_resolution = int(np.log2(resolution))

    for log2_res in range(3, self.log2_resolution+1):
        res = 2**log2_res
        filter_n = self.log2_res_to_filter_size[log2_res]
        self.to_rgb[log2_res] = self.build_to_rgb(res, filter_n)

        input_shape = self.generator_blocks[log2_res-1].output[0].shape
        gen_block = self.build_generator_block(log2_res, input_shape)
        self.generator_blocks[log2_res] = gen_block
```

```
def build_all_generators(self):
    # build all the generator block
    self.to_rgb = {}
    self.generator_blocks = {}
    self.generator_blocks[2] = self.build_generator_base(self.z_dim)
    self.to_rgb[2] = self.build_to_rgb(4, self.log2_res_to_filter_size[2])

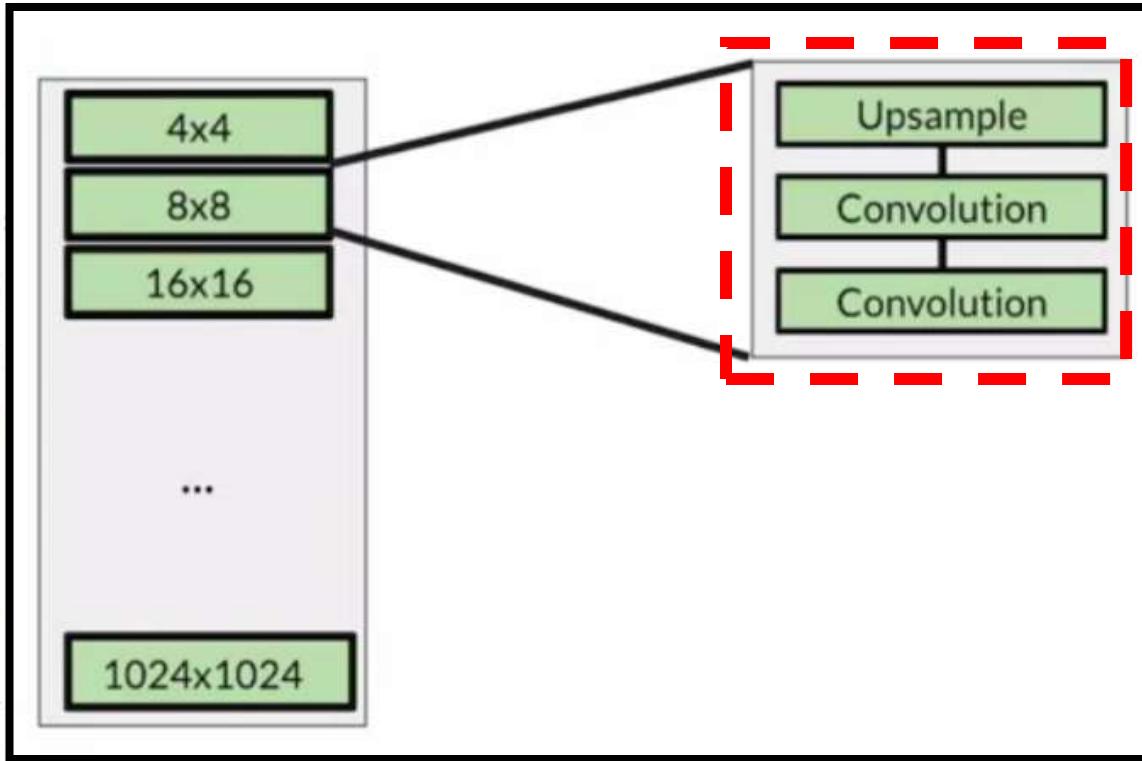
    for log2_res in range(3, self.log2_resolution+1):
        res = 2**log2_res
        filter_n = self.log2_res_to_filter_size[log2_res]
        self.to_rgb[log2_res] = self.build_to_rgb(res, filter_n)

        input_shape = self.generator_blocks[log2_res-1].output[0].shape
        gen_block = self.build_generator_block(log2_res, input_shape)
        self.generator_blocks[log2_res] = gen_block
```

```
def build_all_generator_blocks(self):
    # build all generator blocks
    self.to_rgb = []
    self.generator_blocks = []
    self.generator_blocks.append(self.build_generator_block(4, 1024))
    self.to_rgb.append(self.to_rgb[0])
    log2_res = 1
    res = 2**log2_res
    filter_n = 1024
    self.to_rgb.append(self.to_rgb[0])
    self.generator_blocks.append(self.build_generator_block(log2_res, filter_n))

    for log2_res in range(1, 5):
        res = 2**log2_res
        filter_n = 512
        self.to_rgb.append(self.to_rgb[0])
        self.generator_blocks.append(self.build_generator_block(log2_res, filter_n))

    input_shape = self.generator_blocks[log2_res-1].output[0].shape
    gen_block = self.build_generator_block(log2_res, input_shape)
    self.generator_blocks[log2_res] = gen_block
```



```
def build_generator_block(self, log2_res, input_shape):
    res = 2**log2_res
    res_name = f'{res}x{res}'
    filter_n = self.log2_res_to_filter_size[log2_res]

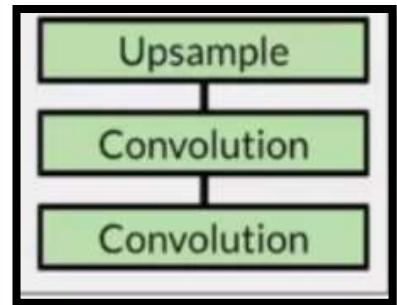
    input_tensor = Input(shape=input_shape)
    x = UpSampling2D((2,2))(input_tensor)

    x = Conv2D(filter_n, 3, name=f'gen_{res_name}_conv1')(x)
    x = LeakyReLU(0.2)(x)
    x = PixelNorm()(x)

    x = Conv2D(filter_n, 3, name=f'gen_{res_name}_conv2')(x)
    x = LeakyReLU(0.2)(x)
    x = PixelNorm()(x)

    return Model(input_tensor, x,
                 name=f'genblock_{res}_x_{res}')
```

```
gen_block = self.build_generator_block(log2_res, input_shape)
self.generator_blocks[log2_res] = gen_block
```



# Progressive Growing: Discriminator

The growing discriminator is done similarly but in the reverse direction

1. At the resolution of the input image, say, 32x32, add `from_rgb` to the discriminator block of 32x32. The output is the activation with a 16x16 feature map.
2. Parallelly, downsample the input image to 16x16, and add `from_rgb` to the 16x16 discriminator block.
3. Fade in the two preceding features and feed that into the next discriminator block of 8x8.
4. Continue adding the discriminator blocks to the base of 4x4, where the output is a single prediction value.

```
def grow_discriminator(self, log2_res):
    res = 2**log2_res

    input_image = Input(shape=(res, res, 3))
    alpha = Input(shape=(1))

    x = self.from_rgb[log2_res](input_image)
    x = self.discriminator_blocks[log2_res](x)

    downsized_image = AveragePooling2D((2,2))(input_image)
    y = self.from_rgb[log2_res-1](downsized_image)

    x = FadeIn()(alpha, x, y)
    for i in range(log2_res-1, 1, -1):
        x = self.discriminator_blocks[i](x)

    self.discriminator = Model([input_image, alpha], x,
                               name=f'discriminator_{res}_x_{res}')
    self.optimizer_discriminator = Adam(**self.opt_init)
```

```
def grow_discriminator(self, log2_res):
    res = 2**log2_res

    input_image = Input(shape=(res, res, 3))
    alpha = Input(shape=(1))

    x = self.from_rgb[log2_res](input_image)
    x = self.discriminator_blocks[log2_res](x)

    downsized_image = AveragePooling2D((2,2))(input_image)
    y = self.from_rgb[log2_res-1](downsized_image)

    x = FadeIn()(alpha, x, y)
    for i in range(log2_res-1, 1, -1):
        x = self.discriminator_blocks[i](x)

    self.discriminator = Model([input_image, alpha], x,
                               name=f'discriminator_{res}_x_{res}')
    self.optimizer_discriminator = Adam(**self.opt_init)
```

```
def grow_discriminator(self, log2_res):
    res = 2**log2_res

def build_all_discriminators(self):
    self.from_rgb = {}
    self.discriminator_blocks = {}

    # all but the final block
    for log2_res in range(self.log2_resolution, 1, -1):
        res = 2**log2_res
        filter_n = self.log2_res_to_filter_size[log2_res]
        self.from_rgb[log2_res] = self.build_from_rgb(res, filter_n)

        input_shape = (res, res, filter_n)
        self.discriminator_blocks[log2_res] = self.build_discriminator_block(
            log2_res, input_shape)

    # last block at 4x4 resolution
    log2_res = 2
    filter_n = self.log2_res_to_filter_size[log2_res]
    self.from_rgb[log2_res] = self.build_from_rgb(4, filter_n)
    res = 2**log2_res
    input_shape = (res, res, filter_n)
    self.discriminator_blocks[log2_res] = self.build_discriminator_base(input_shape)

    self.optimizer_discriminator = Adam(**self.opt_init)
```

```
def grow_discriminator(self, log2_res):
    res = 2**log2_res

def build_all_discriminators(self):
    self.from_rgb = {}
    self.discriminator_blocks = {}

    # all but the final block
    for log2_res in range(self.log2_resolution, 1, -1):
        res = 2**log2_res
        filter_n = self.log2_res_to_filter_size[log2_res]
        self.from_rgb[log2_res] = self.build_from_rgb(res, filter_n)

        input_shape = (res, res, filter_n)
        self.discriminator_blocks[log2_res] = self.build_discriminator_block(
            log2_res, input_shape)

    # last block at 4x4 resolution
    log2_res = 2
    filter_n = self.log2_res_to_filter_size[log2_res]
    self.from_rgb[log2_res] = self.build_from_rgb(4, filter_n)
    res = 2**log2_res
    input_shape = (res, res, filter_n)
    self.discriminator_blocks[log2_res] = self.build_discriminator_base(input_shape)

    self.optimizer_discriminator = Adam(**self.opt_init)
```

```
def build_discriminator_block(self, log2_res, input_shape):  
  
    filter_n = self.log2_res_to_filter_size[log2_res]  
    input_tensor = Input(shape=input_shape)  
  
    # First conv  
    x = Conv2D(filter_n, 3)(input_tensor)  
    x = LeakyReLU(0.2)(x)  
  
    # Second conv + downsample  
    filter_n = self.log2_res_to_filter_size[log2_res-1]  
    x = Conv2D(filter_n, 3, )(x)  
    x = LeakyReLU(0.2)(x)  
    x = AveragePooling2D((2,2))(x)  
  
    res = 2**log2_res  
    return Model(input_tensor, x,  
                name=f'disc_block_{res}_x_{res}')
```

```
input_shape = (res, res, filter_n)
```

```
self.discriminator_blocks[log2_res] = self.build_discriminator_base(input_shape)
```

```
self.optimizer_discriminator = Adam(**self.opt_init)
```

# Summary

- Progressive growing gradually doubles image resolution
- Helps with faster, more stable training for higher resolutions

# Pixel normalization

# How to define a norm layer!

The purpose of normalization in ProGAN is to limit the weight values to prevent them from growing exponentially.

$$b_{x,y} = a_{x,y} / \sqrt{\frac{1}{N} \sum_{j=0}^{N-1} (a_{x,y}^j)^2 + \epsilon}$$

```
class PixelNorm(Layer):
    def __init__(self, epsilon=1e-8):
        super(PixelNorm, self).__init__()
        self.epsilon = epsilon

    def call(self, input_tensor):
        return input_tensor / tf.math.sqrt(
            tf.reduce_mean(input_tensor**2,
                           axis=-1, keepdims=True) +
            self.epsilon)
```

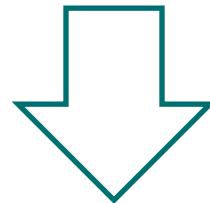
# Minibatch statistics

# Problem!

- **Mode collapse** happens when a GAN generates similar-looking images as it captures only a subset of the variation found in the training data.

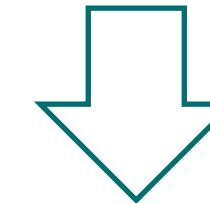
# Encourage image variation with mini-batch statistics

4 3 7



High std

4 4 4



Low std

Use standard deviation in batch to encourage diversity.

# Implementation

The steps to calculate minibatch statistics are as follows:

1. Calculate the standard deviation for each feature in each spatial location over the minibatch.

```
_ , group_var = tf.nn.moments(x, axes=(0), keepdims=False)  
group_std = tf.sqrt(group_var + self.epsilon)
```

2. Calculate the average of these standard deviations across the  $(H, W, C)$  dimensions to arrive at a single scale value.

```
avg_std = tf.reduce_mean(group_std, axis=[1,2,3], keepdims=True)
```

3. Replicate this value across the feature map of  $(H, W)$  and append it to the activation. As a result, the output activation has a shape of  $(N, H, W, C+1)$ .

```
x = tf.tile(avg_std, [self.group_size, h, w, 1])  
tf.concat([input_tensor, x], axis=-1)
```



# Implementation

Before calculating the standard deviation, the activation is first split into groups of 4 or the batch size, whichever is lower. To simplify the code, we assume that the batch size is at least 4 during training. The minibatch layer can be inserted anywhere in the discriminator, but it was found to be more effective toward the end, which is the 4x4 layer.

# Code...

```
class MinibatchStd(Layer):
    def __init__(self, group_size=4, epsilon=1e-8):

        super(MinibatchStd, self).__init__()
        self.epsilon = epsilon
        self.group_size = group_size

    def call(self, input_tensor):

        n, h, w, c = input_tensor.shape
        x = tf.reshape(input_tensor, [self.group_size, -1, h, w, c])
        group_mean, group_var = tf.nn.moments(x, axes=(0), keepdims=False)
        group_std = tf.sqrt(group_var + self.epsilon)
        avg_std = tf.reduce_mean(group_std, axis=[1,2,3], keepdims=True)
        x = tf.tile(avg_std, [self.group_size, h, w, 1])

        return tf.concat([input_tensor, x], axis=-1)
```

# Equalized learning rate

# Problem: Vanishing gradient

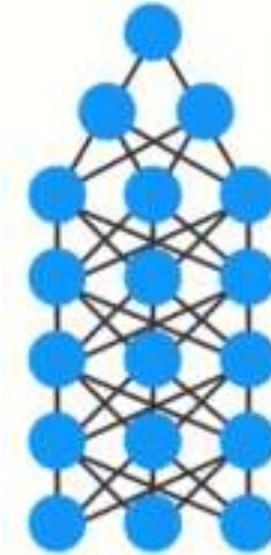
- vanishing gradient
  - The layers closer to the generator input will receive less gradient

$$w = w - \alpha \frac{\partial L}{\partial w}$$

=

$$w = w - \alpha \nabla$$

loss(Pred, Truth) = E



CLASS.  
vision

# Solution

- What if we want a layer to receive more gradient?

$$w = w - m \alpha \nabla$$

# Solution

- What if we want a layer to receive more gradient?

$$w = w - m \alpha \nabla$$

- Different **multiplier constants** for different layers to effectively have different learning rates.

# He or Kaiming initialization

- In ProGAN, these multiplier constants are calculated from He's initializer

$$std = \sqrt{\frac{2}{fan\ in}}$$

- **Fan in** is the multiplication of the weights dimension except for the output channel.

# He or Kaiming initialization

- In ProGAN, these multiplier constants are calculated from He's initializer

$$std = \sqrt{\frac{2}{fan\ in}}$$

- For a convolution weight with the shape (kernel, kernel, channel\_in, channel\_out), the *fan in* is the multiplication of kernel x kernel x channel\_in.

```
fan_in = self.kernel*self.kernel*self.in_channels  
self.scale = tf.sqrt(self.gain/fan_in)
```

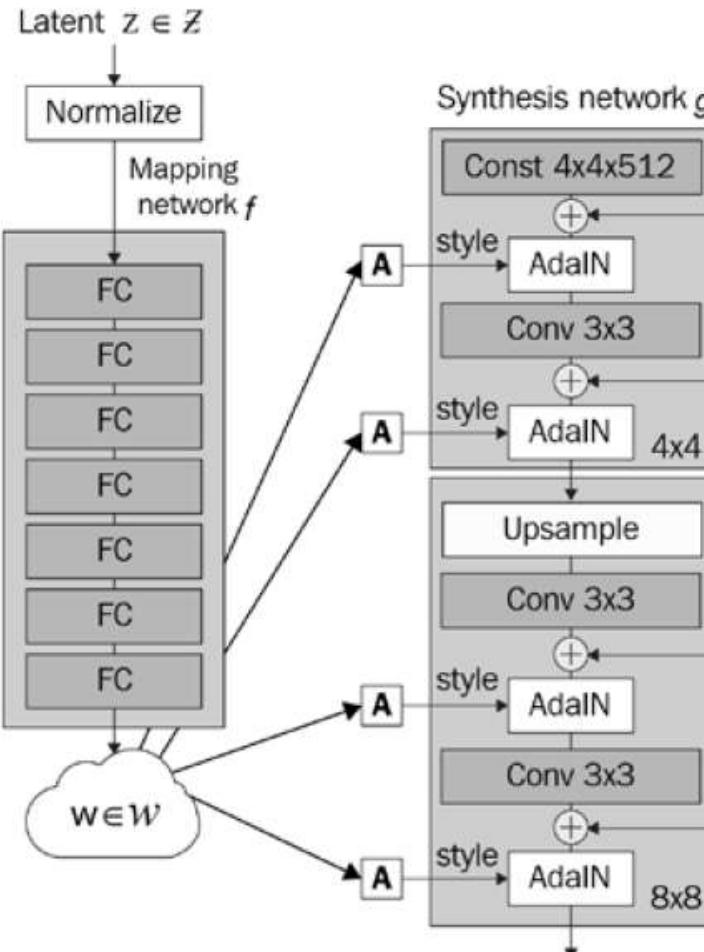
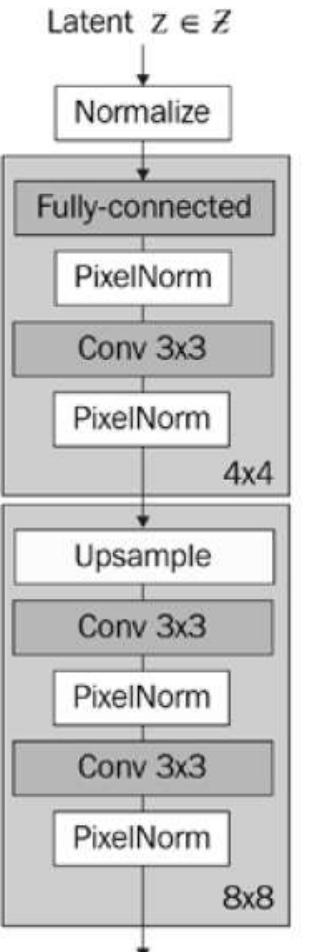


# implementation

- [https://colab.research.google.com/github/Alireza-Akhavan/GAN\\_tutorial/blob/main/progressive\\_gan.ipynb](https://colab.research.google.com/github/Alireza-Akhavan/GAN_tutorial/blob/main/progressive_gan.ipynb)

# StyleGAN

# Style-based generator



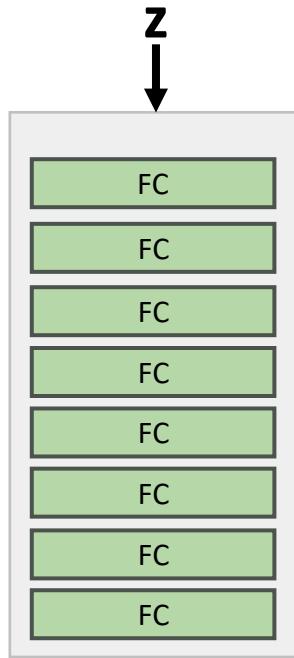
The diagram compares the generator architectures of ProGAN and StyleGAN

# Noise Mapping Network

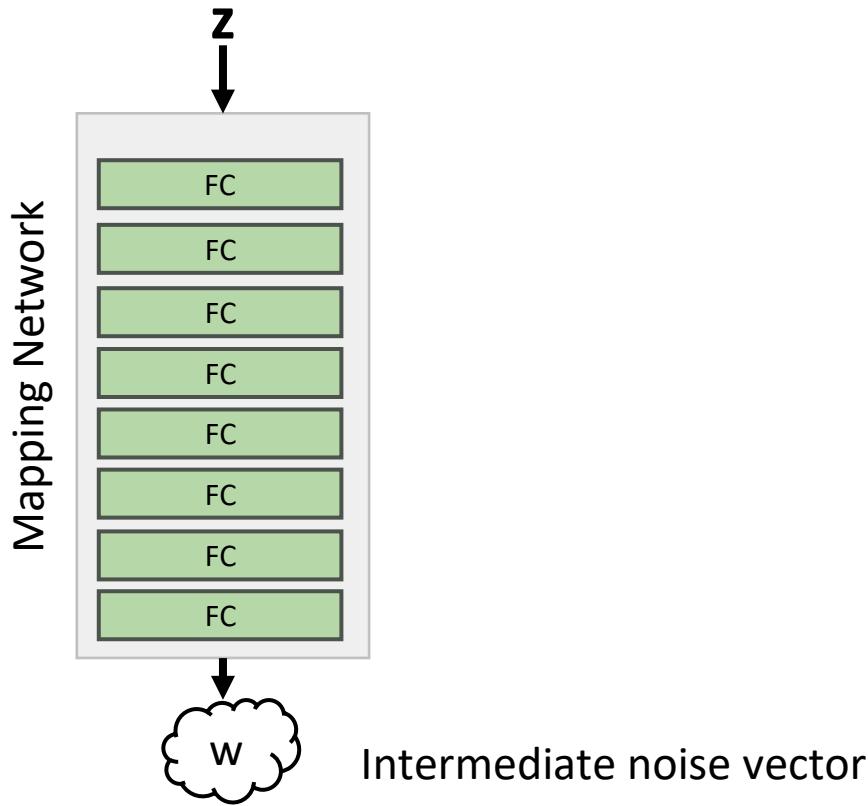
# Outline

- Noise mapping network structure
- Motivation behind the noise mapping network
- Where its output  $W$  goes

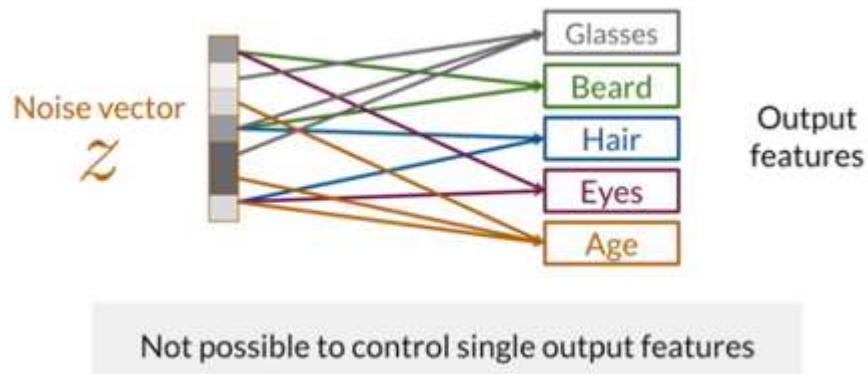
# Noise Mapping Network



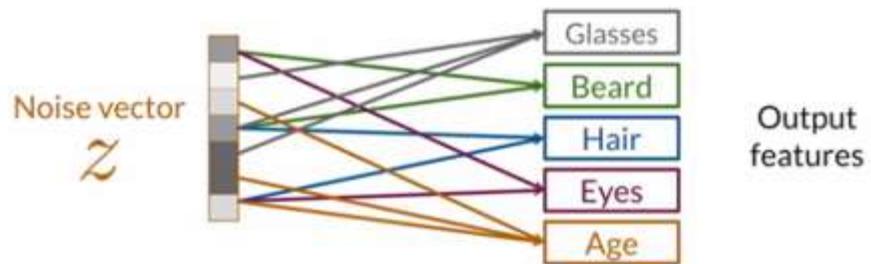
# Noise Mapping Network



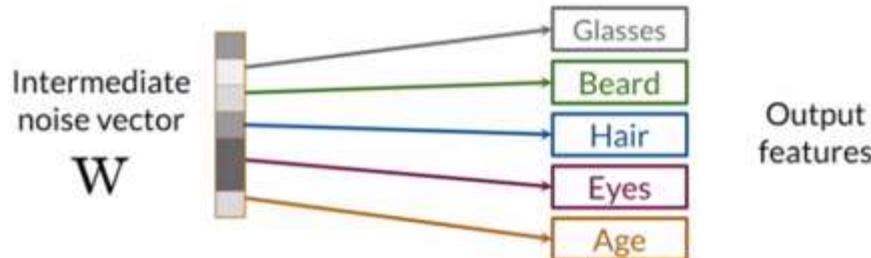
# Z-space Entanglement



# W-space: Less Entangled

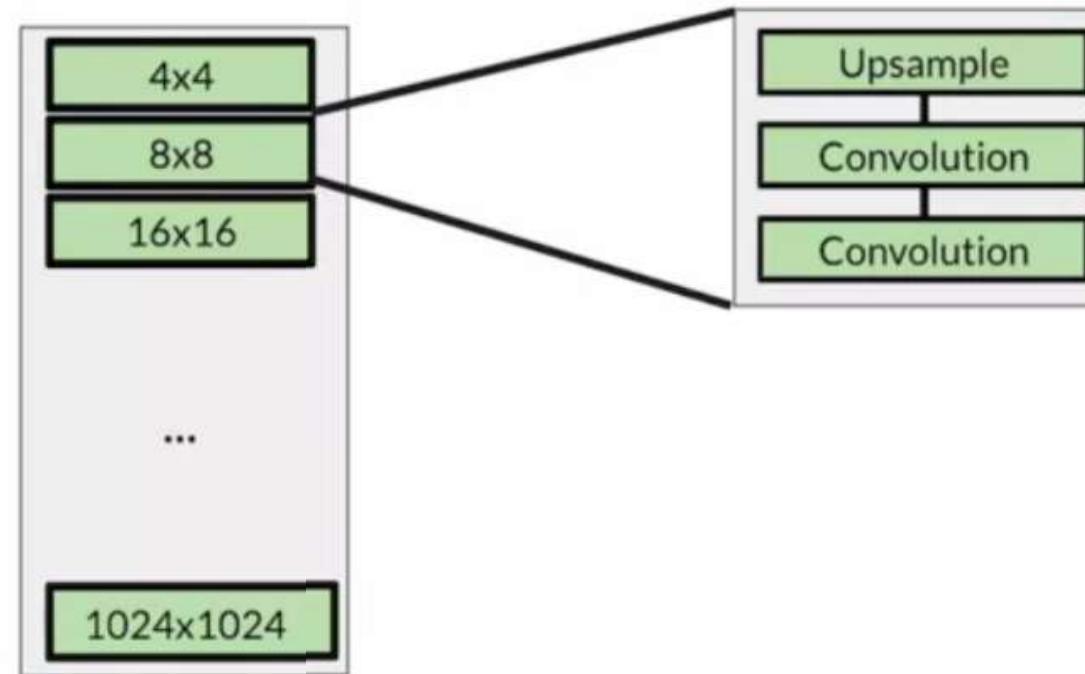


Not possible to control single output features

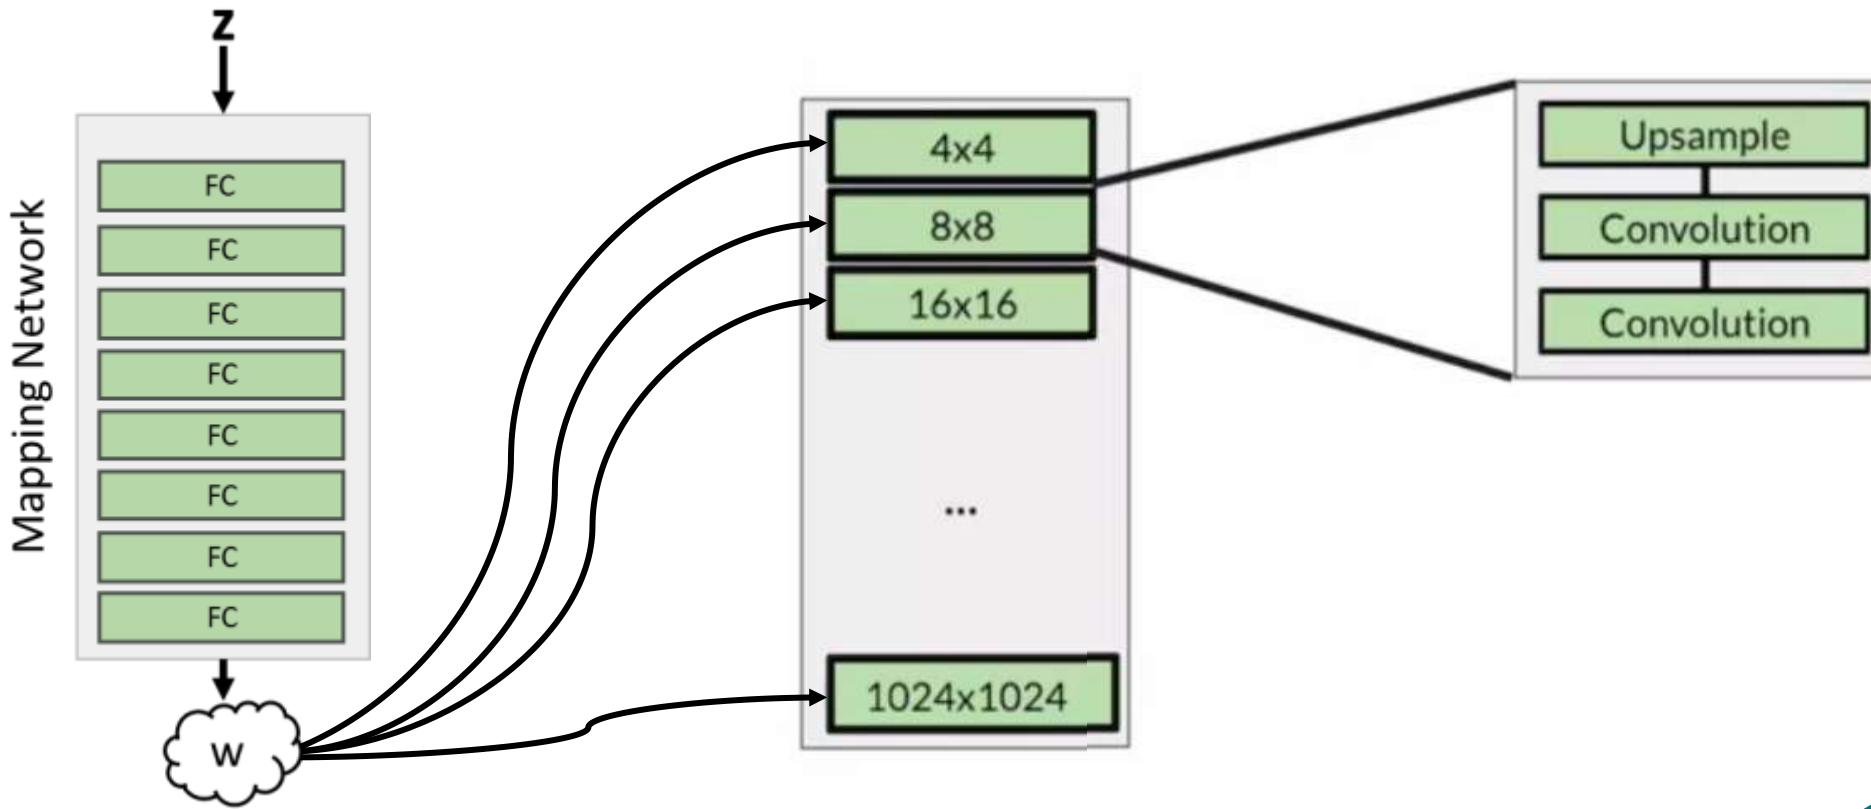


More possible to control single output features

# Mapping Network in Context



# Mapping Network in Context



# Summary

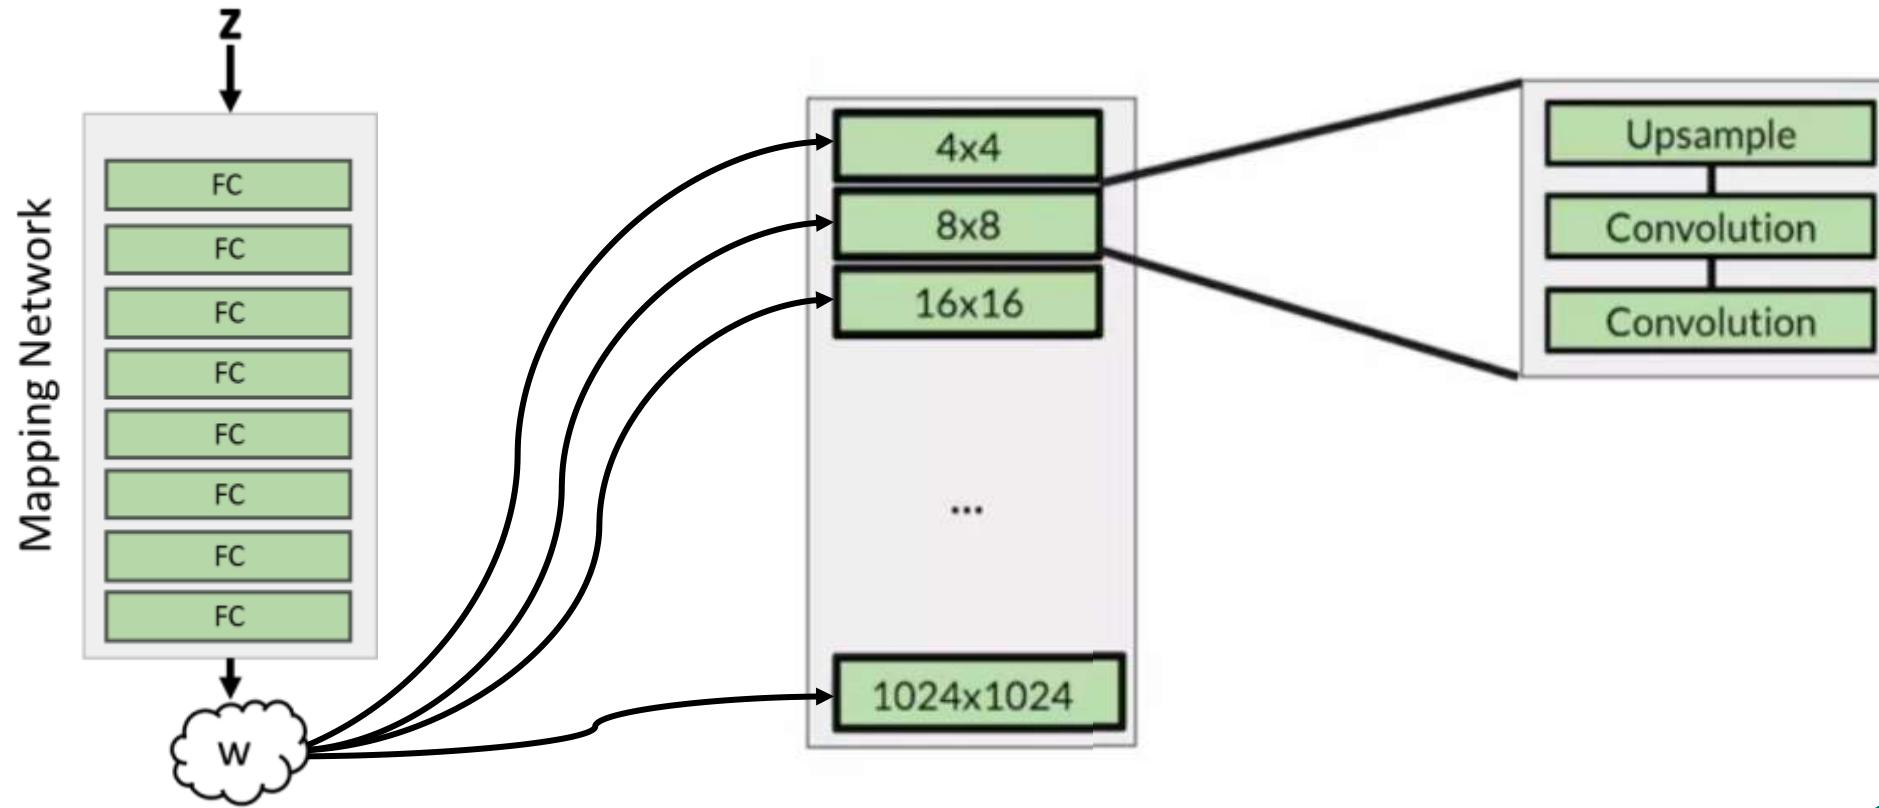
- Noise mapping allows for a more disentangled noise space
- The intermediate noise vector is used as input to the generator

# Adaptive Instance Normalization (AdaIN)

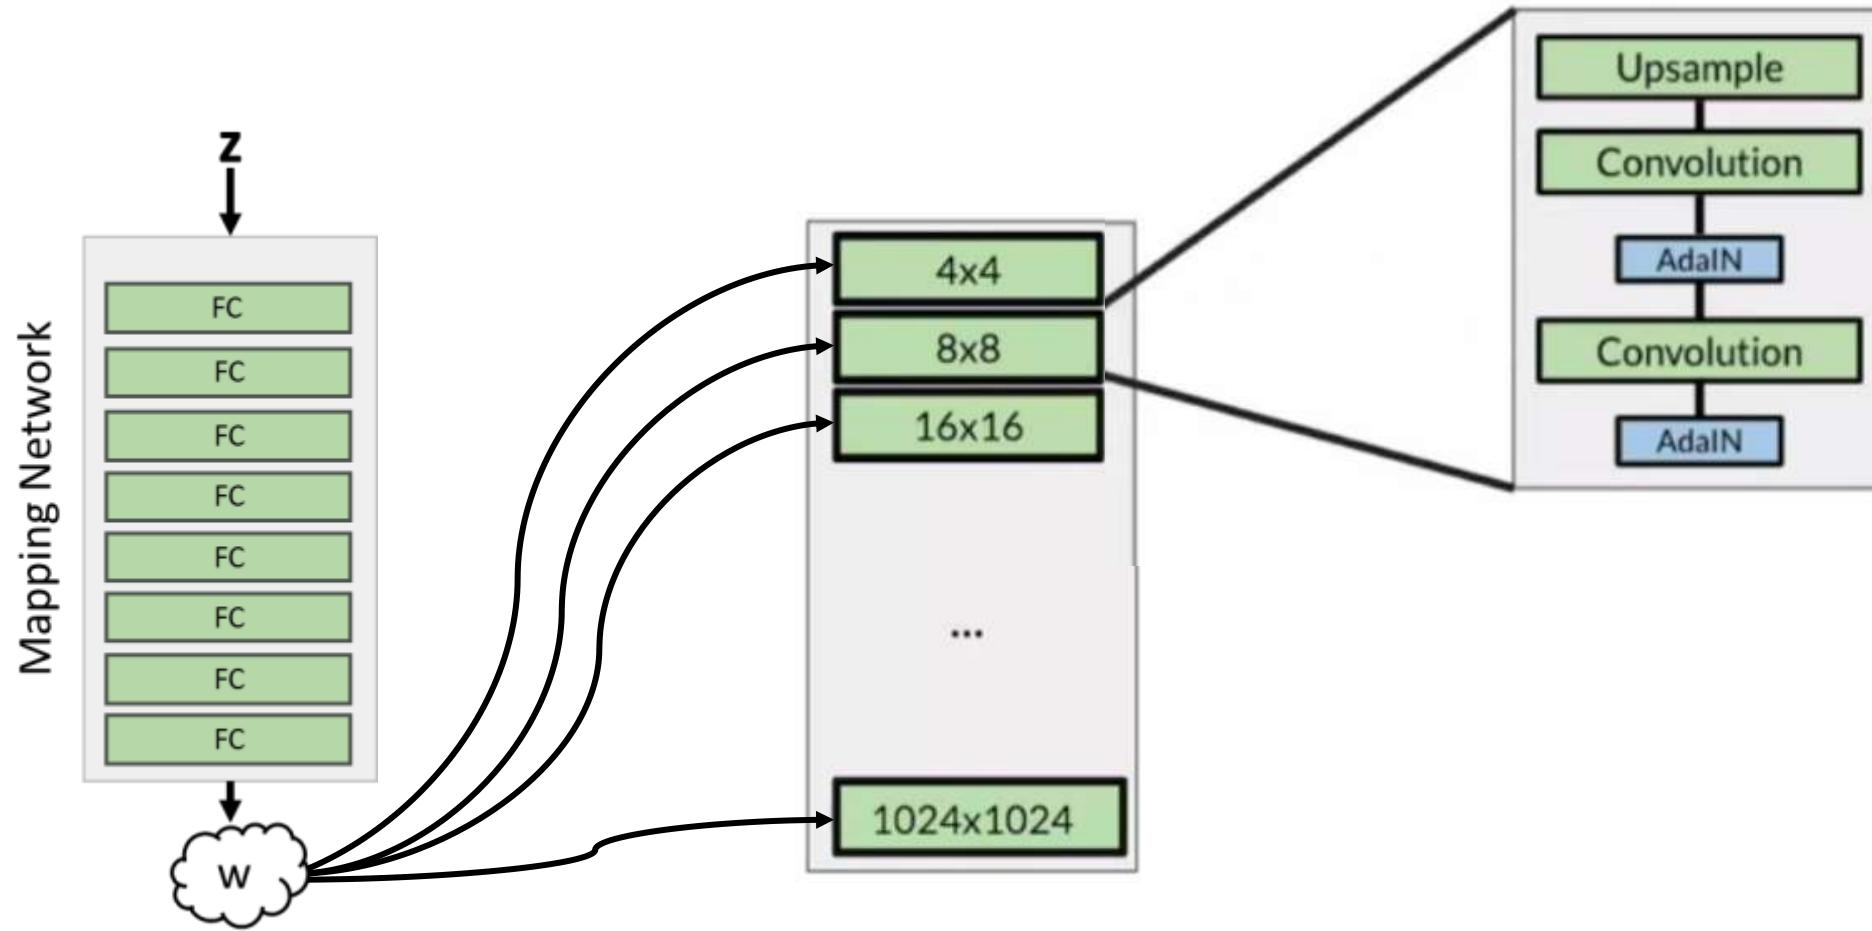
# Outline

- Instance Normalization
- Adaptive Instance Normalization
- Where and why AdaIN is used

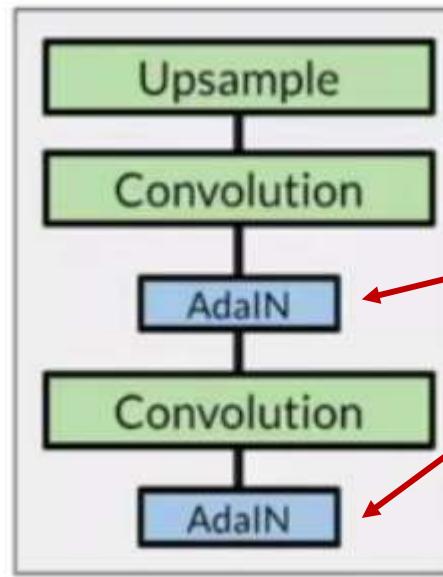
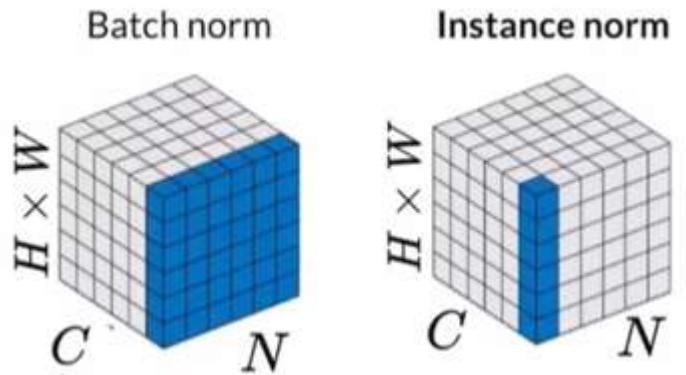
# AdaIN in Context



# AdaIN in Context



# AdaIN

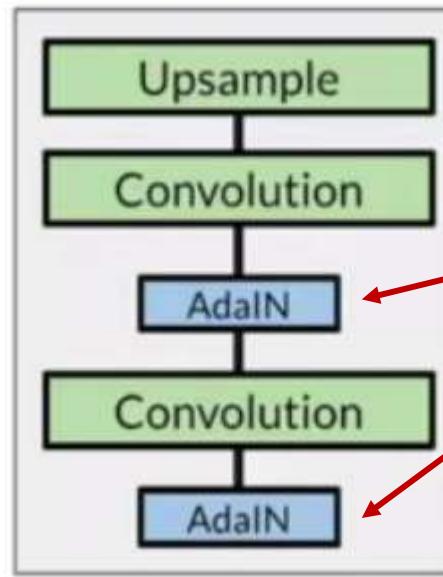
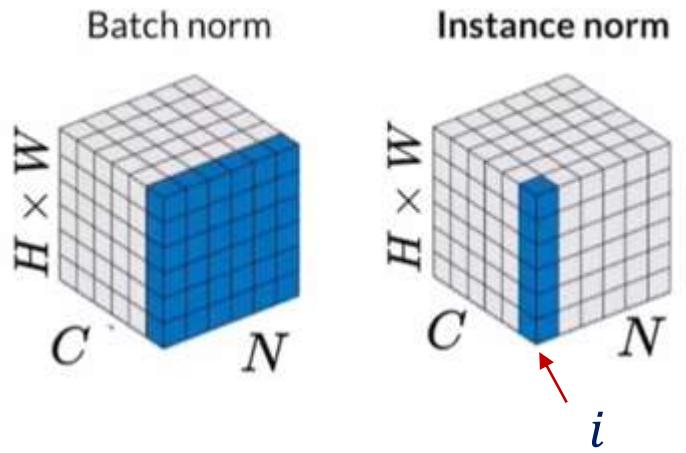


$$\frac{x - \mu(x)}{\delta(x)}$$

Step1: Normalize convolution outputs using Instance Normalization



# AdaIN

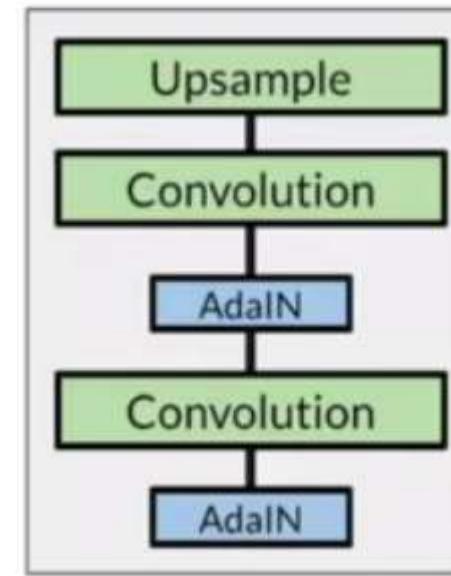


$$\frac{x_i - \mu(x_i)}{\delta(x_i)}$$

Step1: Normalize convolution outputs using Instance Normalization

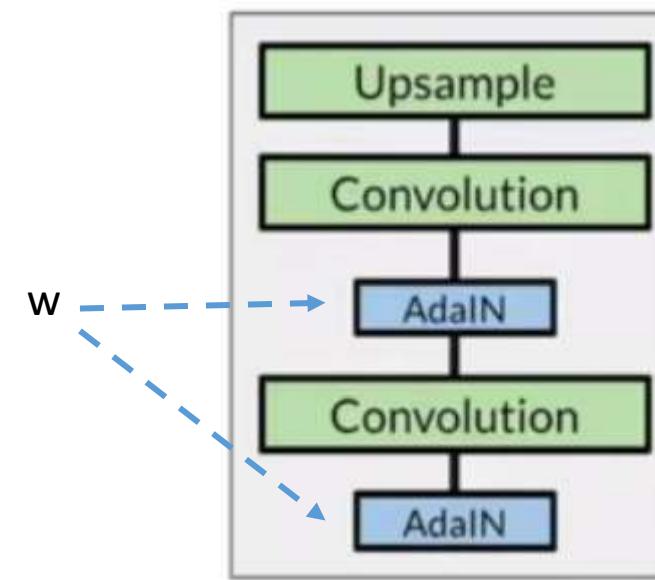
# AdaIN

$z \rightarrow$  Mapping Network  $\rightarrow w$



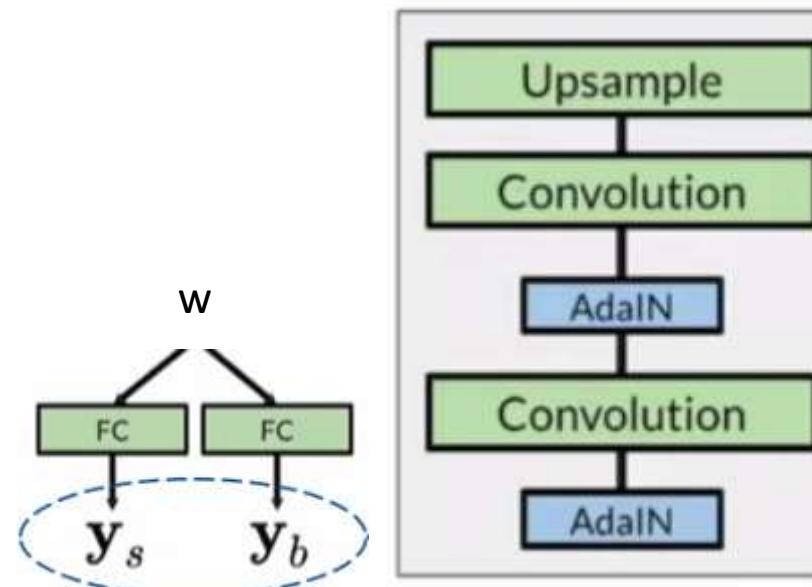
Step 2: Adaptive styles using the intermediate noise vector

# AdaIN



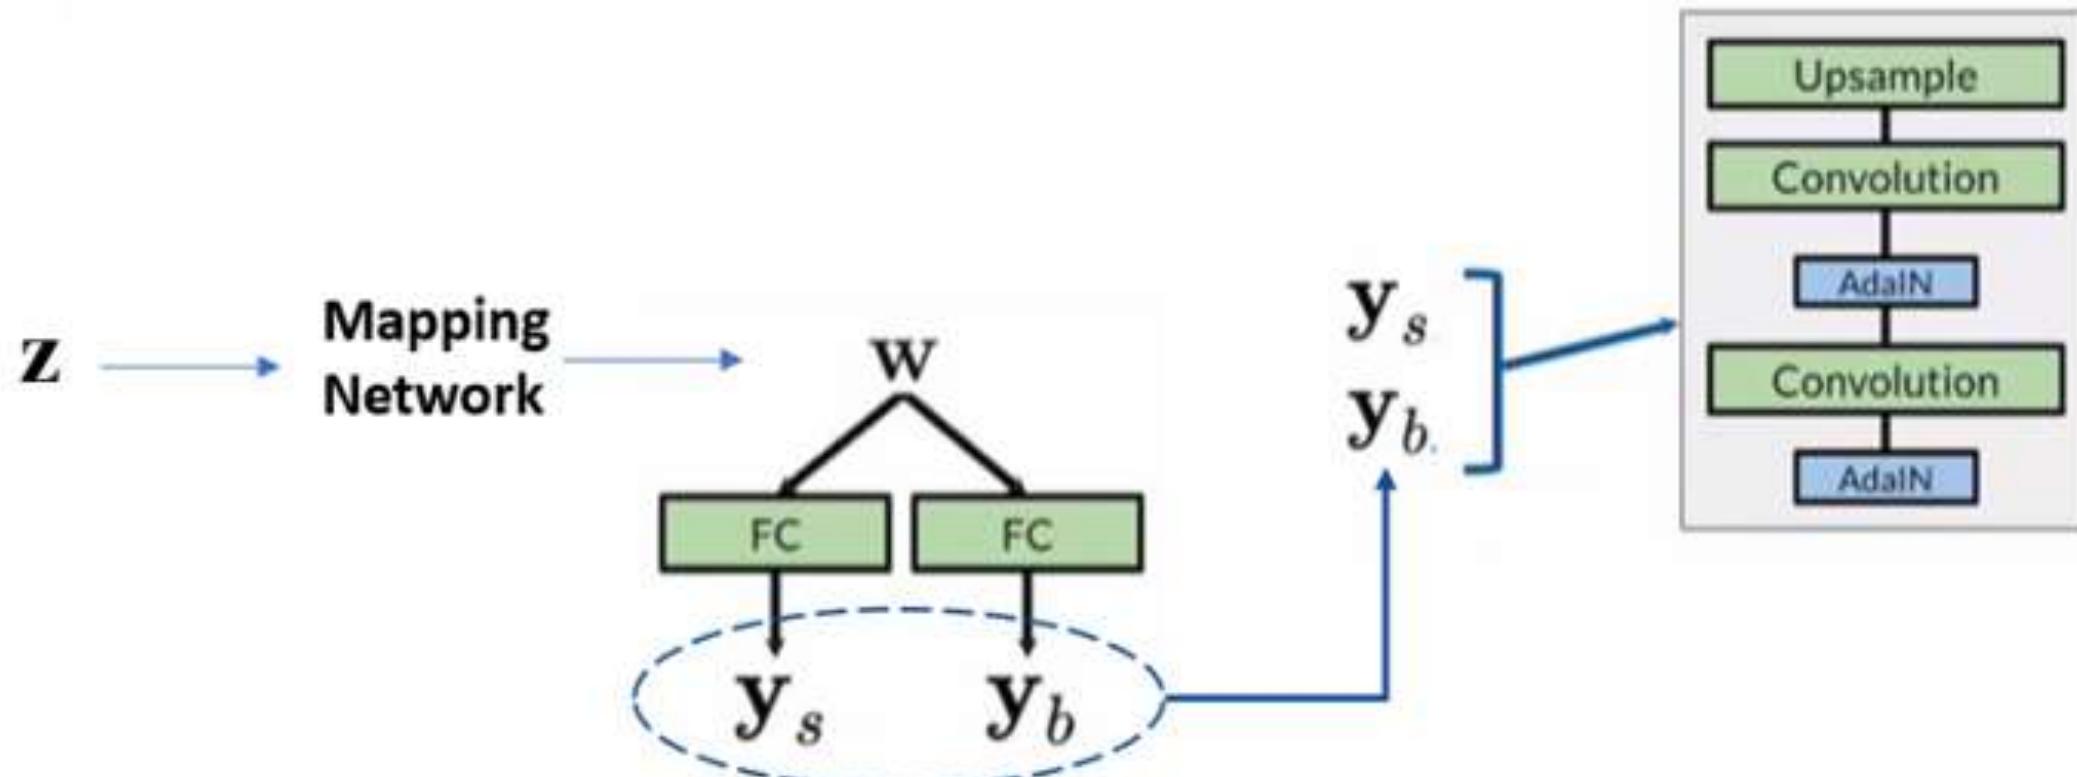
Step 2: Adaptive styles using the intermediate noise vector

# AdaIN



Step 2: Adaptive styles using the intermediate noise vector

# AdaIN



Step 2: Adaptive styles using the intermediate noise vector

# AdaIN

$$AdaIN(x_i, y) = y_{s,i} \frac{x_i - \mu(x_i)}{\delta(x_i)} + y_{b,i}$$

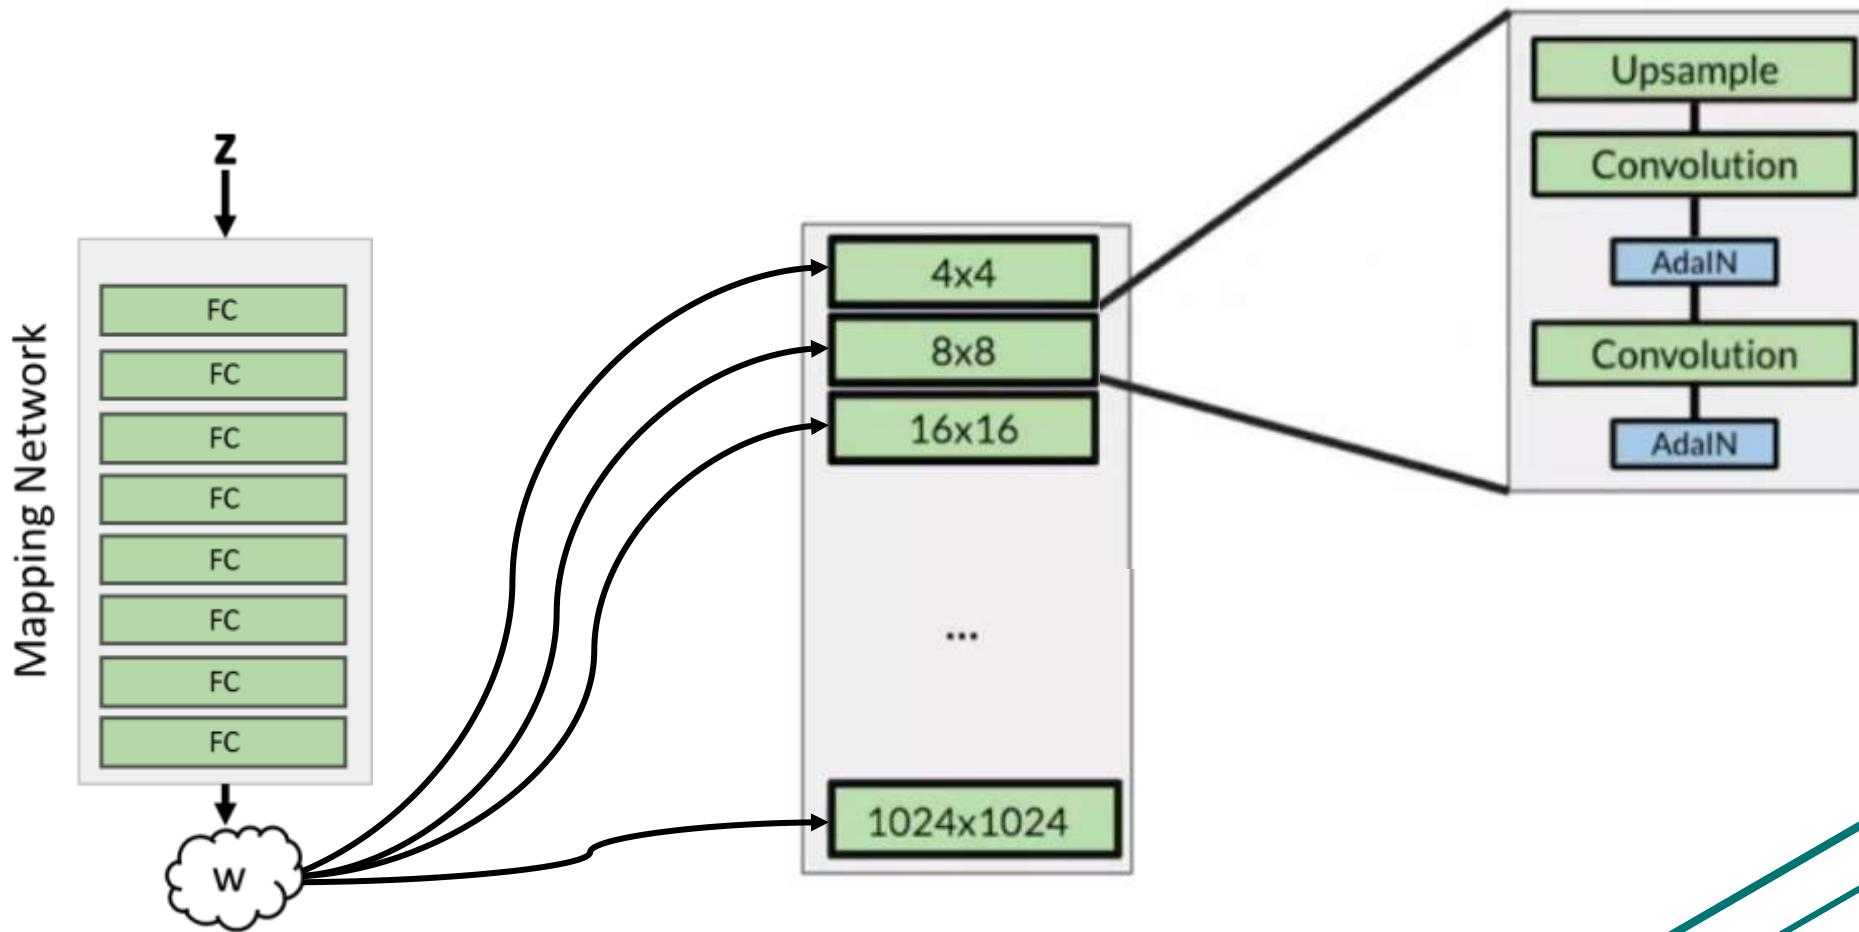
Step 1: Instance  
Normalization

# AdaIN

$$AdaIN(x_i, y) = y_{s,i} \frac{x_i - \mu(x_i)}{\delta(x_i)} + y_{b,i}$$

Step 2: Adaptive styles

# AdaIN in Context



# Summary

- AdaIN transfers style information onto the generated image from the intermediate noise vector  $w$
- Instance Normalization is used to normalize individual examples before apply style statistics from  $w$

# Style and Stochastic Variation

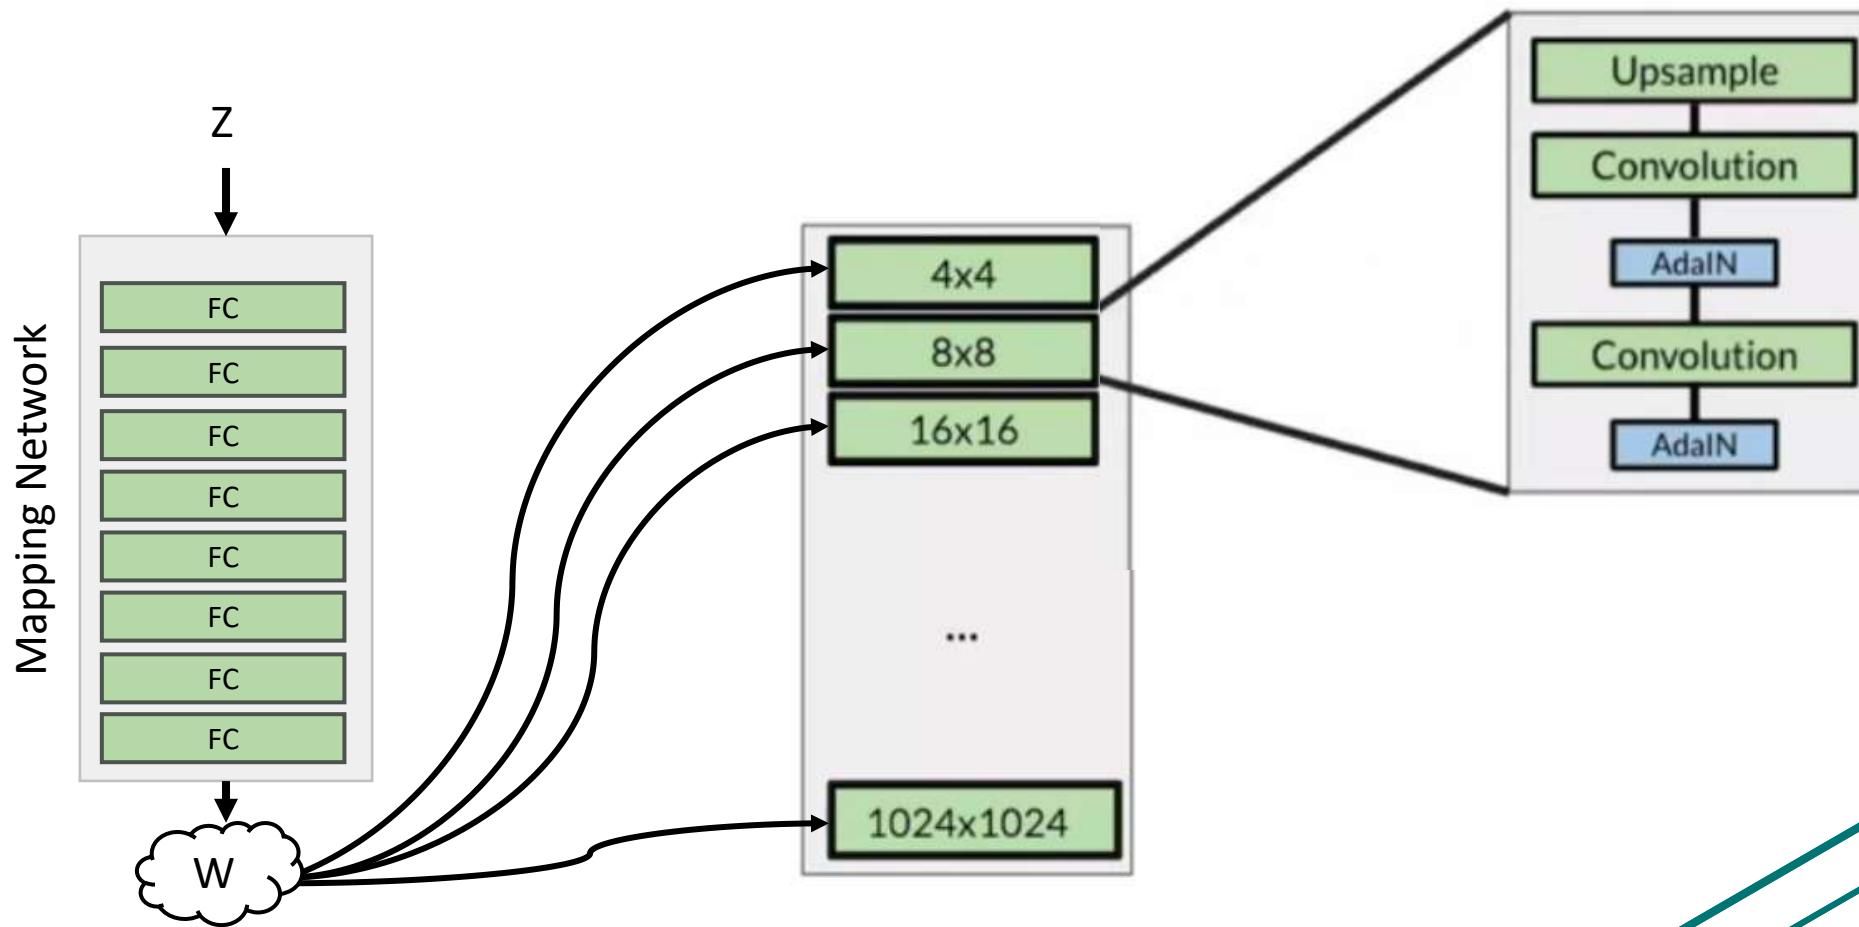
# Outline

- Controlling coarse and fine styles with StyleGAN
- Style mixing for increased diversity during training/inference
- Stochastic noise for additional variation

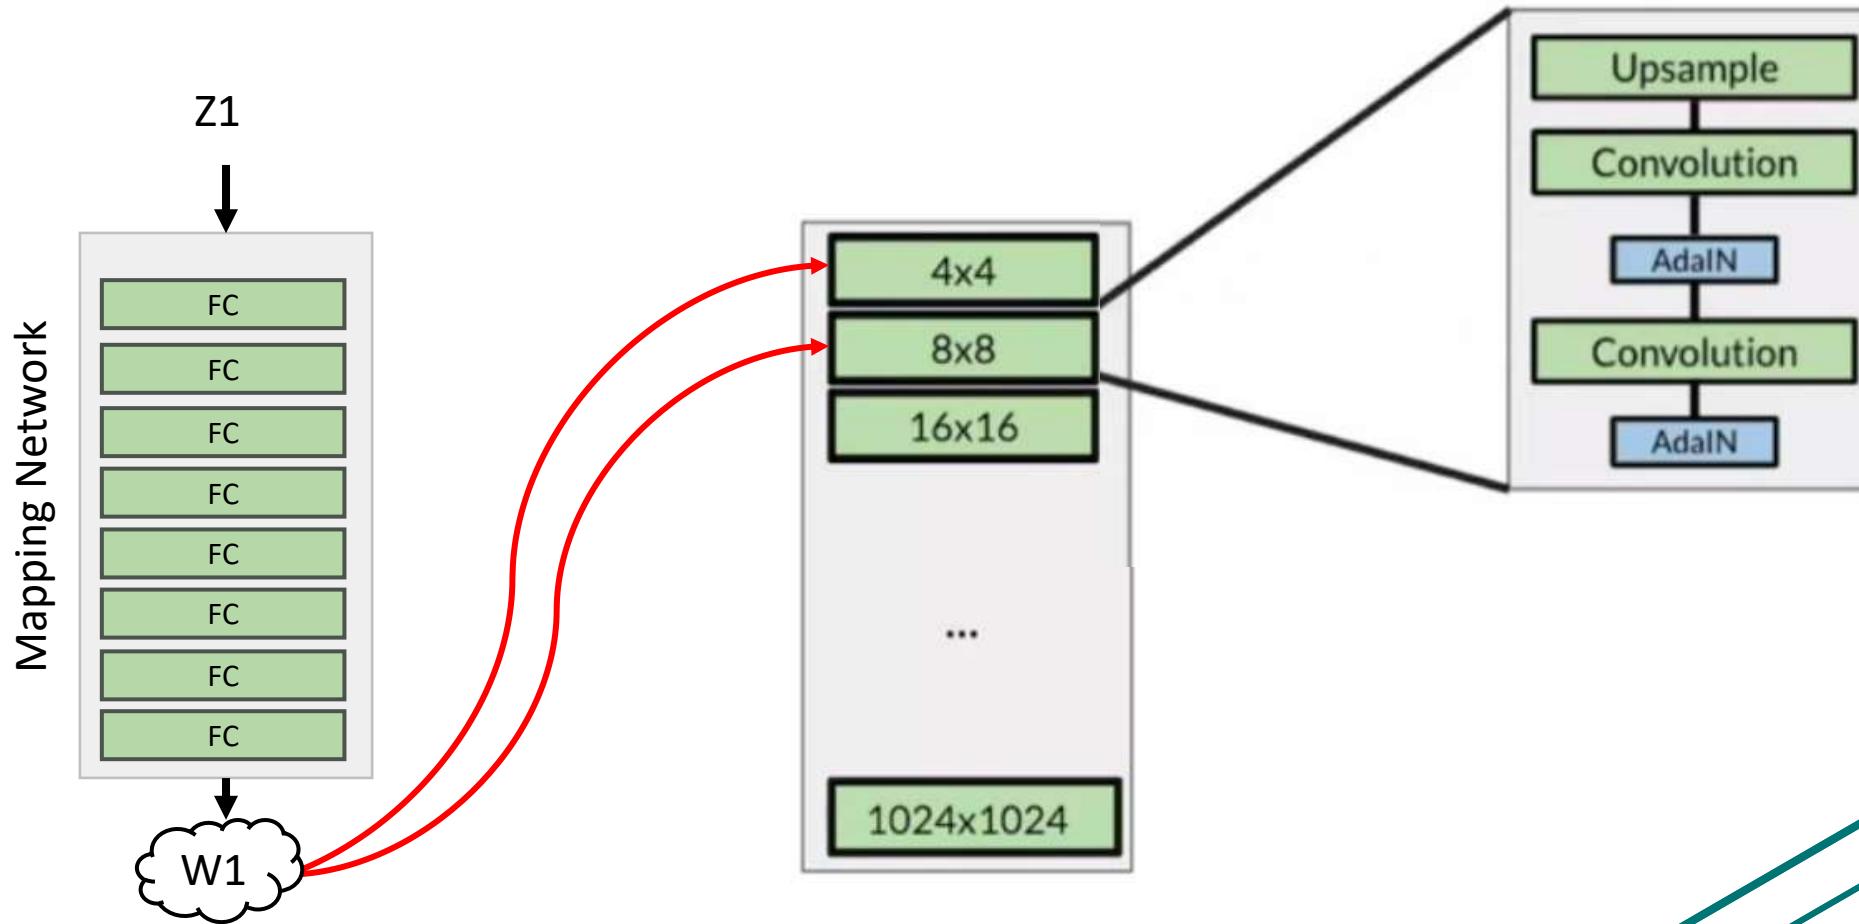
# Style mixing



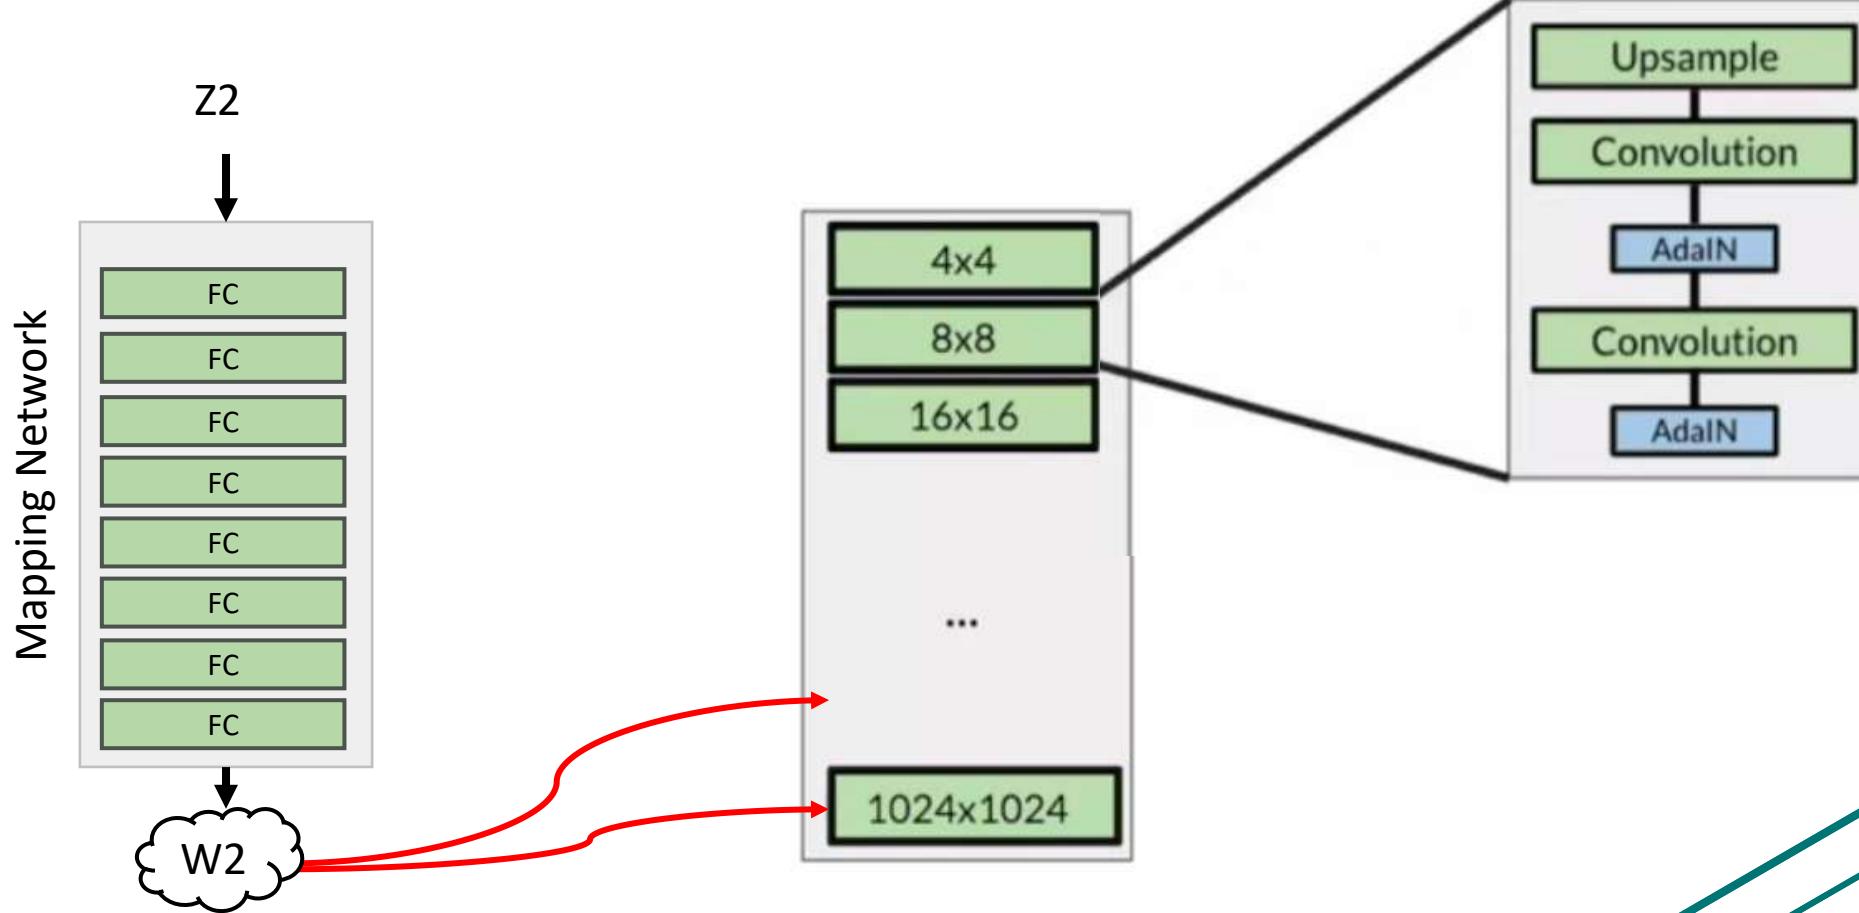
# Style mixing in context



# Style mixing in context

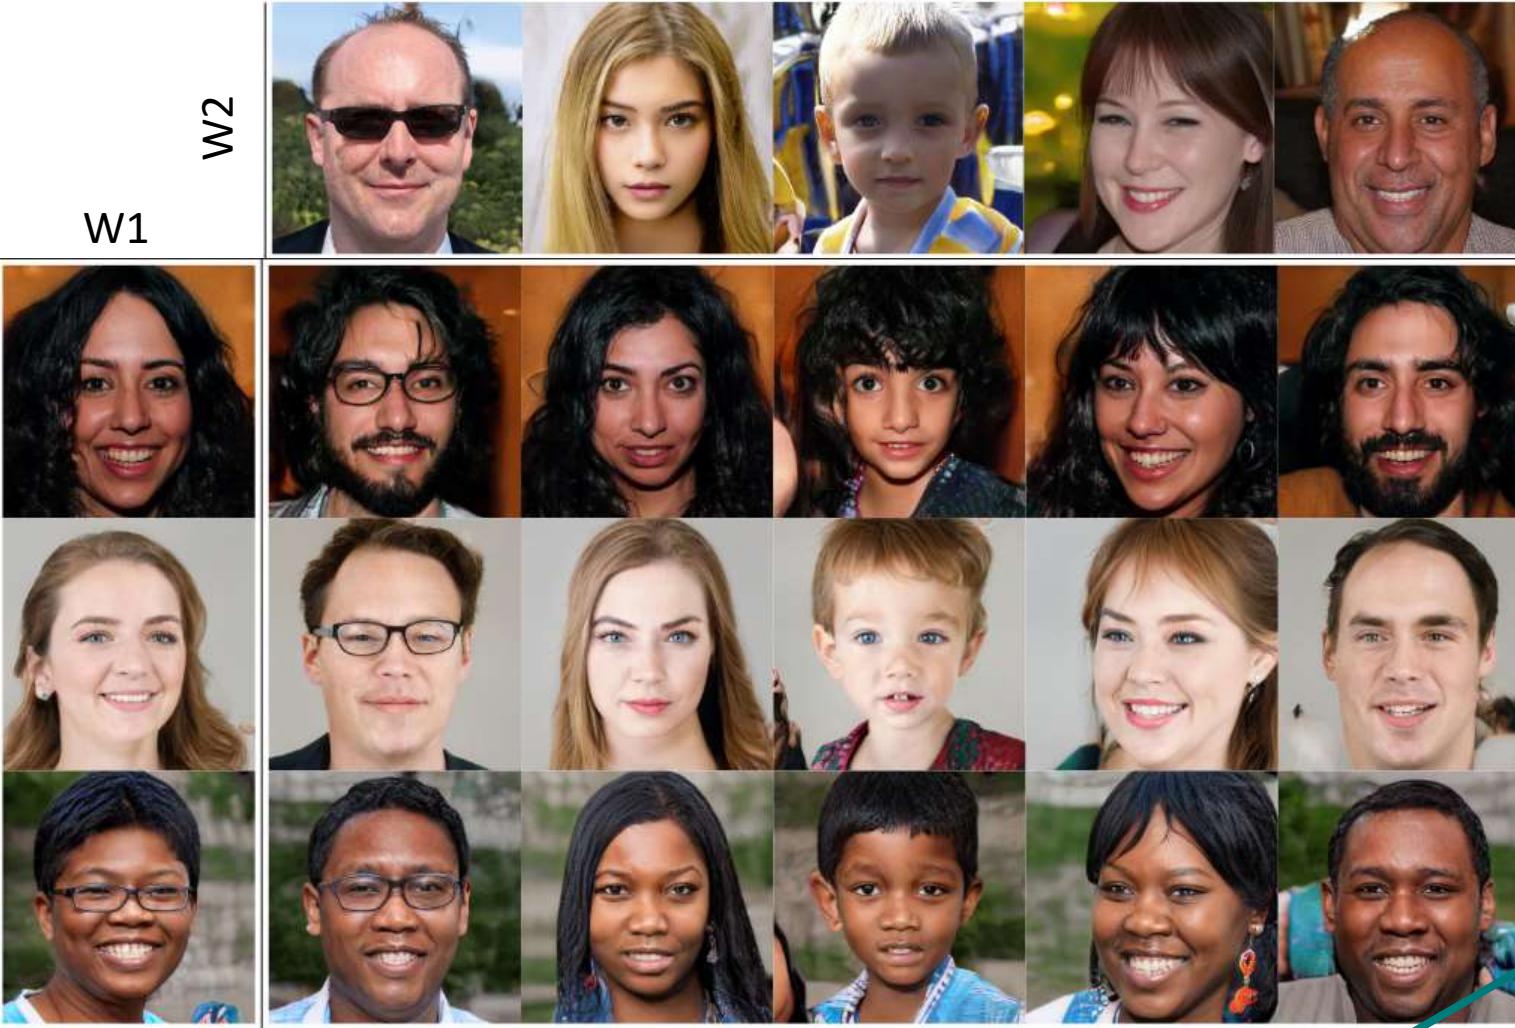


# Style mixing in context

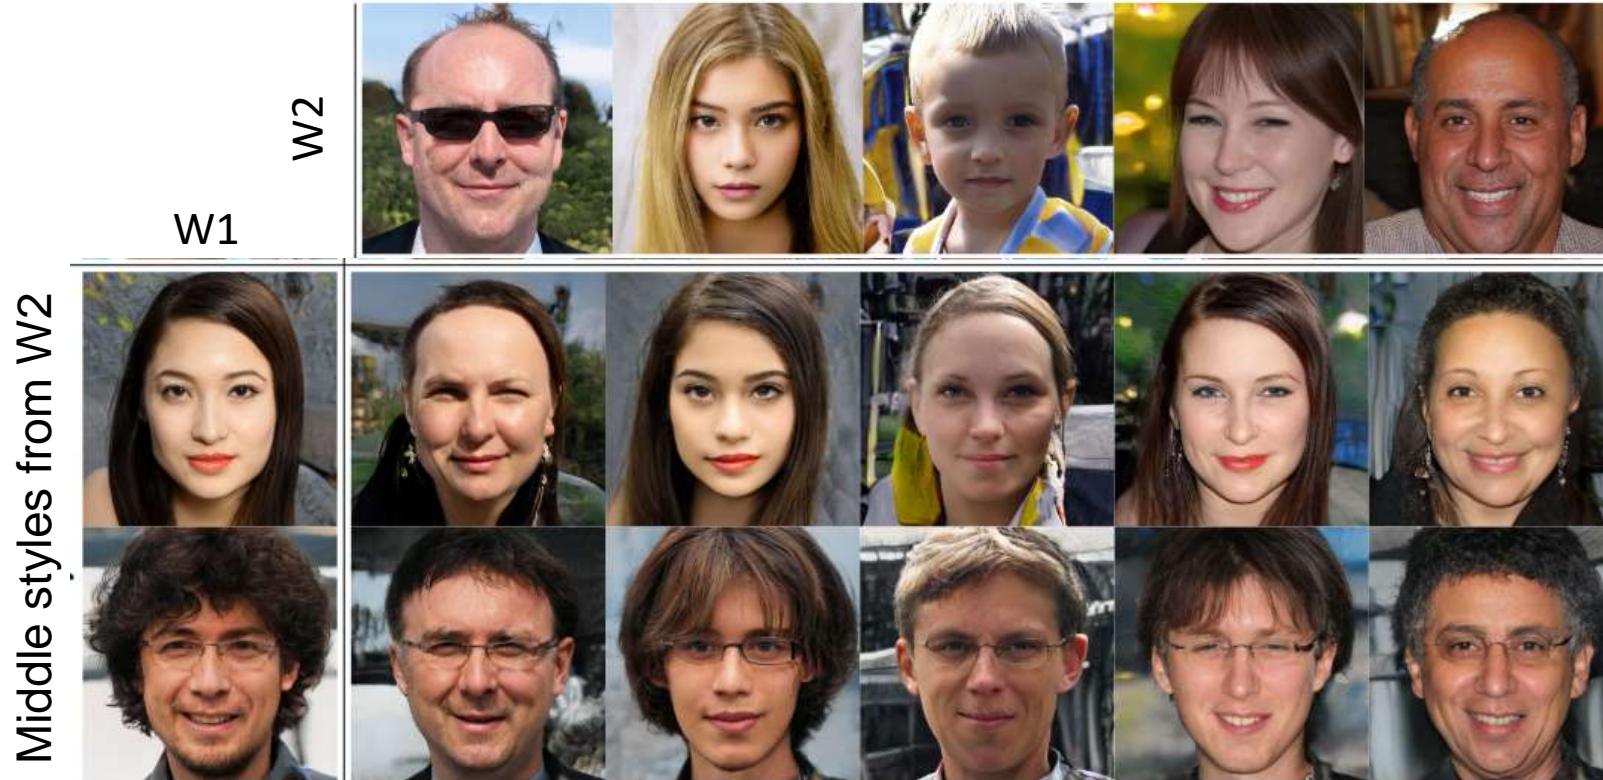


# Style mixing: Course style

Coarse styles from W2



# Style mixing: Middle styles

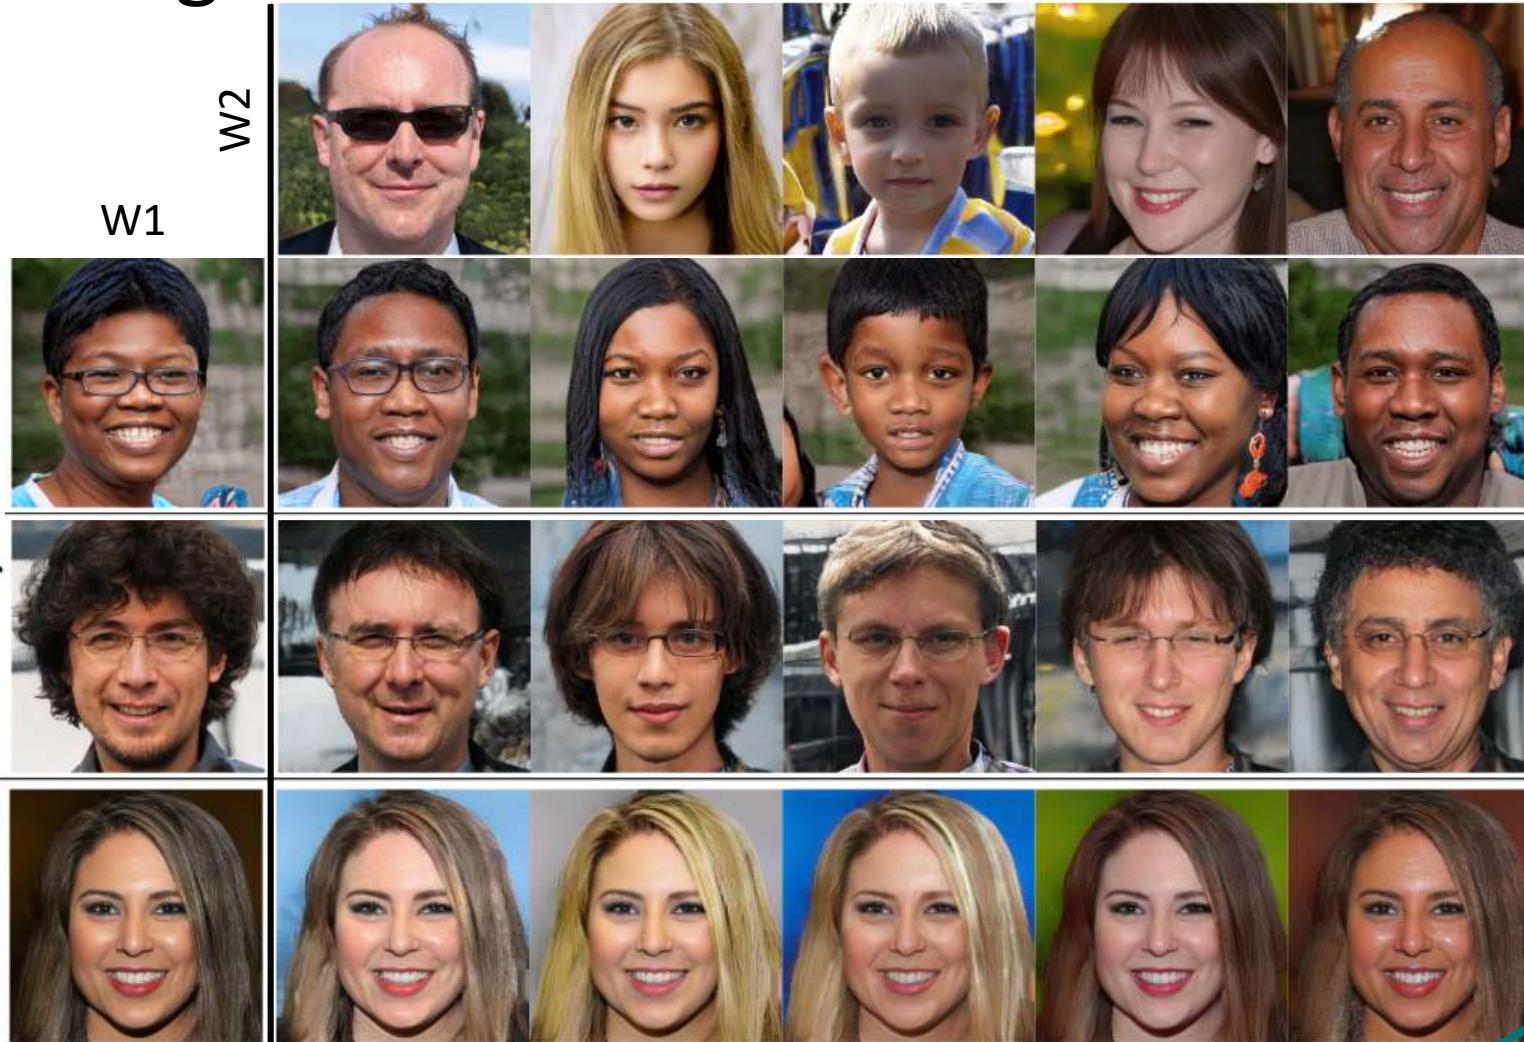


# Style mixing: Fine styles



# Style mixing

Coarse styles  
from W2



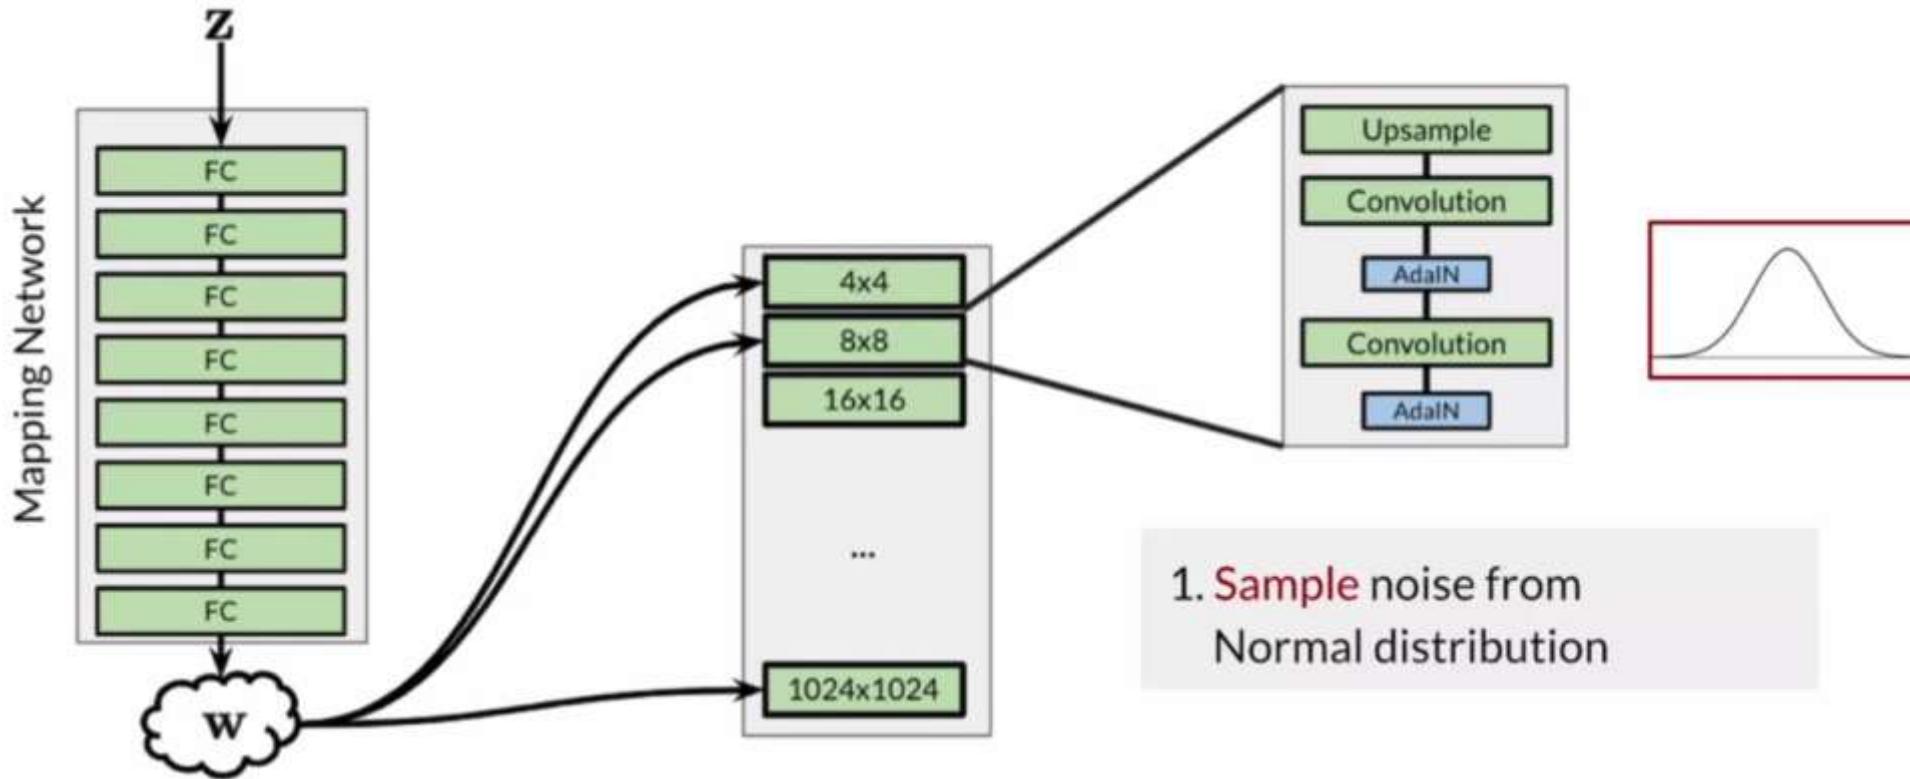
# Stochastic variation

Fine layers



Coarse layers

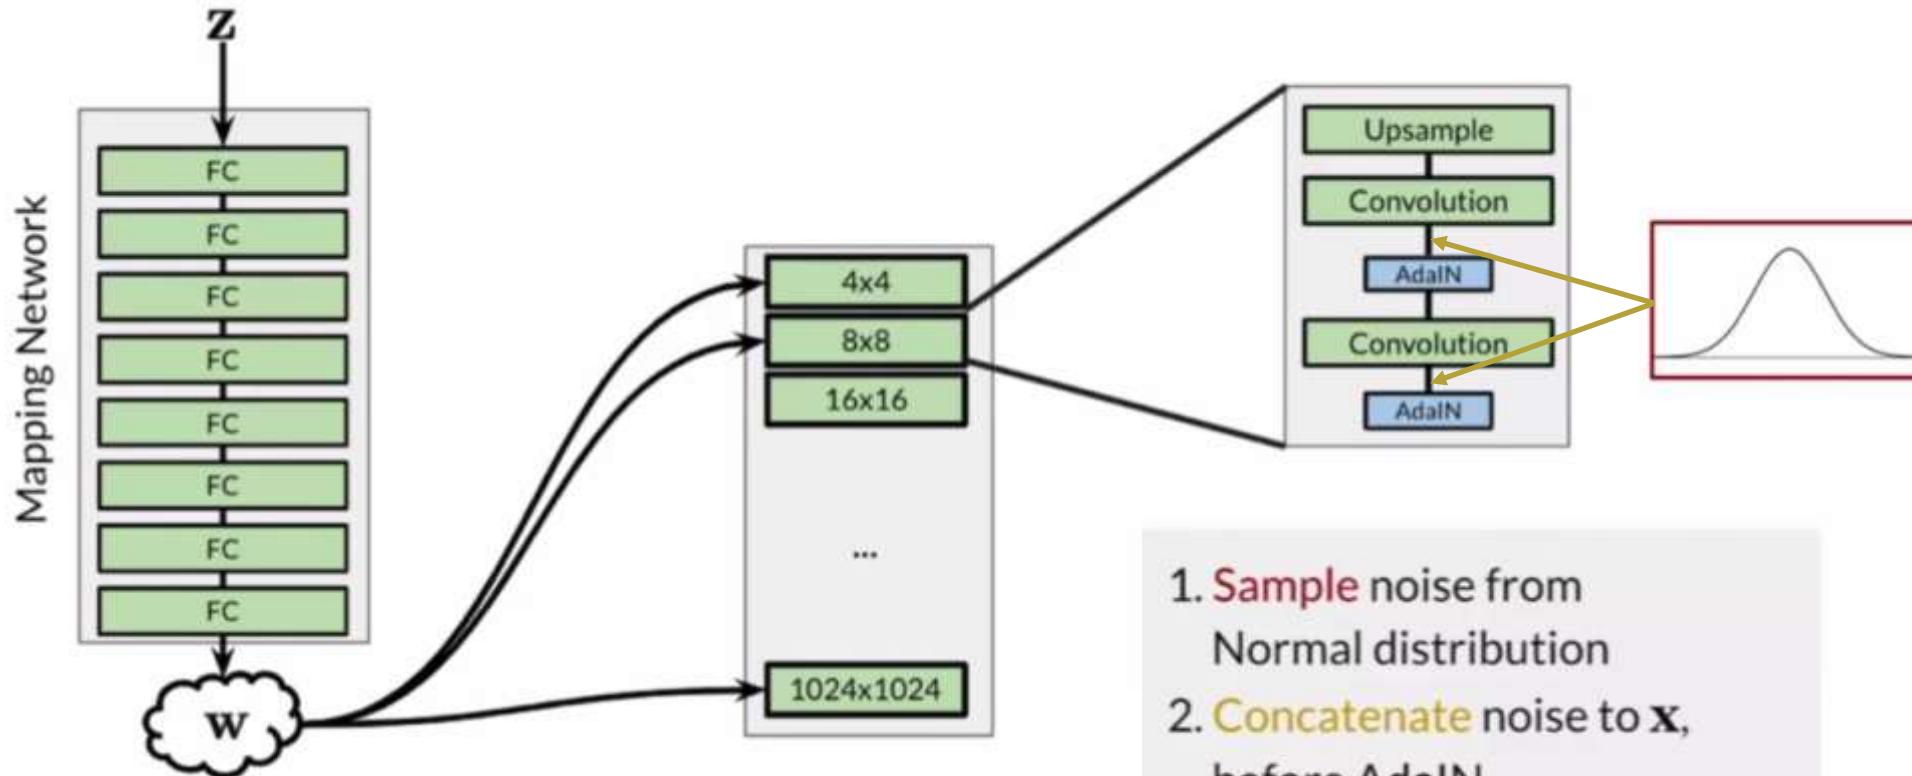
# Stochastic noise in context



1. Sample noise from  
Normal distribution



# Stochastic noise in context



# Stochastic Variation

Small details: hair strands,  
wrinkles, etc.

Different extra noise values  
Create stochastic variation



# Summary

- Style mixing increases diversity that the model sees during training
- Stochastic noise causes small variations to output
- Coarse or fineness depends where in the network style or noise is added
  - Earlier for coarser variation
  - Later for finer variation

# Implementation

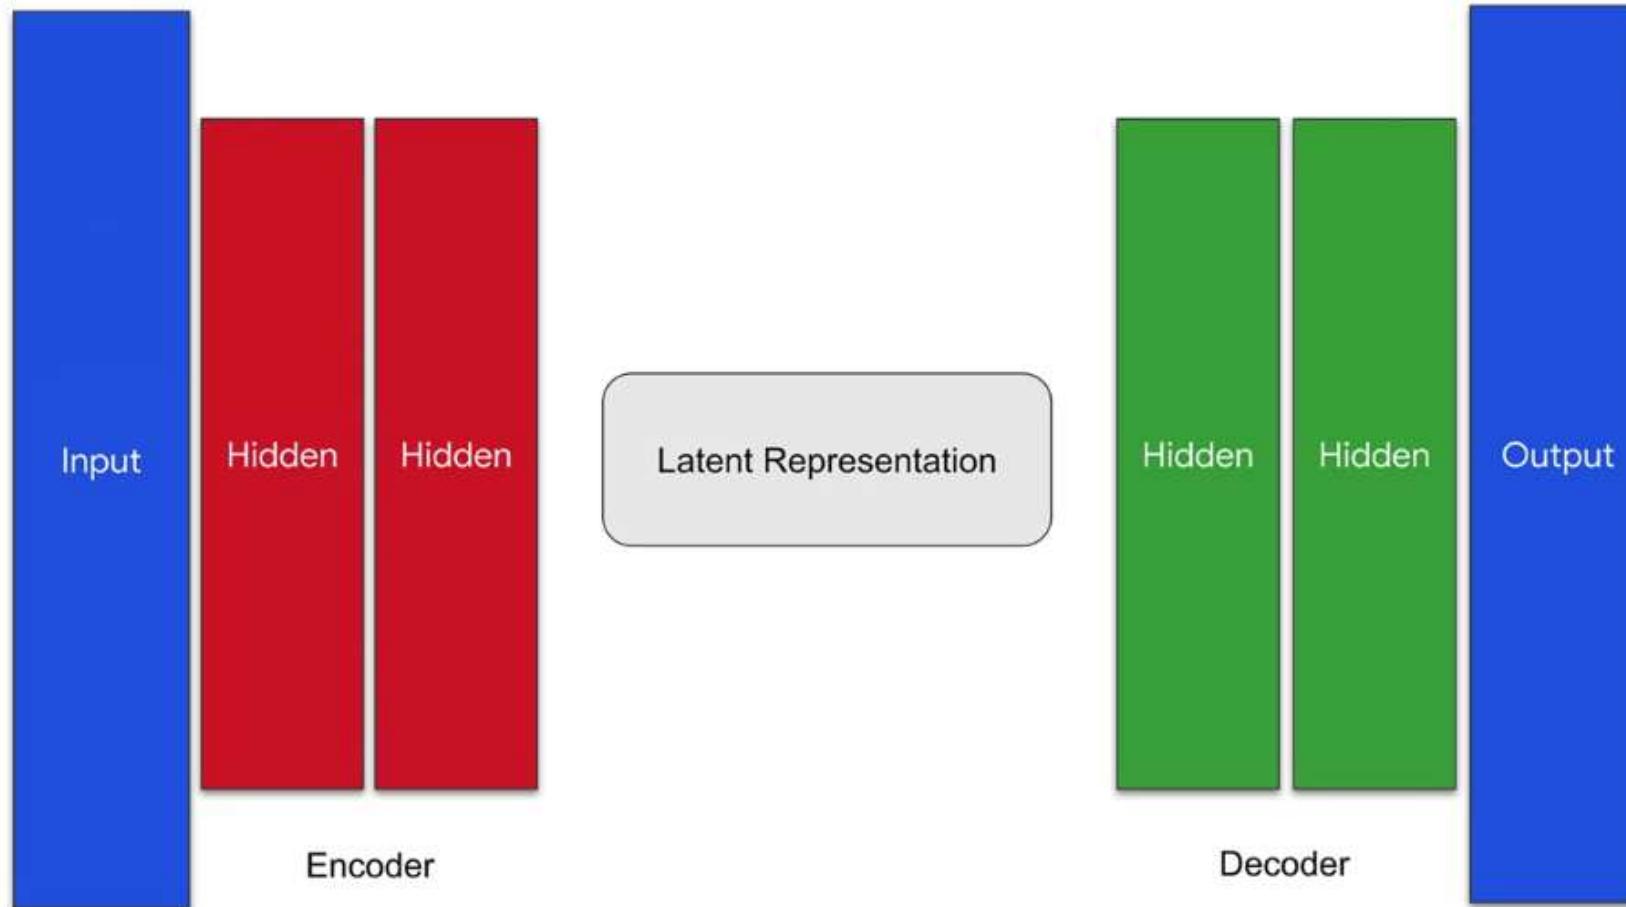
- <https://keras.io/examples/generative/stylegan/>
- <https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/stylegan.ipynb>

# VAE

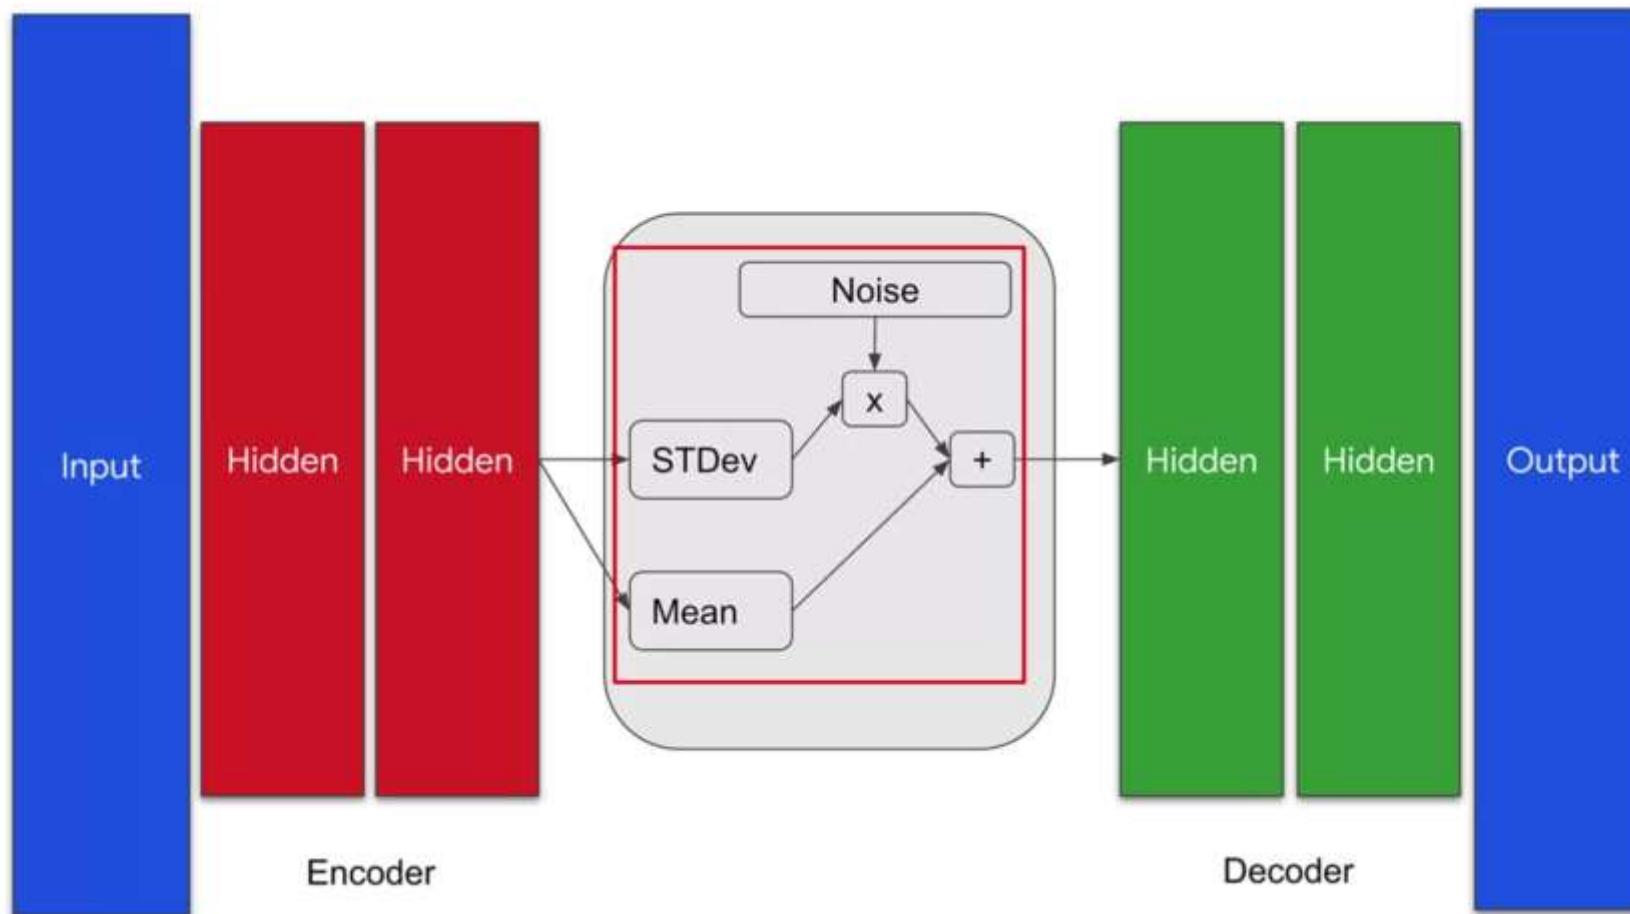
# VAE

- Unlike a traditional autoencoder, which maps the input onto a latent vector, a **VAE** maps the input data into the parameters of a probability distribution, such as the **mean** and **variance** of a Gaussian.

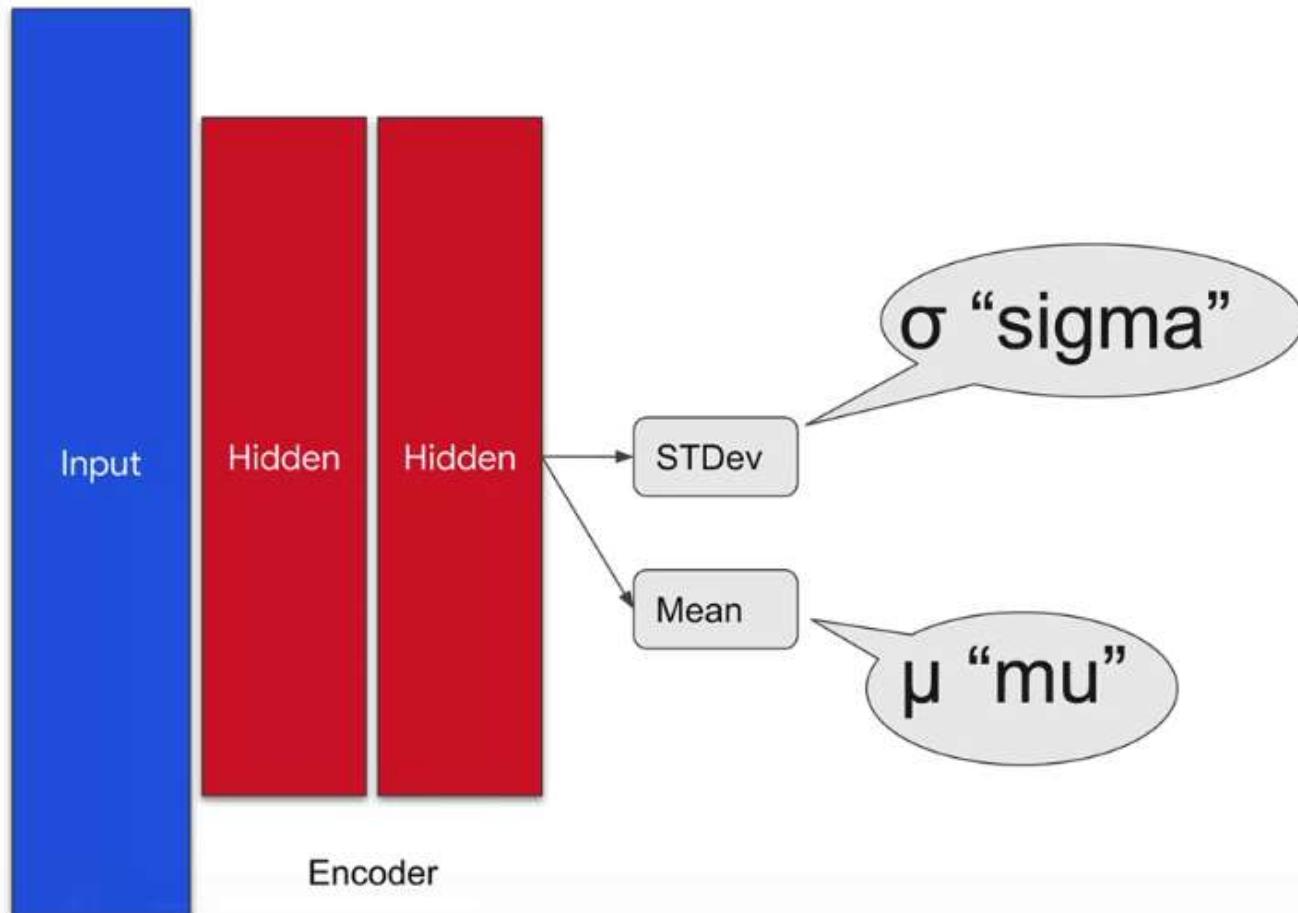
# AE



# VAE



# VAE

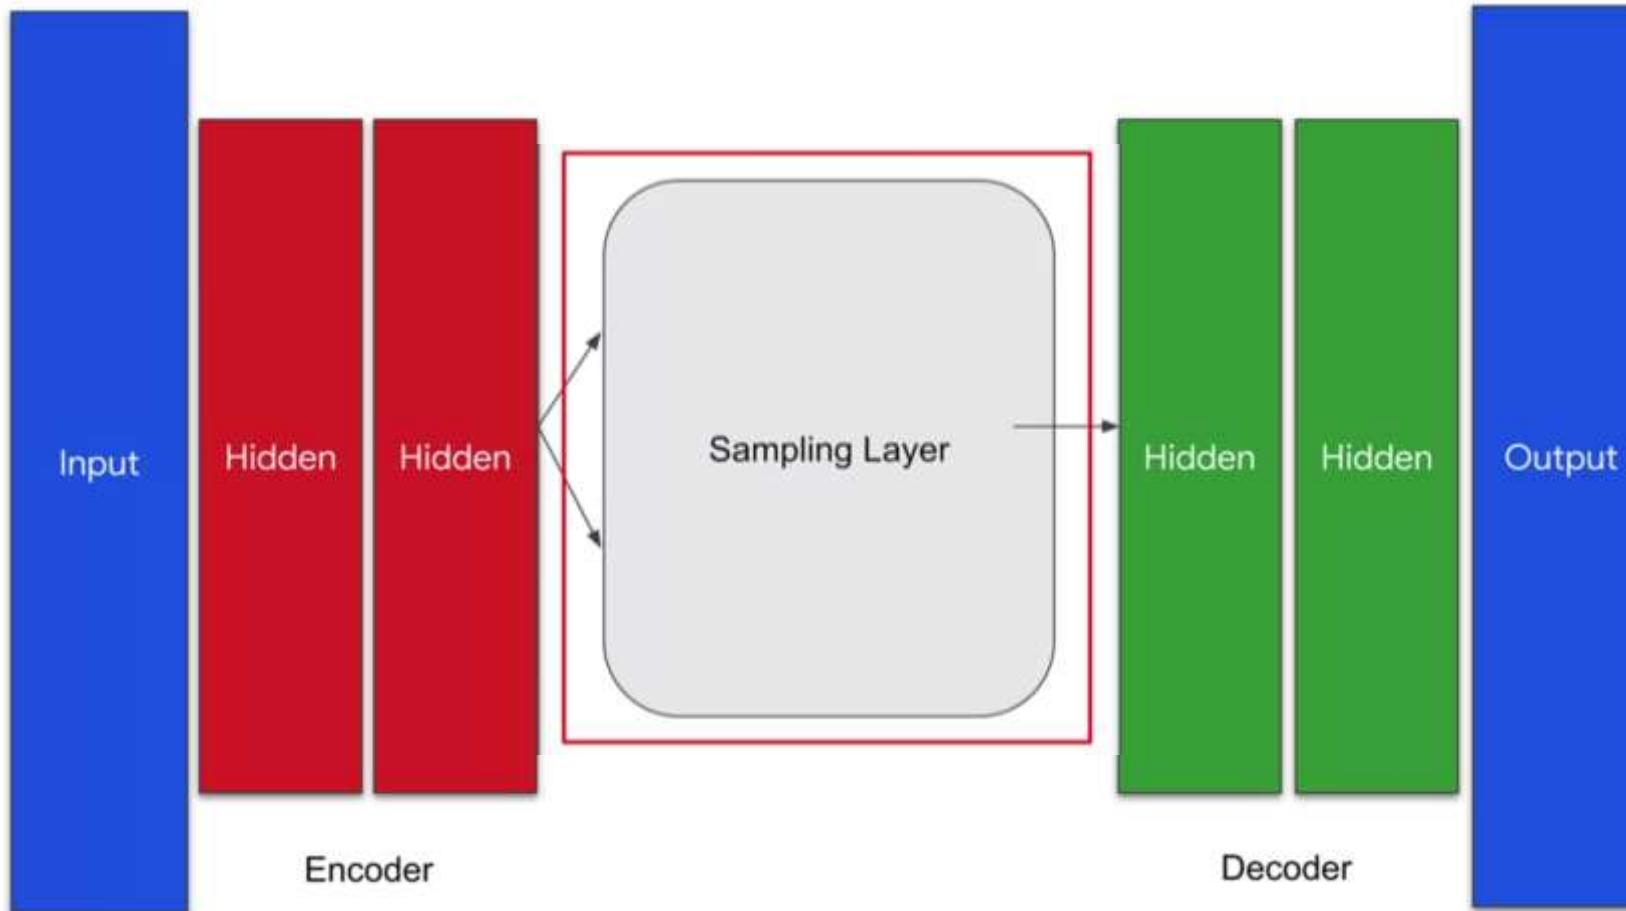


# Build the encoder

```
latent_dim = 2

encoder_inputs = keras.Input(shape=(28, 28, 1))
x = layers.Conv2D(32, 3, activation="relu", strides=2, padding="same")(encoder_inputs)
x = layers.Conv2D(64, 3, activation="relu", strides=2, padding="same")(x)
x = layers.Flatten()(x)
x = layers.Dense(16, activation="relu")(x)
z_mean = layers.Dense(latent_dim, name="z_mean")(x)
z_log_var = layers.Dense(latent_dim, name="z_log_var")(x)
z = Sampling()([z_mean, z_log_var])
encoder = keras.Model(encoder_inputs, [z_mean, z_log_var, z], name="encoder")
encoder.summary()
```

# Sampling Layer and Encoder



# Create a sampling layer

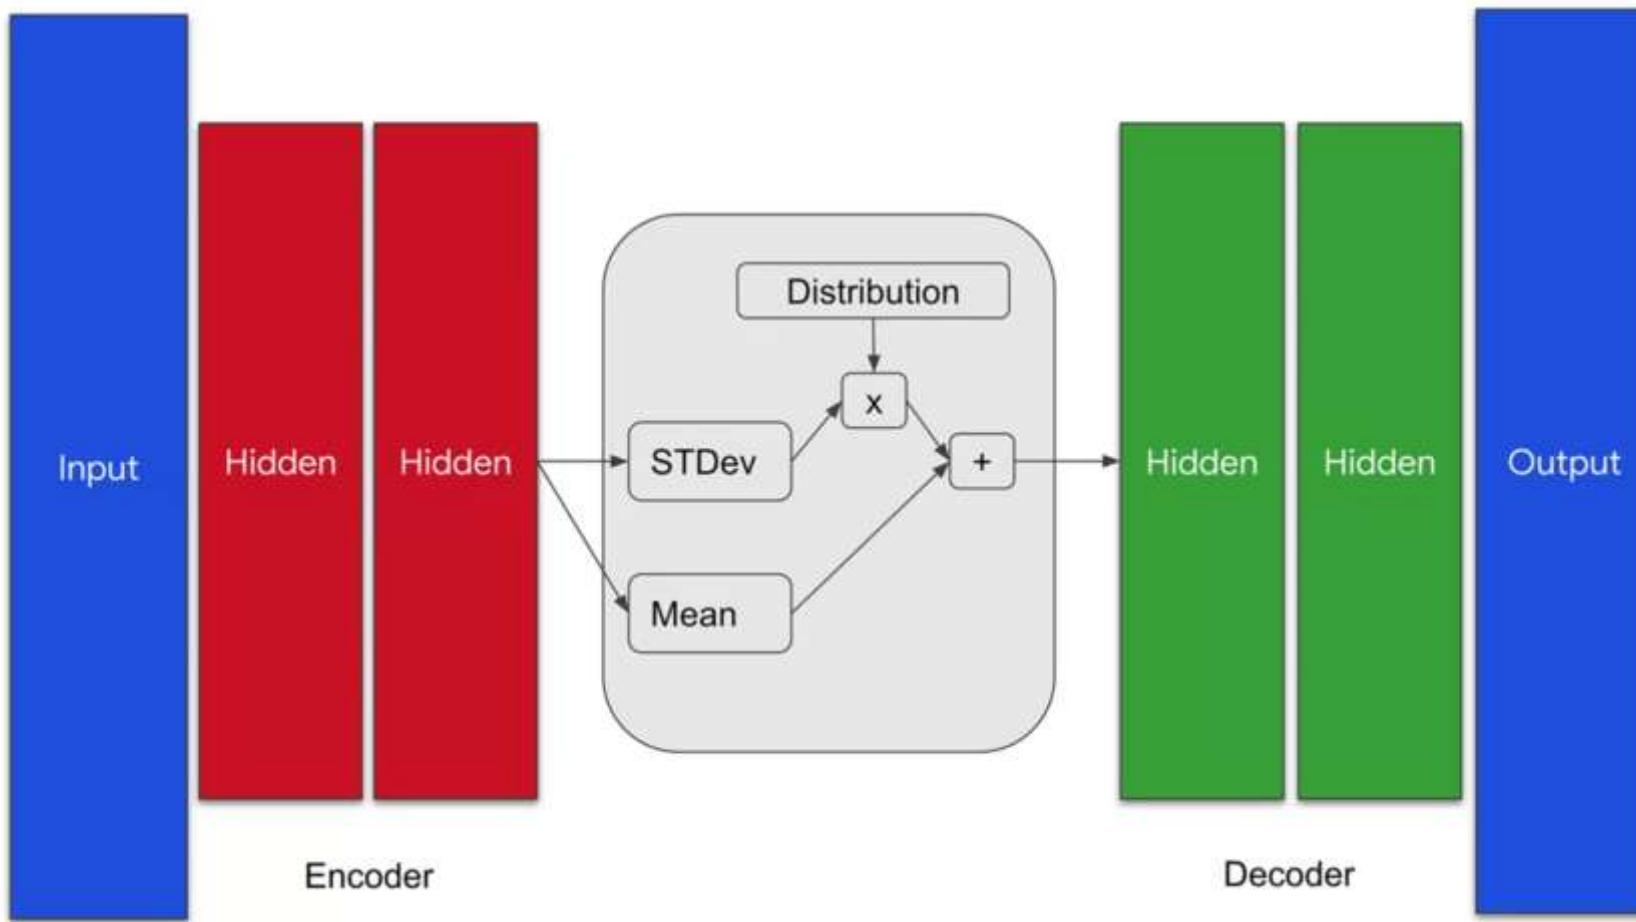
```
class Sampling(layers.Layer):
    """Uses (z_mean, z_log_var) to sample z, the vector encoding a digit."""

    def call(self, inputs):
        z_mean, z_log_var = inputs
        batch = tf.shape(z_mean)[0]
        dim = tf.shape(z_mean)[1]
        epsilon = tf.keras.backend.random_normal(shape=(batch, dim))
        return z_mean + tf.exp(0.5 * z_log_var) * epsilon
```

# Build the decoder

```
latent_inputs = keras.Input(shape=(latent_dim,))
x = layers.Dense(7 * 7 * 64, activation="relu")(latent_inputs)
x = layers.Reshape((7, 7, 64))(x)
x = layers.Conv2DTranspose(64, 3, activation="relu", strides=2, padding="same")(x)
x = layers.Conv2DTranspose(32, 3, activation="relu", strides=2, padding="same")(x)
decoder_outputs = layers.Conv2DTranspose(1, 3, activation="sigmoid", padding="same")(x)
decoder = keras.Model(latent_inputs, decoder_outputs, name="decoder")
decoder.summary()
```

# Loss Function and Model Definition



# Kullback–Leibler divergence

```
kl_loss = -0.5 * (1 + z_log_var - tf.square(z_mean) - tf.exp(z_log_var))  
kl_loss = tf.reduce_mean(tf.reduce_sum(kl_loss, axis=1))
```

# Kullback–Leibler divergence

```
kl_loss = -0.5 * (1 + z_log_var - tf.square(z_mean) - tf.exp(z_log_var))  
kl_loss = tf.reduce_mean(tf.reduce_sum(kl_loss, axis=1))
```

[https://en.wikipedia.org/wiki/Kullback%E2%80%93Leibler\\_divergence](https://en.wikipedia.org/wiki/Kullback%E2%80%93Leibler_divergence)

# منابع

- <https://www.coursera.org/specializations/generative-adversarial-networks-gans>
- [Progressive Growing of GANs for Improved Quality, Stability, and Variation](#)
- [A Style-Based Generator Architecture for Generative Adversarial Networks](#)
- [Hands-On Image Generation with TensorFlow By Soon Yau Cheong](#)
- <https://keras.io/examples/generative/>